

CSE150
Operating Systems
Lecture 3

Processes, Threads and Address Spaces
(contd.)

Threads (Review)

- Thread: Single unique execution context
 - Program Counter, Registers, Stack

Value	Description
0x00000001	Thread's object memory is not de-allocated
0x00000002	Thread's stack is not freed
0x00000004	Thread is terminated
0x00000008	Thread cannot be resumed
0x00000010	Interrupt service thread
0x00000020	Thread has used FPU/MMX/XMM
0x00000040	Timer service thread
0x00000080	Shutdown service thread


Threads (Review)

- Thread: Single unique execution context
 - Program Counter, Registers, Stack
- A thread is executing on a processor when it is resident in the processor registers.
- Program Counter (PC) register holds the address of executing instruction in the thread
- Certain registers hold the *context* of thread
 - Stack pointer holds the address of the top of stack
 - » Other conventions: Frame pointer, Heap pointer, Data
 - May be defined by the instruction set architecture or by compiler conventions
- Registers hold the root state of the thread.
 - The rest is “in memory”

The Basic Problem of Concurrency (Review)

- The basic problem of concurrency involves resources:
 - Hardware: single CPU, single DRAM, single I/O devices
 - Multiprogramming API: users think they have exclusive access to shared resources
- OS Has to coordinate all activity
 - Multiple users, I/O interrupts, ...
 - How can it keep all these things straight?
- Basic Idea: Use Virtual Machine abstraction
 - Simple machine abstraction for processes
 - Then, worry about multiplexing these abstract machines

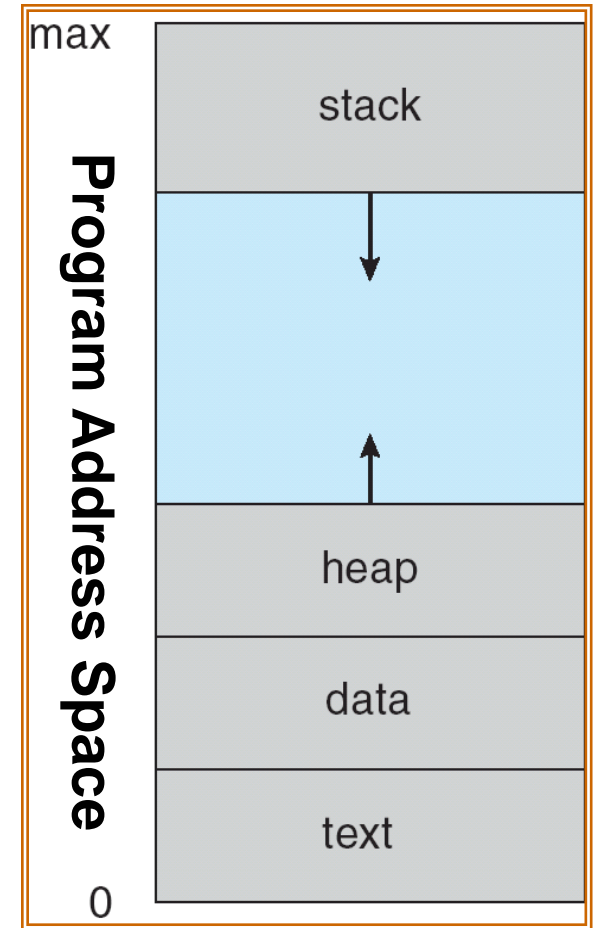
Illusion of multiple processors? (Review)



- Assume a single processor. How do we provide the illusion of multiple processors?
 - Multiplex in time!
- Each virtual “CPU” needs a structure to hold:
 - Program Counter (PC), Stack Pointer (SP)
 - Registers (Integer, Floating point, others...?)
- How switch from one CPU to the next?
 - Save PC, SP, and registers in current state block
 - Load PC, SP, and registers from new state block
- What triggers switch?
 - Timer, voluntary yield, I/O, other things

Program's Address Space (Review)

- Address space \Rightarrow the set of accessible addresses + state associated with them:
 - For a 32-bit processor there are $2^{32} = 4$ billion addresses.



Today: Fundamental OS Concepts (contd.)

- Thread
 - Single unique execution context: fully describes program state
 - Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Programs execute in an *address space* that is distinct from the memory space of the physical machine
- Process
 - An instance of an executing program is *a process consisting of an address space and one or more threads of control*
- Dual mode operation / Protection
 - Only the “system” has the ability to access certain resources
 - The OS and the hardware are protected from user programs and user programs are isolated from one another by *controlling the translation* from program virtual addresses to machine physical addresses

Third OS Concept: Process

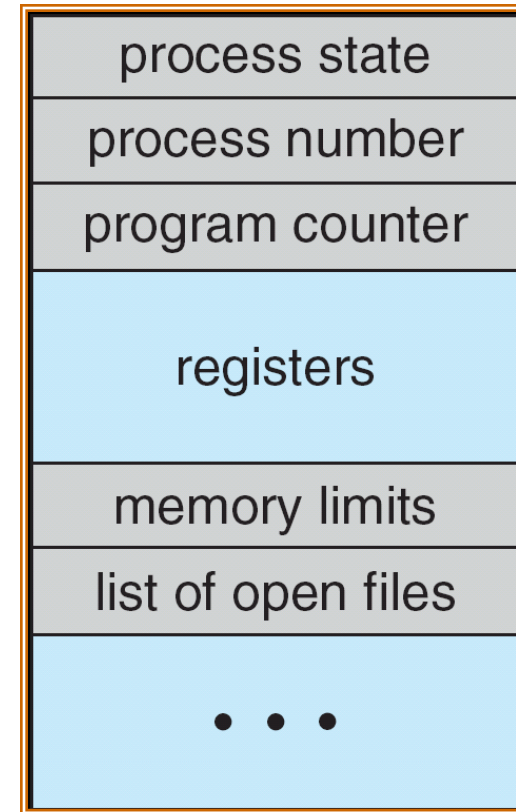
- **Process:** execution environment with Restricted Rights
 - Address Space with One or More Threads
 - Owns memory (address space)
 - Owns file descriptors, file system context, ...
 - Encapsulate one or more threads sharing process resources
- Why processes?
 - Protected from each other!
 - OS Protected from them
 - Threads more efficient than processes (later)
- Fundamental tradeoff between protection and efficiency
 - Communication easier *within* a process
 - Communication harder *between* processes

Traditional UNIX Process

- Process: *Operating system abstraction to represent what is needed to run a single program*
 - Often called a “HeavyWeight Process”
 - Formally: a single, sequential stream of execution in its *own* address space
- Two parts:
 - Sequential Program Execution Stream
 - » Code executed as a *single, sequential* stream of execution
 - » Includes State of CPU registers
 - Protected Resources:
 - » Main Memory State (contents of Address Space)
 - » I/O state (i.e. file descriptors)

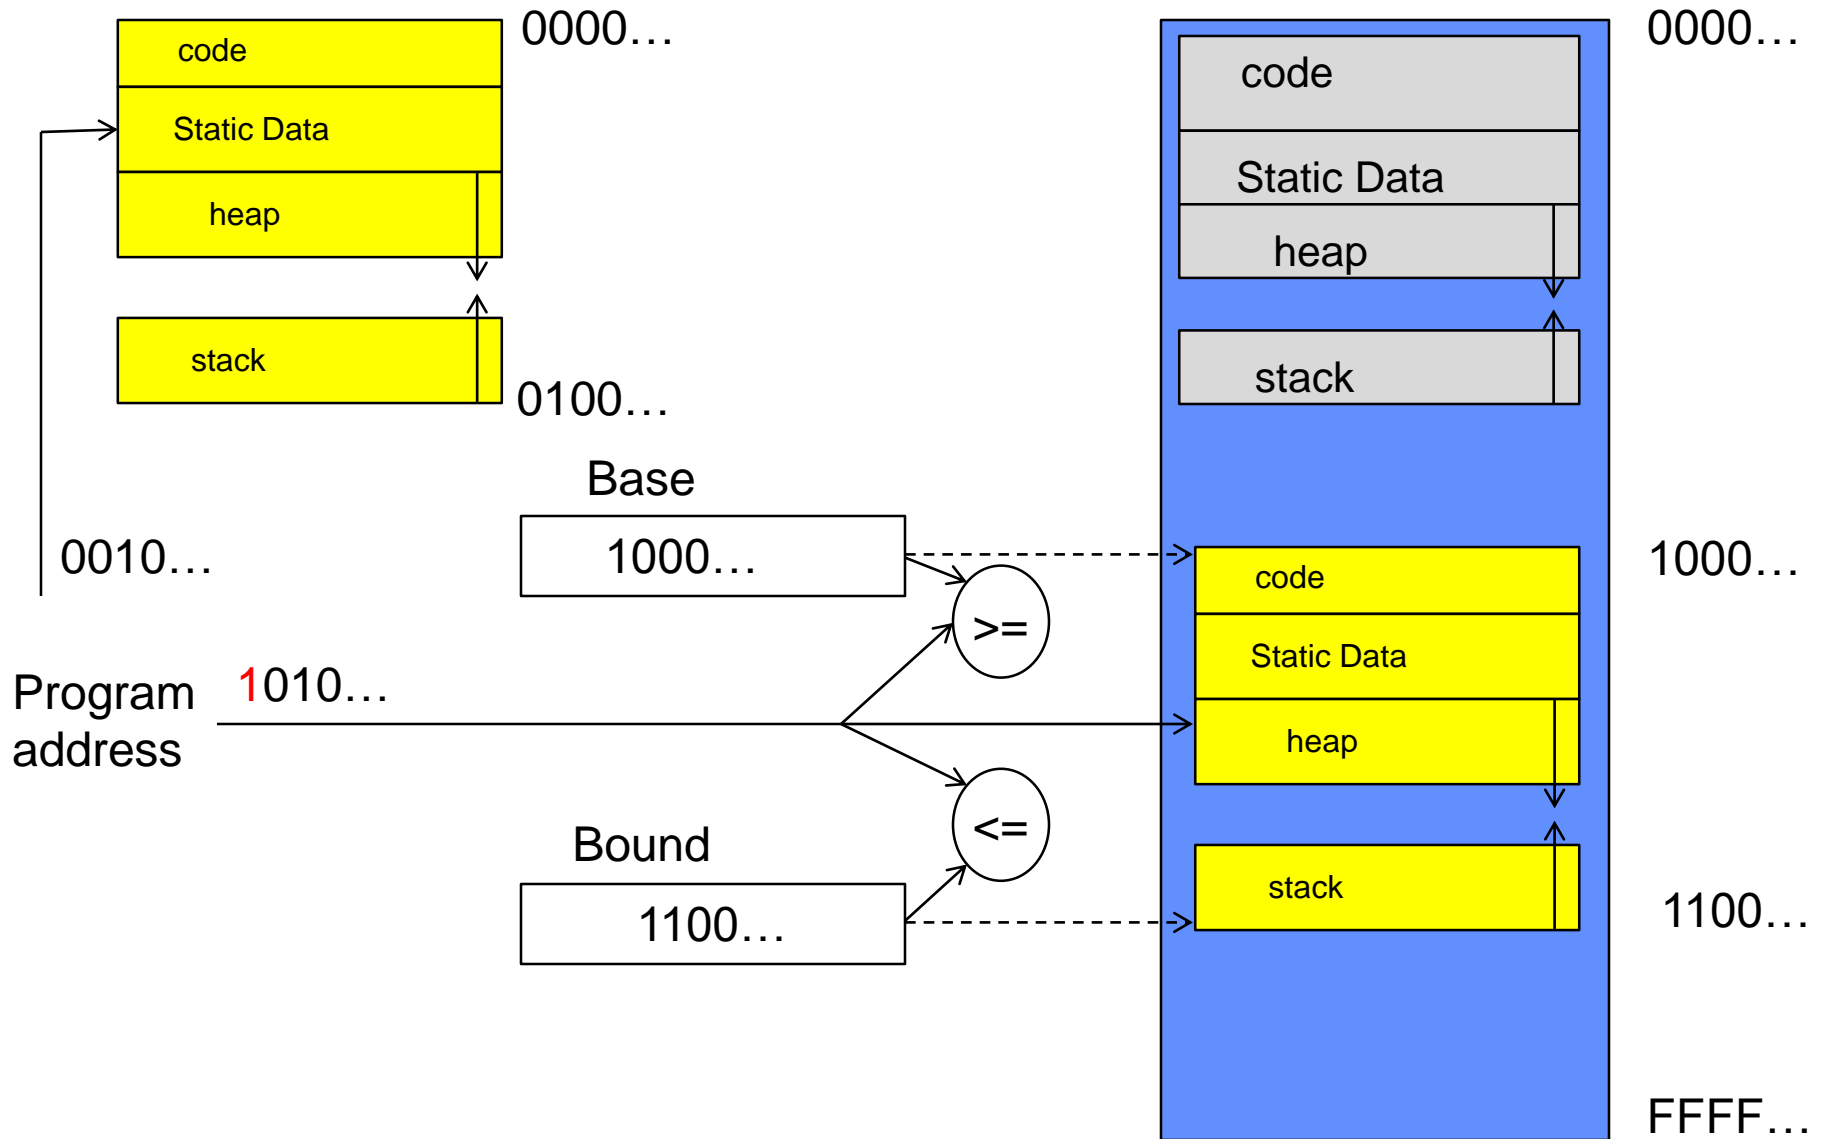
How do we multiplex processes?

- The current state of process held in a process control block (PCB):
 - This is a “snapshot” of the execution and protection environment
 - Only one PCB active at a time
- Give out CPU time to different processes (**Scheduling**):
 - Only one process “running” at a time
 - Give more time to important processes
- Give pieces of resources to different processes (**Protection**):
 - Controlled access to non-CPU resources
 - Sample mechanisms:
 - » Memory Mapping: Give each process their own address space
 - » Kernel/User duality: Arbitrary multiplexing of I/O through system calls

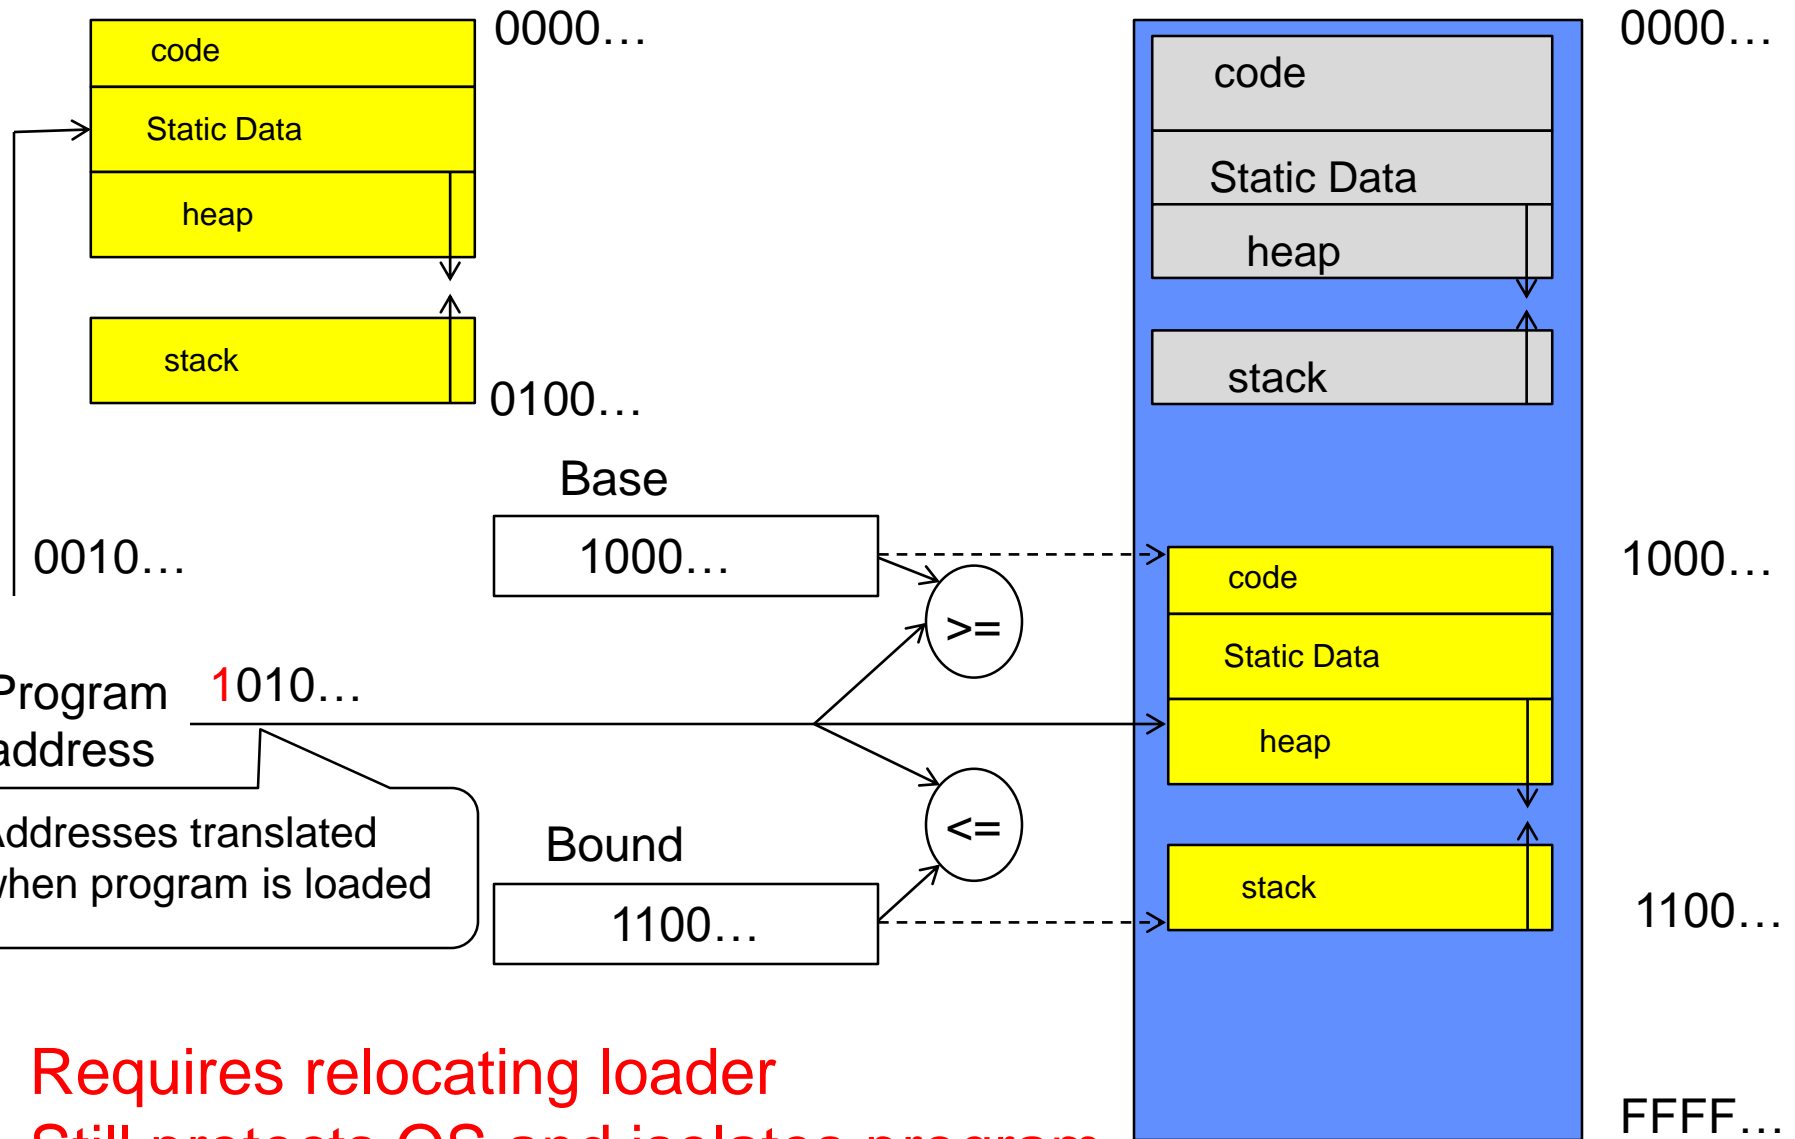


**Process
Control
Block**

Simple Protection: Base and Bound (B&B)

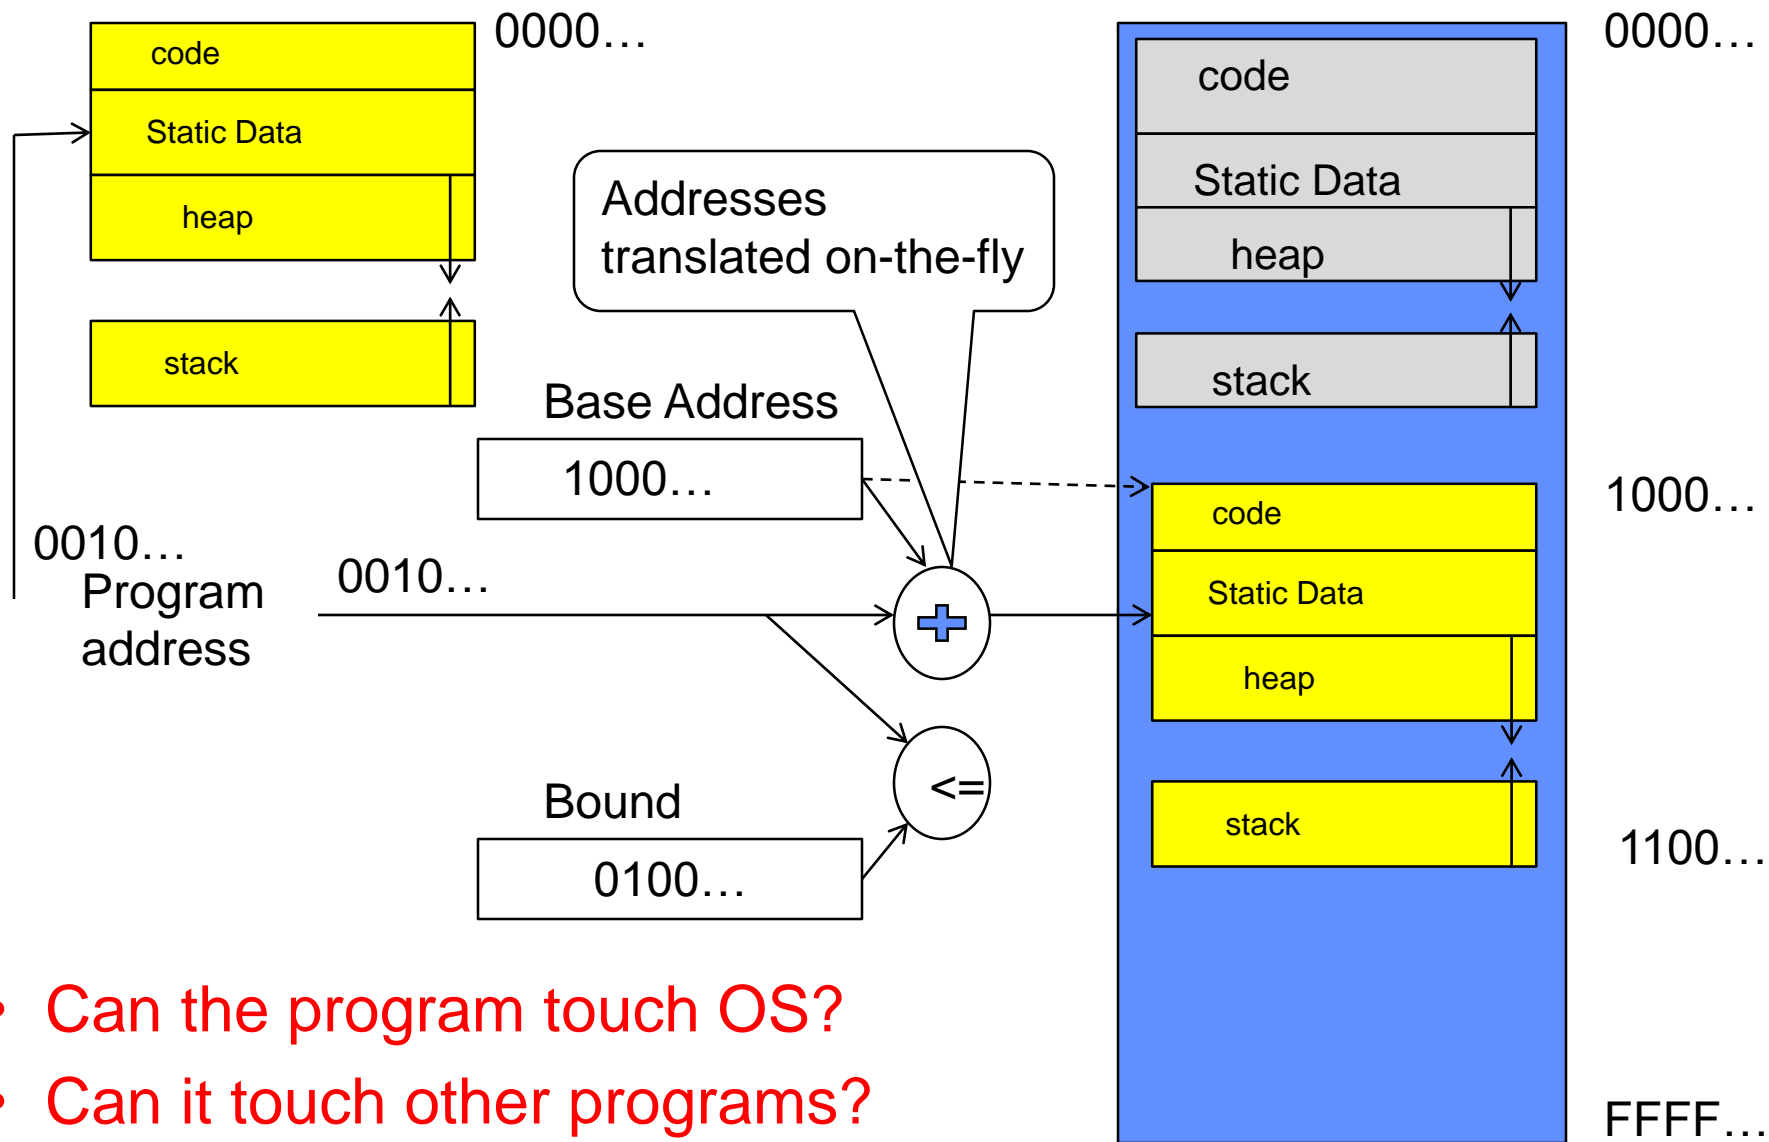


Simple Protection: Base and Bound (B&B)



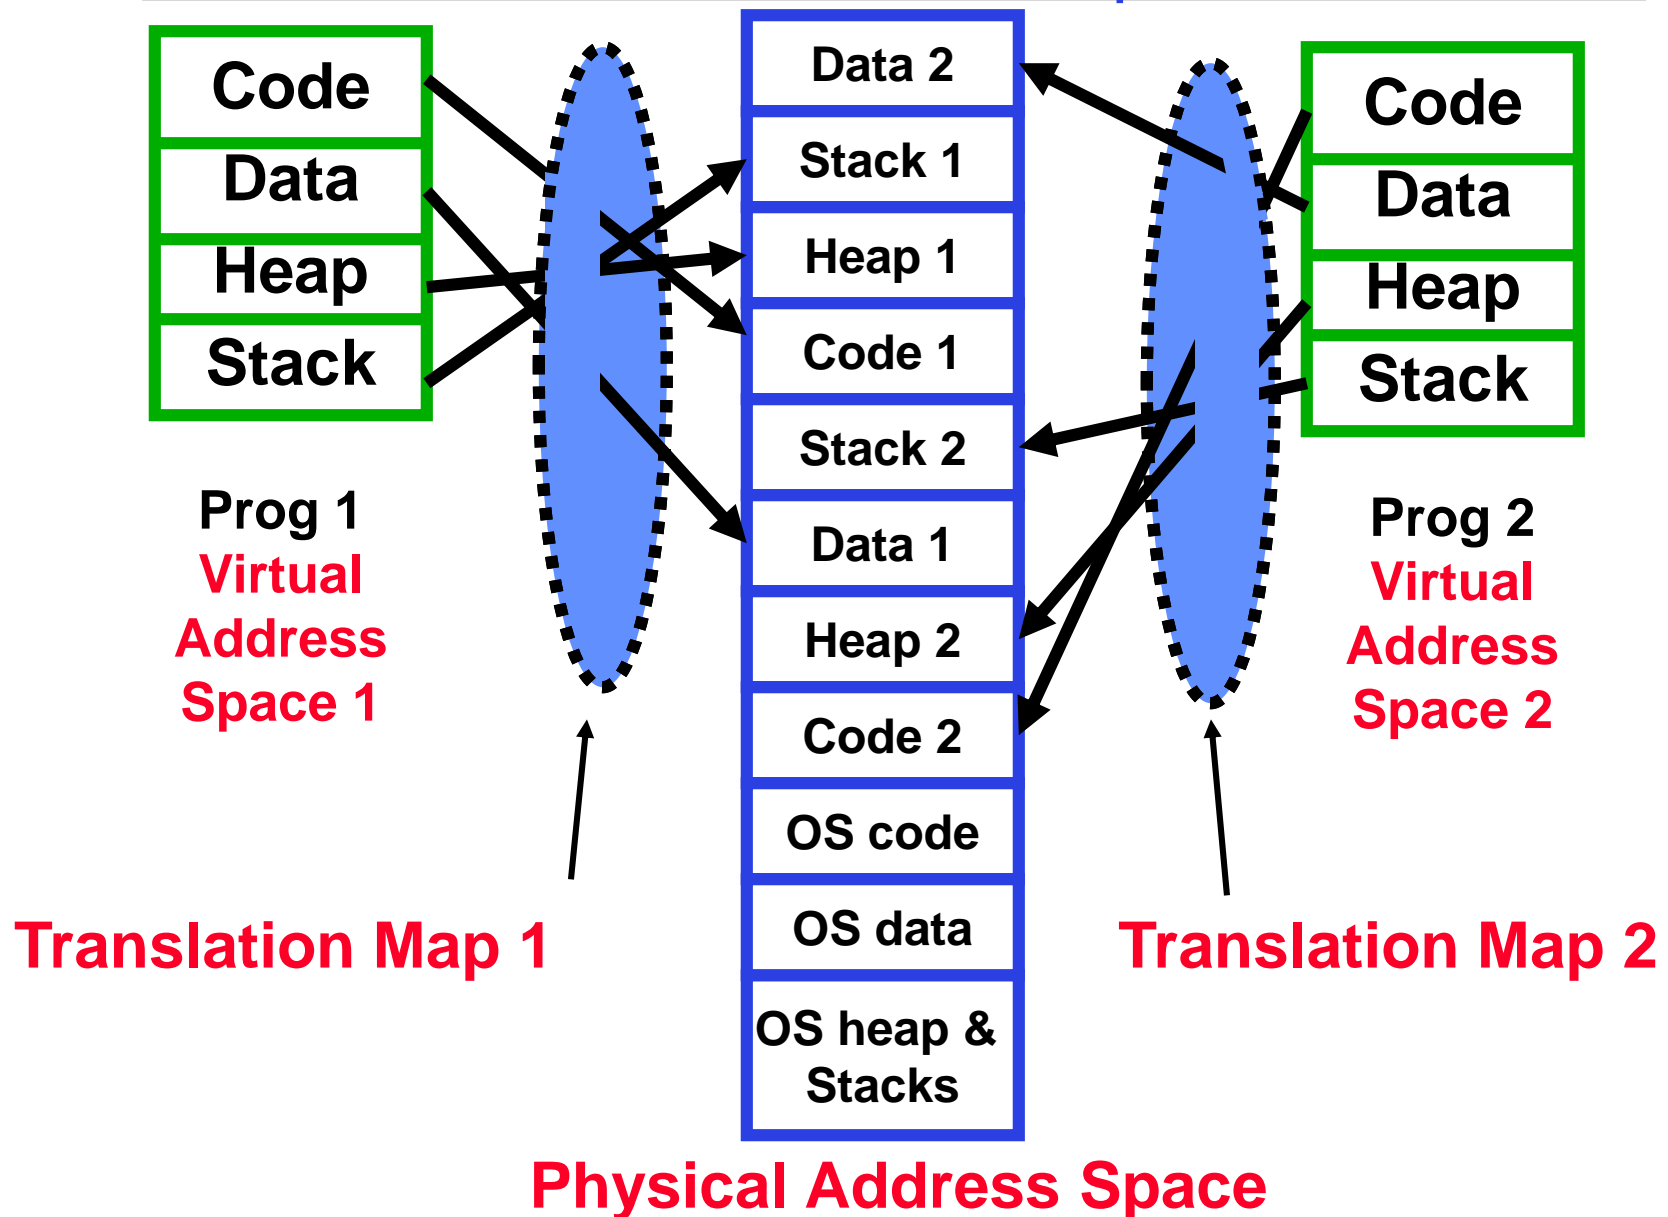
- Requires relocating loader
- Still protects OS and isolates program

A simple address translation with Base and Bound

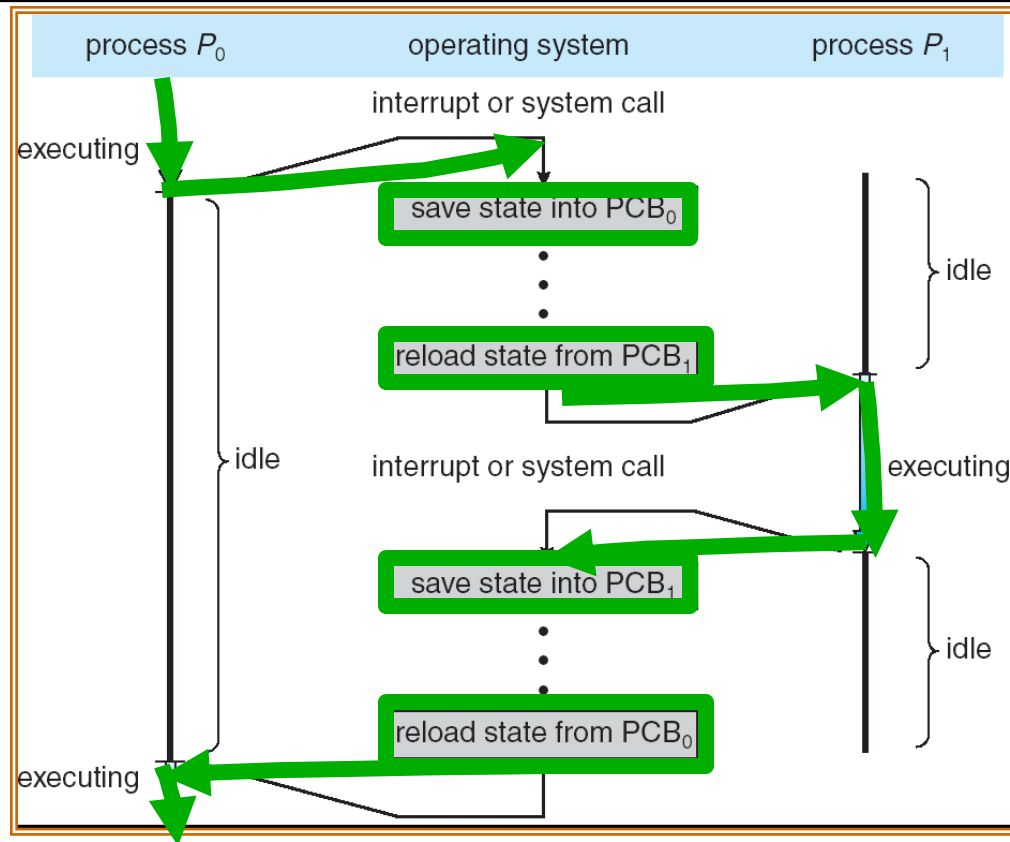


- Can the program touch OS?
- Can it touch other programs?

Providing Illusion of Separate Address Space: Load new Translation Map on Switch



CPU Switch From Process to Process




- This is also called a “context switch”
- Code executed in kernel above is overhead
 - Overhead sets minimum practical switching time

Fourth OS Concept: Dual Mode Operation

- **Hardware** provides at least two modes:
 - “Kernel” mode (or “supervisor” or “protected”)
 - “User” mode: Normal programs executed
- What is needed in the hardware to support “dual mode” operation?
 - a bit of state (user/system mode bit)
 - Certain operations / actions only permitted in system/kernel mode
 - » In user mode they fail
 - **User→Kernel** transition sets system mode AND saves the user PC
 - » Operating system code carefully puts aside user state then performs the necessary operations
 - **Kernel→User** transition clears system mode AND restores appropriate user PC
 - » return-from-interrupt

User/Kernel (Privileged) Mode

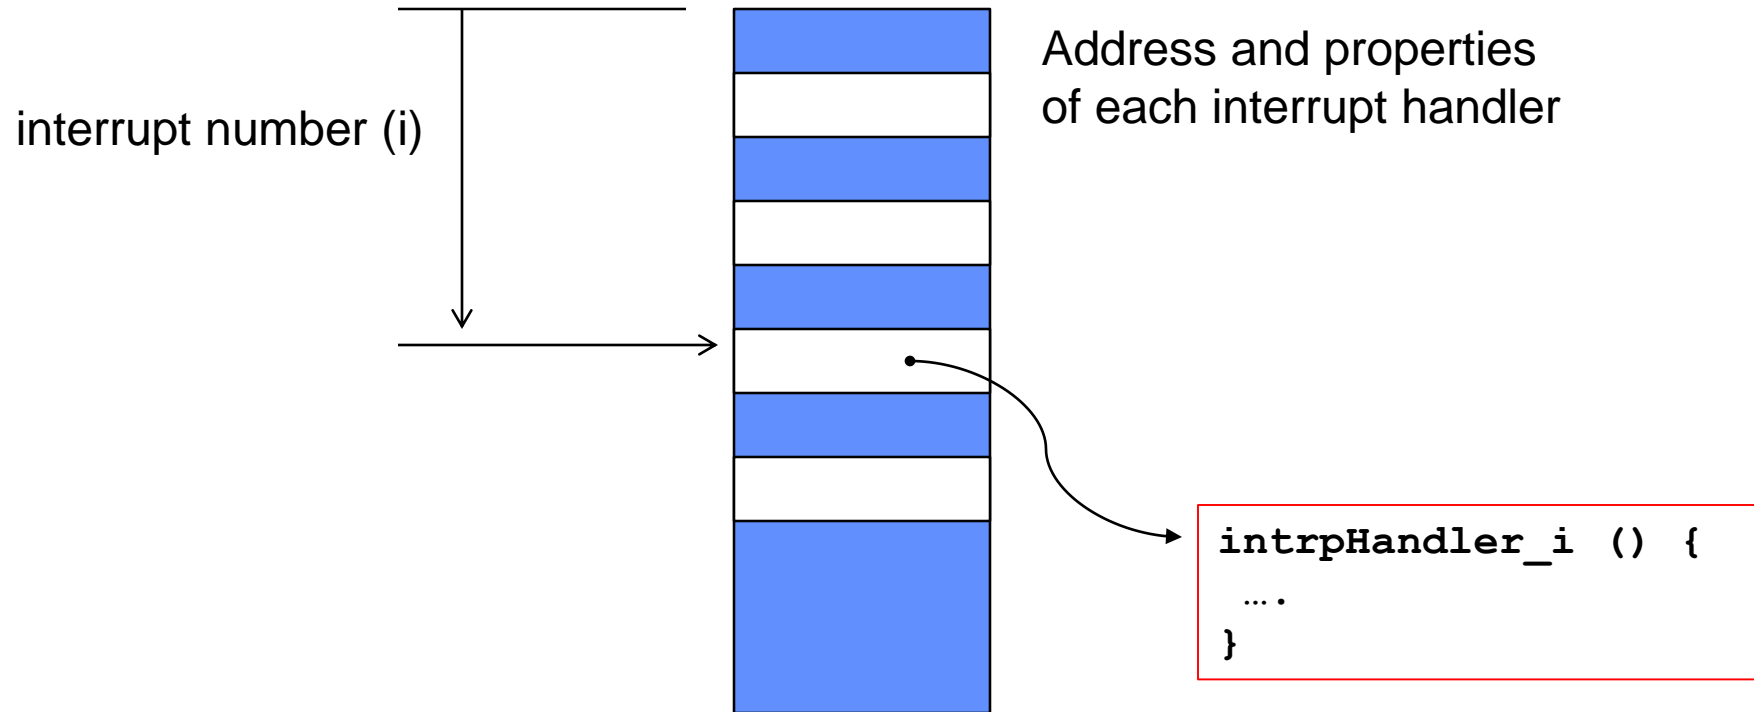


3 types of Mode Transfer

- Syscall
 - Process requests a system service, e.g., exit
 - Like a function call, but “outside” the process
 - Does not have the address of the system function to call
 - Put the syscall id and args in registers and exec syscall
- Interrupt
 - External asynchronous event triggers context switch
 - eg. Timer, I/O device
 - Independent of user process
- Trap or Exception
 - e.g., Protection violation (segmentation fault), Divide by zero, ...
- All 3 are an UNPROGRAMMED CONTROL TRANSFER
 - Where does it go?

How do we get the system target address of the “unprogrammed control transfer?”

Interrupt Vector



Summary

- Two more OS concepts
 - Processes (Threads + Address Space... more or less)
 - Dual Mode Operation (Hardware support for Protection)
- Base and Bounds
 - A simple protection mechanism used in Cray 1 using two registers
 - Modern systems use a translation map per process
- Book talks about processes
 - When this concerns concurrency, really talking about thread portion of a process
 - When this concerns protection, talking about address space portion of a process
- User/Kernel Mode
 - Must (carefully) control User→Kernel transitions
 - Transitions controlled via Syscalls, Interrupts, Traps/Exceptions