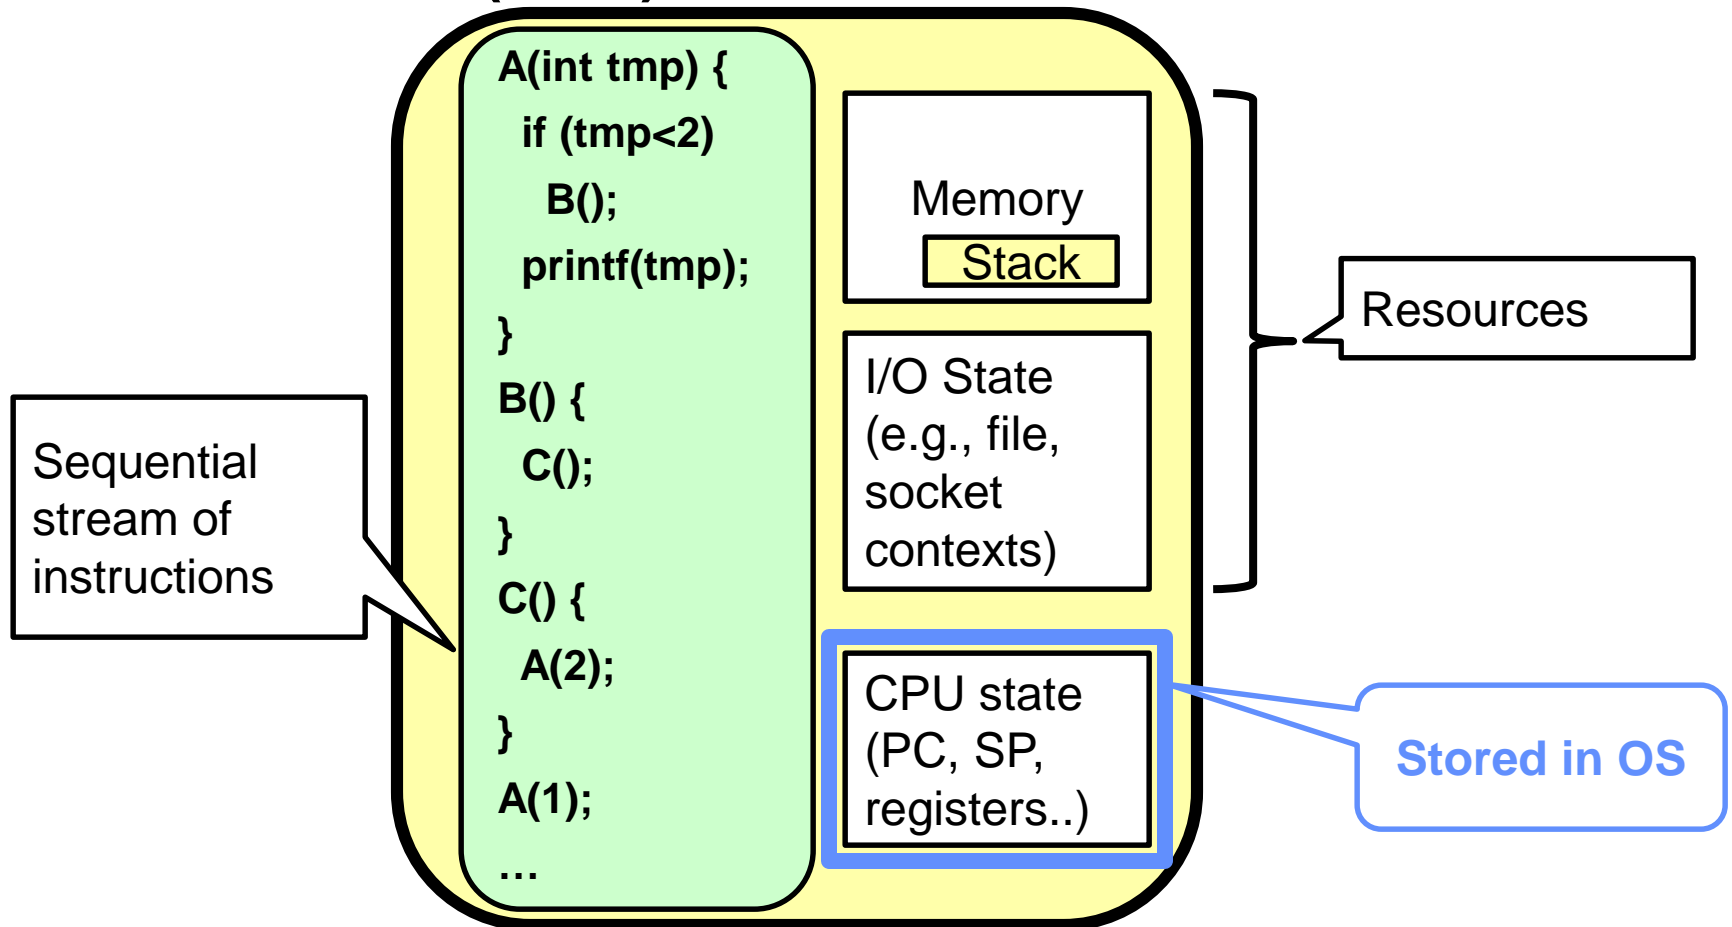


CSE150
Operating Systems
Lecture 5

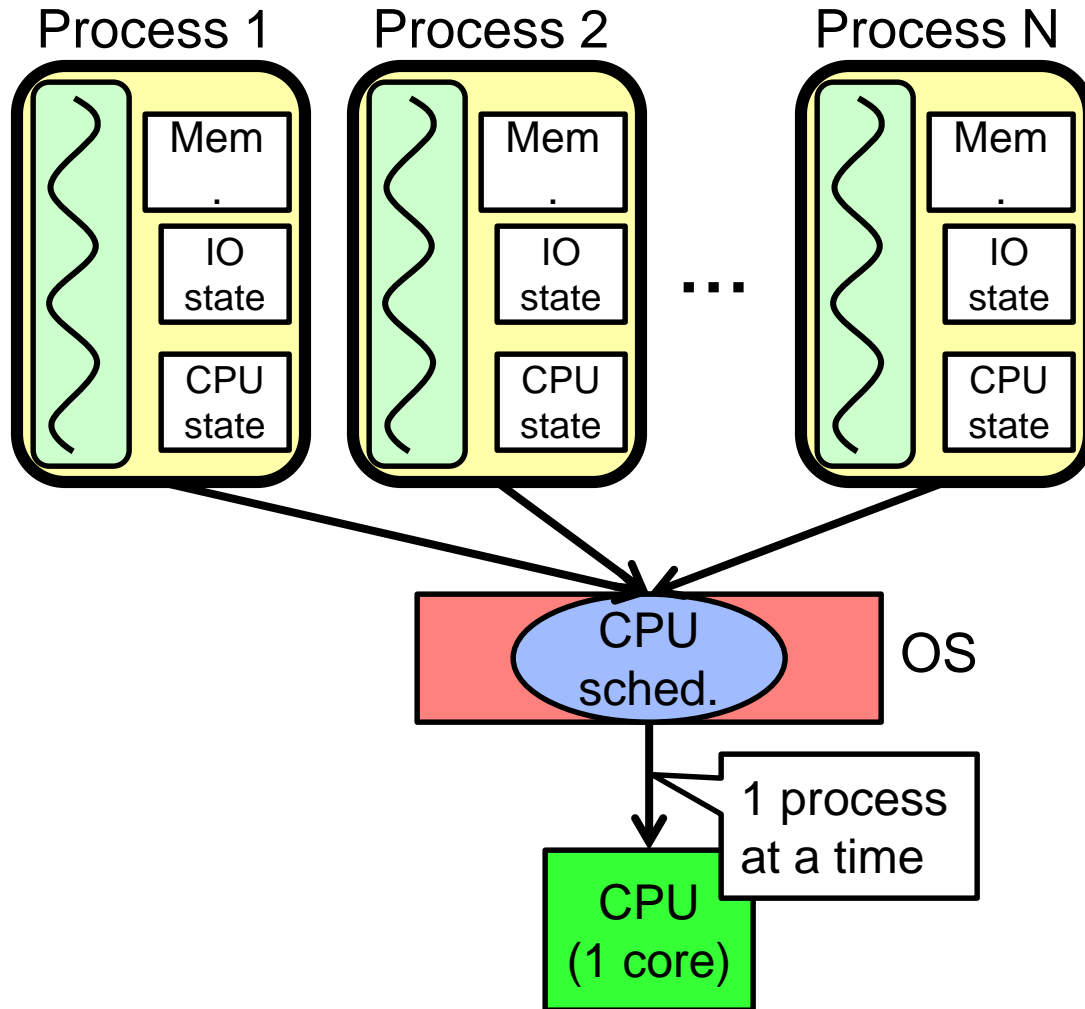
Concurrency and Thread Dispatching
(contd.)

Process (Review)

(Unix) Process

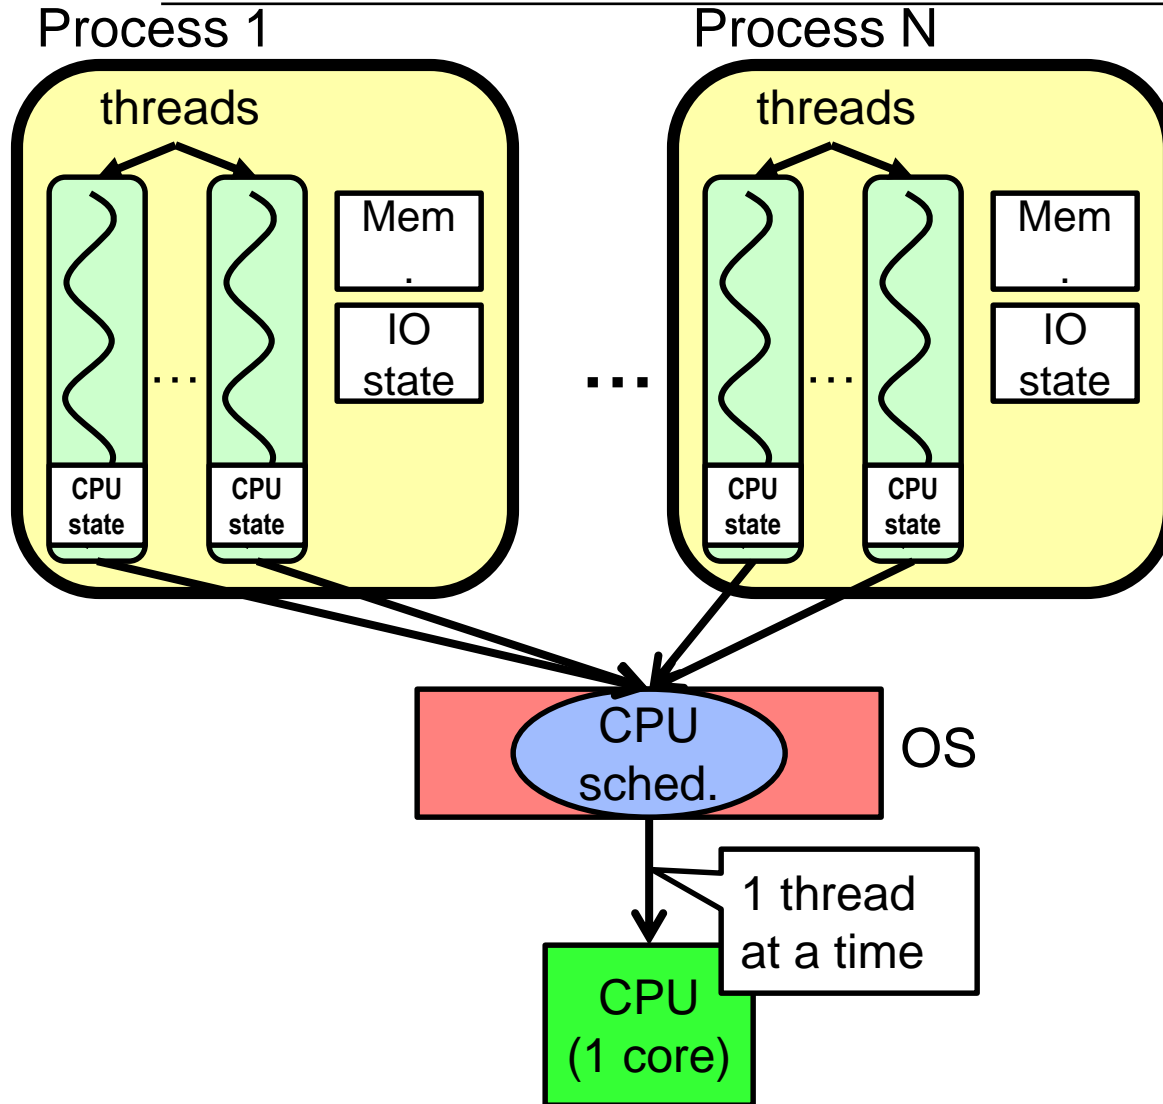


Processes (Review)



- Switch overhead: **high**
 - CPU state: **low**
 - Memory/IO state: **high**
- Process creation: **high**
- Protection
 - CPU: **yes**
 - Memory/IO: **yes**
- Sharing overhead: **high** (involves at least a context switch)

Multithreaded Processes (Review)



- Switch overhead: **low** (only CPU state)
- Thread creation: **low**
- Protection
 - CPU: **yes**
 - Memory/IO: **No**
- Sharing overhead: **low** (thread switch overhead low)

Why Processes & Threads? (Review)

Goals:

- **Multiprogramming:** Run multiple applications concurrently
- **Protection:** Don't want a bad application to crash system!

Solution:

Process: unit of execution and allocation

- Virtual Machine abstraction: give process illusion it owns machine (i.e., CPU, Memory, and IO device multiplexing)

Challenge:

- Need concurrency within same app (e.g., web server)
- Process creation, communication & switching expensive

Solution:

Thread: Decouple allocation and execution


- Run multiple threads within same process

Thread State (Review)

- State shared by all threads in process/address space
 - Content of memory (global variables, heap)
 - I/O state (file system, network connections, etc)
- State “private” to each thread
 - Kept in TCB \equiv Thread Control Block
 - CPU registers (including, program counter)
 - Execution stack – what is this?
- Execution Stack
 - Parameters, temporary variables
 - Return PCs are kept while called procedures are executing

Execution Stack Example (Review)

```
addrX:  A(int tmp) {  
        .  
        .   if (tmp<2)  
        .   B();  
addrY:  printf(tmp);  
        .  
        .   }  
        .  
        .   B() {  
        .   C();  
addrU:  }  
        .  
        .   C() {  
        .   A(2);  
addrV:  }  
        .  
        .   A(1);  
addrZ:  exit;
```



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Single-Threaded Example (Review)

- Imagine the following C program:

```
main() {  
    ComputePI("pi.txt");  
    PrintClassList("clist.text");  
}
```


- What is the behavior here?
 - Program would never print out class list
 - Why? ComputePI would never finish

Use of Threads (Review)

- Version of program with Threads:

```
main() {  
    CreateThread(ComputePI("pi.txt"));  
    CreateThread(PrintClassList("clist.text"));  
}
```

- What does `CreateThread` do?
 - Start independent thread running given procedure
- What is the behavior here?
 - Now, you would actually see the class list
 - This *should* behave as if there are two separate CPUs



Today


- Further Understanding Threads
 - Dispatching
 - Beginnings of Thread Scheduling

Why allow cooperating threads?

- People cooperate; computers help/enhance people's lives, so computers must cooperate
 - By analogy, the non-reproducibility/non-determinism of people is a notable problem for “carefully laid plans”
- Advantage 1: Share resources
 - One computer, many users
 - One bank balance, many ATMs
 - » What if ATMs were only updated at night?
 - Embedded systems (robot control: coordinate arm & hand)
- Advantage 2: Speedup
 - Overlap I/O and computation
 - Multiprocessors – chop up program into parallel pieces
- Advantage 3: Modularity
 - Chop large problem up into simpler pieces
 - » Makes system easier to extend

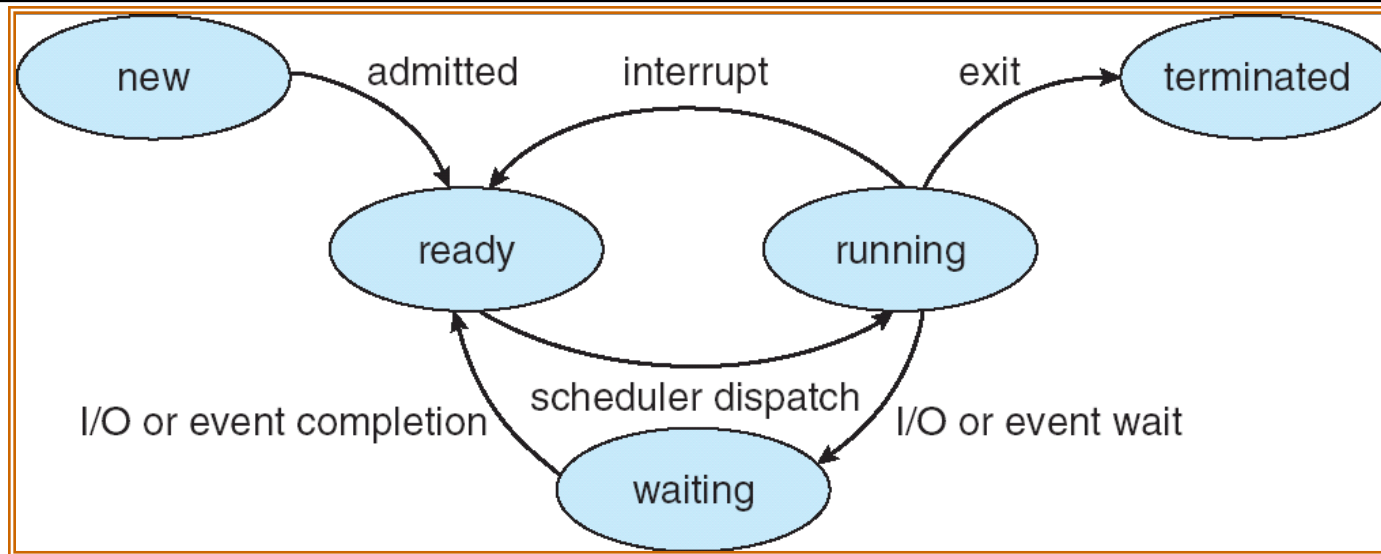
Multithreaded Processes

- PCB points to multiple TCBs:



- Switching threads within a block is a simple thread switch
- Switching threads across blocks requires changes to memory and I/O address tables.

Lifecycle of a Thread (or Process)



- As a thread executes, it changes state:
 - **new**: The thread is being created
 - **ready**: The thread is waiting to run
 - **running**: Instructions are being executed
 - **waiting**: Thread waiting for some event to occur
 - **terminated**: The thread has finished execution
- “Active” threads are represented by their TCBs
 - TCBs organized into queues based on their state

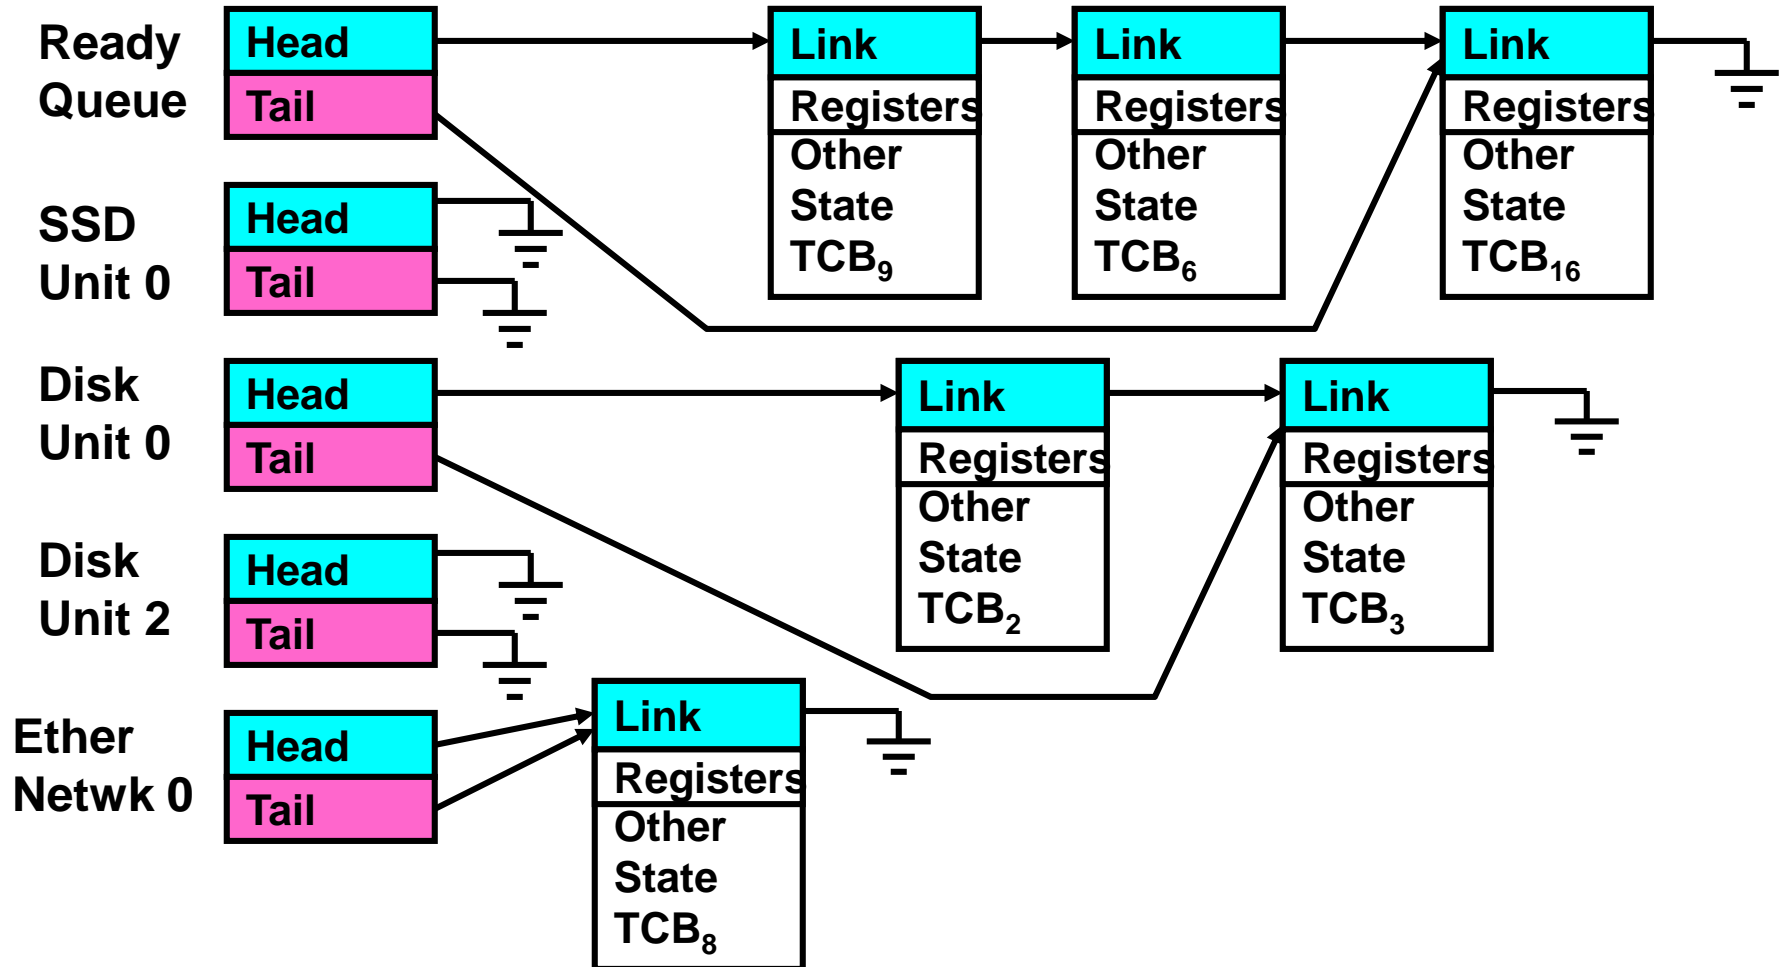
Ready Queues

- Note because of the actual number of live threads in a typical OS, and the (much smaller) number of running threads, most threads will be in a “ready” state.
- Thread not running \Rightarrow TCB is in some scheduler queue



Ready Queue And Various I/O Device Queues

- Thread not running \Rightarrow TCB is in some scheduler queue
 - Separate queue for each device/signal/condition
 - Each queue can have a different scheduler policy



Choosing a Thread to Run

- How does Dispatcher decide what to run?
 - Zero ready threads – dispatcher loops
 - » Alternative is to create an “idle thread”
 - » Can put machine into low-power mode
 - Exactly one ready thread – easy
 - More than one ready thread: use scheduling priorities
- Possible priorities:
 - LIFO (last in, first out):
 - » put ready threads on front of list, remove from front
 - Pick one at random
 - FIFO (first in, first out):
 - » Put ready threads on back of list, pull them from front
 - » This is fair and is what Nachos does
 - Priority queue:
 - » keep ready list sorted by TCB priority field

Per Thread State

- Each Thread has a *Thread Control Block* (TCB)
 - Execution State: CPU registers, program counter (PC), pointer to stack (SP)
 - Scheduling info: state (more later), priority, CPU time
 - Various Pointers (for implementing scheduling queues)
 - Pointer to enclosing process (PCB)
 - Etcetera (add stuff as you find a need)
- OS Keeps track of TCBs in protected memory
 - In Array, or Linked List, or ...

Dispatch Loop

- Conceptually, the dispatching loop of the operating system looks as follows:

```
Loop {  
    RunThread();  
    ChooseNextThread();  
    SaveStateOfCPU(curTCB);  
    LoadStateOfCPU(newTCB);  
}
```

- This is an *infinite* loop
 - One could argue that this is all that the OS does

Running a thread

Consider first portion: `RunThread()`

- How do I run a thread?
 - Load its state (registers, stack pointer) into CPU
 - Load environment (virtual memory space, etc)
 - Jump to the PC
- How does the dispatcher get control back?
 - Internal events: thread returns control voluntarily
 - External events: thread gets *preempted*


Yielding through Internal Events

- Blocking on I/O
 - The act of requesting I/O implicitly yields the CPU (e.g., printf)
- Waiting on a “signal” from other thread (e.g., join)
 - Thread asks to wait and thus yields the CPU
- Thread executes a `yield()`
 - Thread volunteers to give up CPU

```
computePI() {  
    while(TRUE) {  
        ComputeNextDigit();  
        yield();  
    }  
}
```

- Note that `yield()` must be called by programmer frequently enough!

Stack for Yielding a Thread

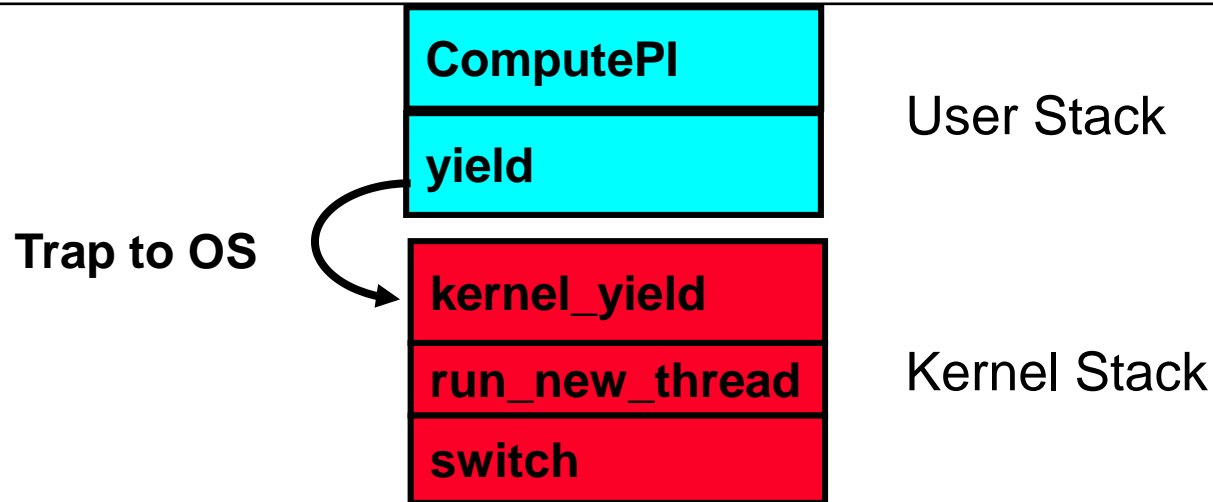


- How do we run a new thread?

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, newThread);  
    ThreadHouseKeeping(); /* deallocates finished  
                           threads */  
}
```

- Finished thread not killed right away.
 - Move them in “exit/terminated” state
 - ThreadHouseKeeping() deallocates finished threads

Stack for Yielding a Thread



- How do we run a new thread?

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, newThread);  
    ThreadHouseKeeping(); /* deallocates finished  
                           threads */  
}
```

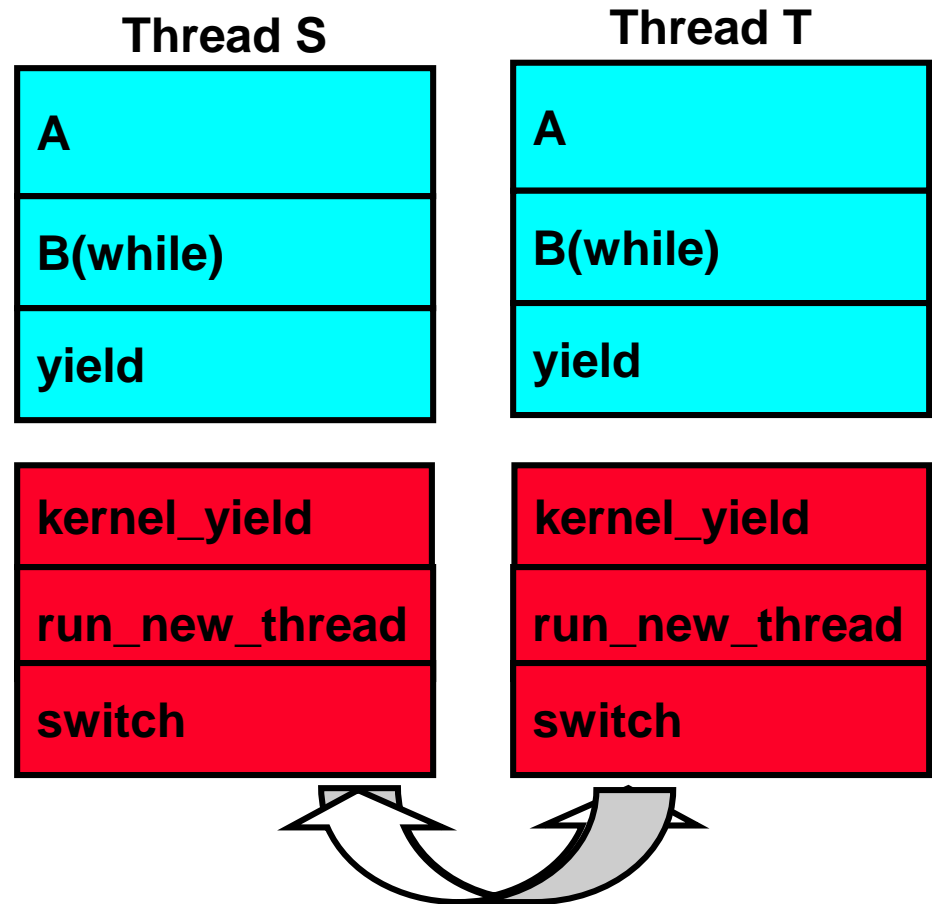
- How does dispatcher switch to a new thread?
 - Save anything next thread may trash: PC, regs, SP
 - Maintain isolation for each thread

Two Thread Yield Example

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```

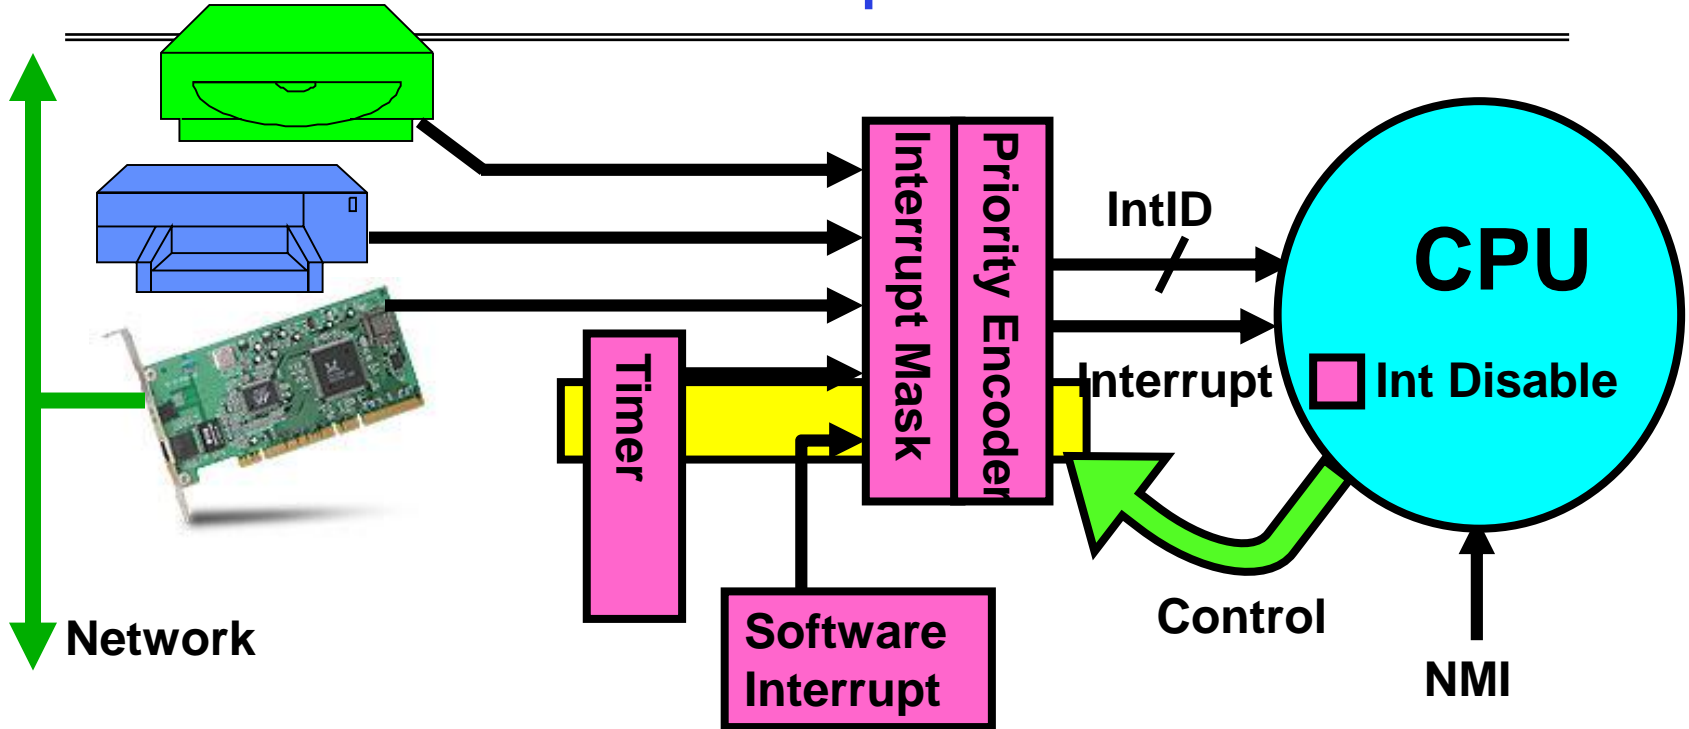
- Suppose we have two threads:
 - Threads S and T



Need for External Events

- What happens if thread never does any I/O, never waits, and never yields control?
 - Could the ComputePI program grab all resources and never release the processor?
 - » What if it didn't print to console?
 - Must find way that dispatcher can regain control!
- Answer: Utilize External Events
 - Interrupts: signals from hardware or software that stop the running code and jump to kernel
 - Timer: like an alarm clock that goes off every some many milliseconds
- If we make sure that external events occur frequently enough, can ensure dispatcher runs

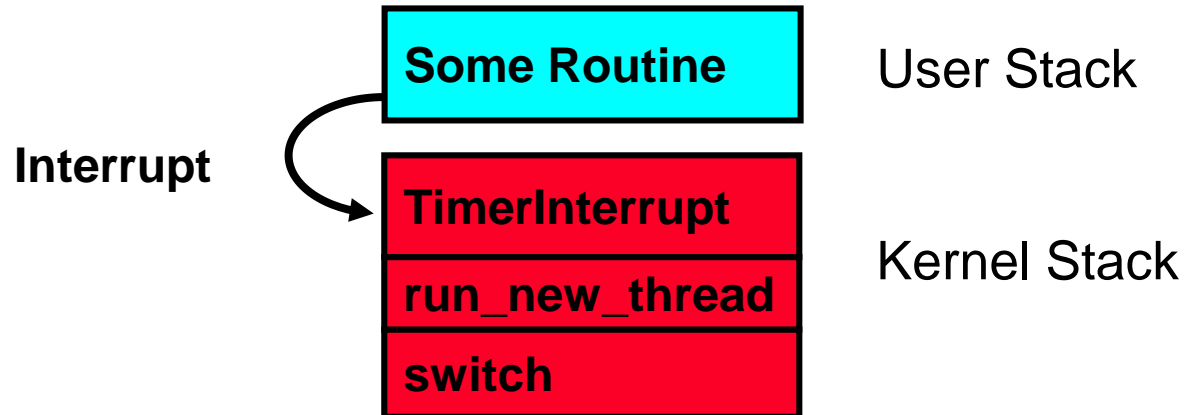
Detour: Interrupt Controller



- Interrupts invoked with interrupt lines from devices
- Interrupt controller chooses interrupt request to honor
 - Mask enables/disables interrupts
 - Priority encoder picks highest enabled interrupt
 - Software Interrupt Set/Cleared by Software
 - Interrupt identity specified with ID line
- CPU can disable all interrupts with internal flag
- Non-maskable interrupt line (NMI) can't be disabled

Preemptive Multithreading

- Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```
TimerInterrupt() {  
    DoPeriodicHouseKeeping();  
    run_new_thread();  
}
```

- This is often called **preemptive multithreading**, since threads are preempted for better scheduling
 - Solves problem of user who doesn't insert `yield()`;

Two important functions

ThreadFork (aNewThread, null) ;

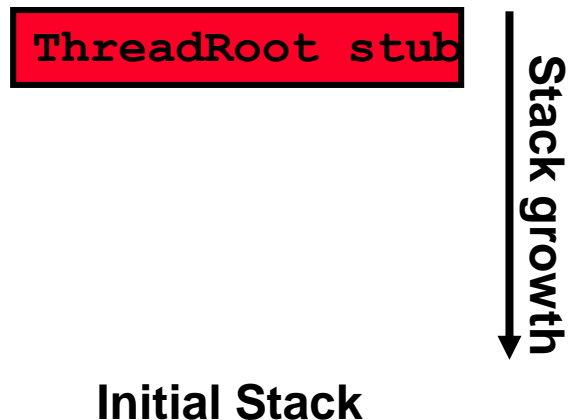
ThreadJoin (tid) ;

ThreadFork() : Create a New Thread

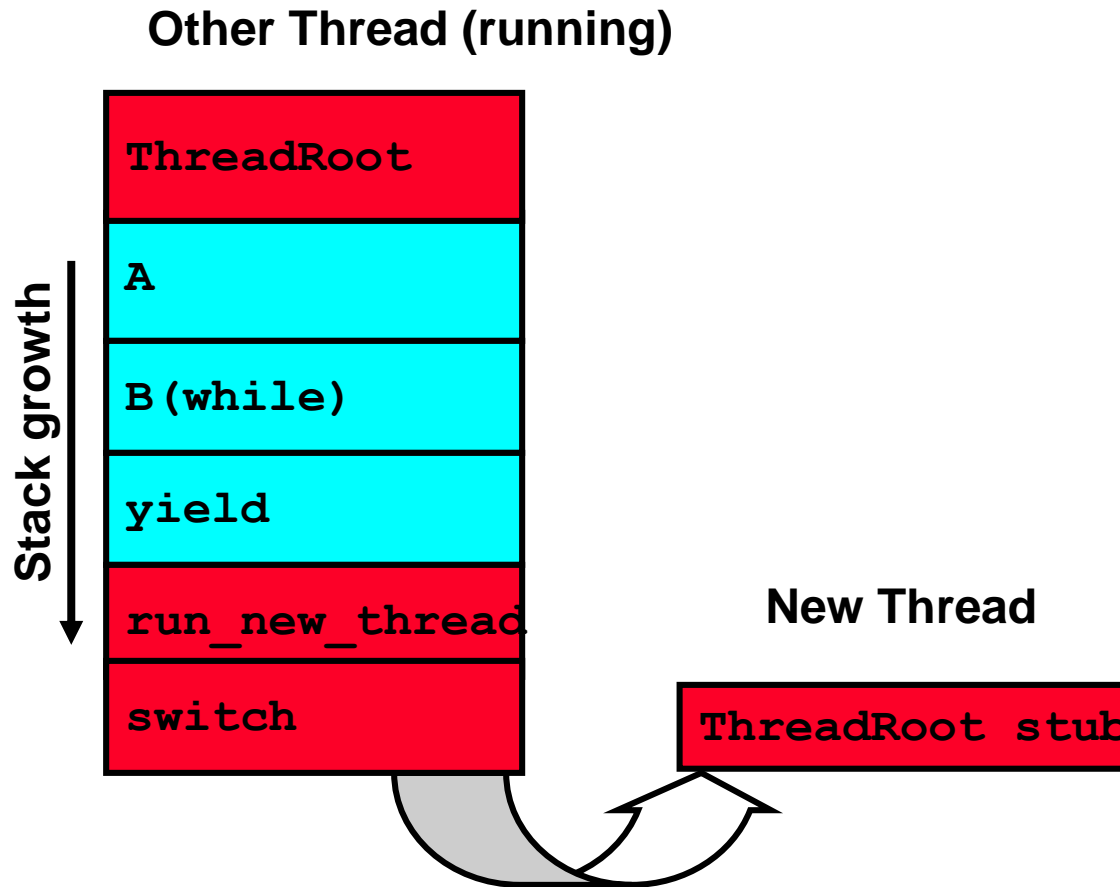
- ThreadFork() is a user-level procedure that creates a new thread and places it on ready queue
 - We called this CreateThread() earlier
- Arguments to ThreadFork()
 - Pointer to application routine (fcnPtr)
 - Pointer to array of arguments (fcnArgPtr)
 - Size of stack to allocate
- Implementation
 - Sanity Check arguments
 - Enter Kernel-mode and Sanity Check arguments again
 - Allocate new Stack and TCB
 - Initialize TCB and place on ready list (Runnable).

How do we initialize TCB and Stack?

- Initialize (4) Register fields of TCB (e.g., MIPS)
 - Stack pointer (sp) made to point at stack
 - PC return address (ra) set to OS (asm) routine `ThreadRoot()`
 - Two argument registers (a0 and a1) initialized to `fcnPtr` and `fcnArgPtr`, respectively
- Initialize stack data?
 - No. Important part of stack frame is in registers (ra)



How does Thread get started?




- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
 - This really starts the new thread

What does ThreadRoot () look like?

- ThreadRoot() is the root for the thread routine:

```
ThreadRoot() {  
    DoStartupHousekeeping();  
    UserModeSwitch(); /* enter user mode */  
    Call fcnPtr(fcnArgPtr);  
    ThreadFinish();  
}
```

- Startup Housekeeping
 - Includes things like recording start time of thread
 - Other Statistics
- Stack will grow and shrink with execution of thread
- Final return from thread returns into ThreadRoot () which calls ThreadFinish ()
 - ThreadFinish () will start at user-level



What does ThreadFinish() do?

- Needs to re-enter kernel mode (system call)
- “Wake up” (place on ready queue) threads waiting for this thread
 - Threads (like the parent) may be on a wait queue waiting for this thread to finish
- Can’t deallocate thread yet
 - We are still running on its stack! A thread cannot deallocate itself while running!
 - Instead, record thread as “waitingToBeDestroyed”
- Call `run_new_thread()` to run another thread:

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, newThread);  
    ThreadHouseKeeping();  
}
```

 - `ThreadHouseKeeping()` notices `waitingToBeDestroyed` and deallocates the finished thread’s TCB and stack

ThreadJoin() system call

- One thread can wait for another to finish with the `ThreadJoin(tid)` call
 - Calling thread will be taken off run queue and placed on waiting queue for thread `tid`

Use of Join for Traditional Procedure Call

- A traditional procedure call is logically equivalent to doing a `ThreadFork()` followed by `ThreadJoin()`
- Consider the following normal procedure call of `B()` by `A()`:

```
A() { B(); }
```

```
B() { Do interesting, complex stuff }
```

- The procedure `A()` is equivalent to `A'()`:

```
A'() {  
    tid = ThreadFork(B, null);  
    ThreadJoin(tid);  
}
```

- Why not do this for every procedure?
 - Context Switch Overhead
 - Memory Overhead for Stacks

Summary

- The state of a thread is contained in the TCB
 - Registers, PC, stack pointer
 - States: New, Ready, Running, Waiting, or Terminated
- Interrupts: hardware mechanism for returning control to operating system
 - Used for important/high-priority events
 - Can force dispatcher to schedule a different thread (preemptive multithreading)
- New Threads Created with `ThreadFork()`
 - Create initial TCB and stack to point at `ThreadRoot()`
 - `ThreadRoot()` calls thread code, then `ThreadFinish()`
 - `ThreadFinish()` wakes up waiting threads then prepares TCB/stack for destruction
- Threads can wait for other threads using `ThreadJoin()`
- Many scheduling options
 - Decision of which thread to run complex enough for complete lecture

Next

- Reading: Chapter 5.