

CSE150
Operating Systems
Lecture 6

Synchronization, Atomic operations, Locks

Why allow cooperating threads?

- People cooperate; computers help/enhance people's lives, so computers must cooperate
 - By analogy, the non-reproducibility/non-determinism of people is a notable problem for “carefully laid plans”
- Advantage 1: Share resources
 - One computer, many users
 - One bank balance, many ATMs
 - » What if ATMs were only updated at night?
 - Embedded systems (robot control: coordinate arm & hand together)
- Advantage 2: Speedup
 - Overlap I/O and computation
 - Multiprocessors – chop up program into parallel pieces
- Advantage 3: Modularity
 - Chop large problem up into simpler pieces
 - » To compile, for instance, gcc calls cpp | cc1 | cc2 | as | ld
 - » Makes system easier to extend

Interrupts and Two Functions

- The state of a thread is contained in the TCB
 - Registers, PC, stack pointer
 - States: New, Ready, Running, Waiting, or Terminated
- Interrupts: hardware mechanism for returning control to operating system
 - Used for important/high-priority events
 - Can force dispatcher to schedule a different thread (preemptive multithreading)
- New Threads Created with `ThreadFork()`
 - Create initial TCB and stack to point at `ThreadRoot()`
 - `ThreadRoot()` calls thread code, then `ThreadFinish()`
 - `ThreadFinish()` wakes up waiting threads then prepares TCB/stack for destruction
- Threads can wait for other threads using `ThreadJoin()`
- Many scheduling options
 - Decision of which thread to run complex enough for complete lecture

Today

- Concurrency examples and sharing
- Synchronization


Announcements - Important time points

- Lab sessions start in **the week 02/10.**
- Project 1: 4 weeks
- Project 2: 4 weeks
- Three or Four homework (Feb. 13th; Feb 27th, March 10th, April 9th)
 - One week to finish each.
- In-class exam (March. 19th)
- Final exam (11:30 AM - 2:30 PM May 13th)

Correctness for systems with concurrent threads

- **Independent Threads (ideal):**
 - No state shared with other threads
 - Deterministic \Rightarrow Input state determines results
 - Reproducible \Rightarrow Can recreate Starting Conditions, I/O
 - Scheduling order doesn't matter (if `switch()` works!!!)
- **Cooperating Threads:**
 - Shared State between multiple threads
 - Non-deterministic
 - Non-reproducible (performance??)
- Non-deterministic and Non-reproducible means that bugs can be intermittent

ATM Bank Server



- ATM server problem:
 - Service a set of requests
 - Do so without corrupting database
 - Don't hand out too much money

ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```
BankServer() {  
    while (TRUE) {  
        ReceiveRequest(&op, &acctId, &amount);  
        ProcessRequest(op, acctId, amount);  
    }  
}  
  
ProcessRequest(op, acctId, amount) {  
    if (op == deposit) Deposit(acctId, amount);  
    else if ...  
}  
  
Deposit(acctId, amount) {  
    acct = GetAccount(acctId); /* may use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct); /* Involves disk I/O */  
}
```

- How could we speed this up?
 - More than one request being processed at once
 - Multiple threads (multi-proc, or overlap comp and I/O)

Can Threads Help?

- One thread per request!
- Requests proceed to completion, blocking as required:

```
Deposit(acctId, amount) {  
    acct = GetAccount(actId); /* May use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct);      /* Involves disk I/O */  
}
```

- Unfortunately, shared state can get corrupted:

Thread 1

```
load r1, acct->balance  
  
add r1, amount1  
store r1, acct->balance
```

Thread 2

```
load r1, acct->balance  
add r1, amount2  
store r1, acct->balance
```

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, What about (Initially, $y = 12$):

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of x ?

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2$

$x = 13$

Preemption can occur
at any time!

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, What about (Initially, $y = 12$):

Thread A

$x = 1;$

$x = y + 1;$

Thread B


$y = 2;$

$y = y * 2;$

- What are the possible values of x ?

Thread A

Thread B



$x=5$

Preemption can occur
at any time!

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, What about (Initially, $y = 12$):

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of x ?

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

$x = 3$

Preemption can occur
at any time!

Atomic Operations

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- **Atomic Operation**: an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
- Many instructions are not atomic
 - Double-precision floating point store often not atomic
 - VAX and IBM 360 had an instruction to copy a whole array

Concurrency Challenges

- Multiple computations (threads) executing in parallel to
 - share resources, and/or
 - share data
- Examples:
 - Sharing CPU for 10ms vs. 1min
 - Sharing a database at the row vs. table granularity
- Fine grain sharing:
 - ↑ increase concurrency → better performance
 - ↓ more complex
- Coarse grain sharing:
 - ↑ Simpler to implement
 - ↓ Lower performance

Definitions

- **Synchronization**: using atomic operations to ensure cooperation between threads
 - For now, only loads and stores are atomic
 - We will show that it is hard to build anything useful with only atomic reads and writes
- **Critical Section**: piece of code that only one thread can execute at once.
- **Mutual Exclusion**: ensuring that only one thread executes a critical section
 - One thread *excludes* the other while doing its task
 - Critical section and mutual exclusion are two ways of describing the same thing.

Motivation: “Too much milk”



- Great thing about OS's – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

Lock

- **Lock**: prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
- » Important idea: all synchronization involves waiting
- For example: fix the milk problem by putting a key on the refrigerator
 - Lock it and take key if you are going to go buy milk
 - Fixes too much (coarse granularity): roommate angry if only wants OJ



- Of Course – We don't know how to make a lock yet


Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
 - Always write down ***desired*** behavior first
 - Impulse is to start coding first, then when it doesn't work, pull hair out
 - Instead, think first, then code
- What are the correctness properties for the “Too much milk” problem?
 - Never more than one person buys
 - Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



- Result?

Too Much Milk: Solution #1

- Still too much milk **but only occasionally!**

<u>Thread A</u>	<u>Thread B</u>
<pre>if (noMilk) if (noNote) {</pre>	
	<pre>if (noMilk) if (noNote) {</pre>
<pre> leave Note; buy milk; remove note; }</pre>	
<pre>}</pre>	<pre> leave Note; buy milk; ...</pre>

- Thread can get context switched after checking milk and note but before leaving note!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...
 - Must work despite what the thread dispatcher does!

Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
    }  
}  
remove note;
```

- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys milk



Too Much Milk Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note B;
```

- Does this work?

Too Much Milk Solution #2

- Possible for neither thread to buy milk!

Thread A

```
leave note A;
```

Thread B


```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy Milk;  
    }  
}
```

```
if (noNote B) {  
    if (noMilk) {  
        buy Milk;  
    }  
    ...  
}
```

```
remove note B;
```

- Really insidious:
 - **Unlikely** that this would happen, but will at worse possible time

Too Much Milk Solution #2: problem!



- *I'm not getting milk, You're getting milk*
- This kind of lockup is called "starvation!"

Too Much Milk Solution #3

- Here is a possible two-note solution:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) { //X	if (noNote A) { //Y
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	}
}	remove note B;
remove note A;	

- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - if no note B, safe for A to buy,
 - otherwise wait to find out what will happen
- At Y:
 - if no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

Solution #3 discussion

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - » Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - » Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - » This is called “busy-waiting”
- There's a better way

Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock (more in a moment).
 - `Lock.Acquire()` – wait until lock is free, then grab
 - `Lock.Release()` – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
millock.Acquire();  
if (noMilk)  
    buy milk;  
millock.Release();
```
- Once again, section of code between `Acquire()` and `Release()` called a “**Critical Section**”

Summary

- Concurrent threads are a very useful abstraction
 - Allow transparent overlapping of computation and I/O
 - Allow use of parallel processing when available
- Concurrent threads introduce problems when accessing shared data
 - Programs must be insensitive to arbitrary interleaving
 - Without careful design, shared variables can become completely inconsistent
- Important concept: Atomic Operations
 - An operation that runs to completion or not at all
 - These are the primitives on which to construct various synchronization primitives
- How to protect a critical section with only atomic load and store
⇒ pretty complex!

Next

- Reading: Chapter 6.