

CSE150
Operating Systems
Lecture 7

Mutual Exclusion

Synchronization problem with Threads (Review)

- One thread per transaction, each running:

```
Deposit(acctId, amount) {  
    acct = GetAccount(actId); /* May use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct);      /* Involves disk I/O */  
}
```

- Unfortunately, shared state can get corrupted:

Thread 1

```
load r1, acct->balance
```

```
add r1, amount1  
store r1, acct->balance
```

Thread 2

```
load r1, acct->balance  
add r1, amount2  
store r1, acct->balance
```

- **Atomic Operation**: an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle

Too Much Milk Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note B;
```

- Does this work?

Too Much Milk Solution #2.5

- Algorithm looks like this:

Thread A

```
if (noNote B) {  
  
    leave note A;  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
  
        buy Milk;  
    }  
}  
remover note B;
```

- Does this work?

Too Much Milk Solution #3 (Review)

- Here is a possible two-note solution:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) {\\X	if (noNote A) {\\Y
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	}
}	remove note B;
remove note A;	

- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - if no note B, safe for A to buy,
 - otherwise wait to find out what will happen
- At Y:
 - if no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

Solution #3 discussion (Review)

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - » Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - » Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - » This is called “busy-waiting”
- There's a better way

Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock (more in a moment).
 - `Lock.Acquire()` – wait until lock is free, then grab
 - `Lock.Release()` – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
millock.Acquire();  
if (noMilk)  
    buy milk;  
millock.Release();
```
- Once again, section of code between `Acquire()` and `Release()` called a “**Critical Section**”

Why a thread joins another thread, but not does the job by its own?

- Multithreading
 - Share resources among multiple tasks
 - Modularity
 - Parallel on multiple CPU cores
- One thread may wait for the result of another thread
- One thread may wait for the result of two other threads that run in parallel.

Today

- Locks
- Higher-level Synchronization Abstractions
 - Semaphores
- Programming paradigms for concurrent programs

How to implement Locks?

- **Lock**: prevents someone from accessing/doing something
 - Lock before entering critical section (e.g., before accessing shared data)
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - » Important idea: all synchronization involves waiting
 - » Should *sleep* if waiting for a long time
- Hardware Lock instruction
 - Is this a good idea?
 - What about putting a task to sleep?
 - » How do you handle the interface between the hardware and scheduler?
 - Complexity?
 - » Each feature makes hardware more complex and slow



Naïve use of Interrupt Enable/Disable

- How can we build multi-instruction atomic operations?
 - Recall: dispatcher gets control in two ways.
 - » Internal: Thread does something to relinquish the CPU
 - » External: Interrupts cause dispatcher to take CPU
 - On a uniprocessor, can avoid context-switching by:
 - » Avoiding internal events
 - » Preventing external events by disabling interrupts
- Consequently, naïve implementation of locks:

```
LockAcquire { disable ints; }  
LockRelease { enable ints; }
```

Naïve use of Interrupt Enable/Disable

- **Can't let user do this!** Consider following:

```
LockAcquire();  
While(TRUE) {;}
```

- Real-Time system—no guarantees on timing!
 - Critical Sections might be arbitrarily long
- What happens with I/O or other important events?
 - “Reactor about to meltdown. Help?”



Better Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable (in memory) and impose mutual exclusion only during operations (access/update) on that variable

```
int value = FREE;
```



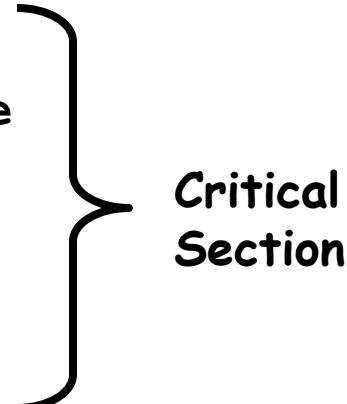
```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue  
        and Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        Take thread off wait queue  
        Put at front of ready queue  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

New Lock Implementation: Discussion

- Disable interrupts: avoid interrupting between checking and setting lock value
 - Otherwise two threads could think that they both have lock

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue  
        and Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```





Critical
Section

- Note: unlike previous solution, critical section very short
 - User of lock can take as long as they like in their own critical section
 - Critical interrupts taken in time

Interrupt re-enable in going to sleep

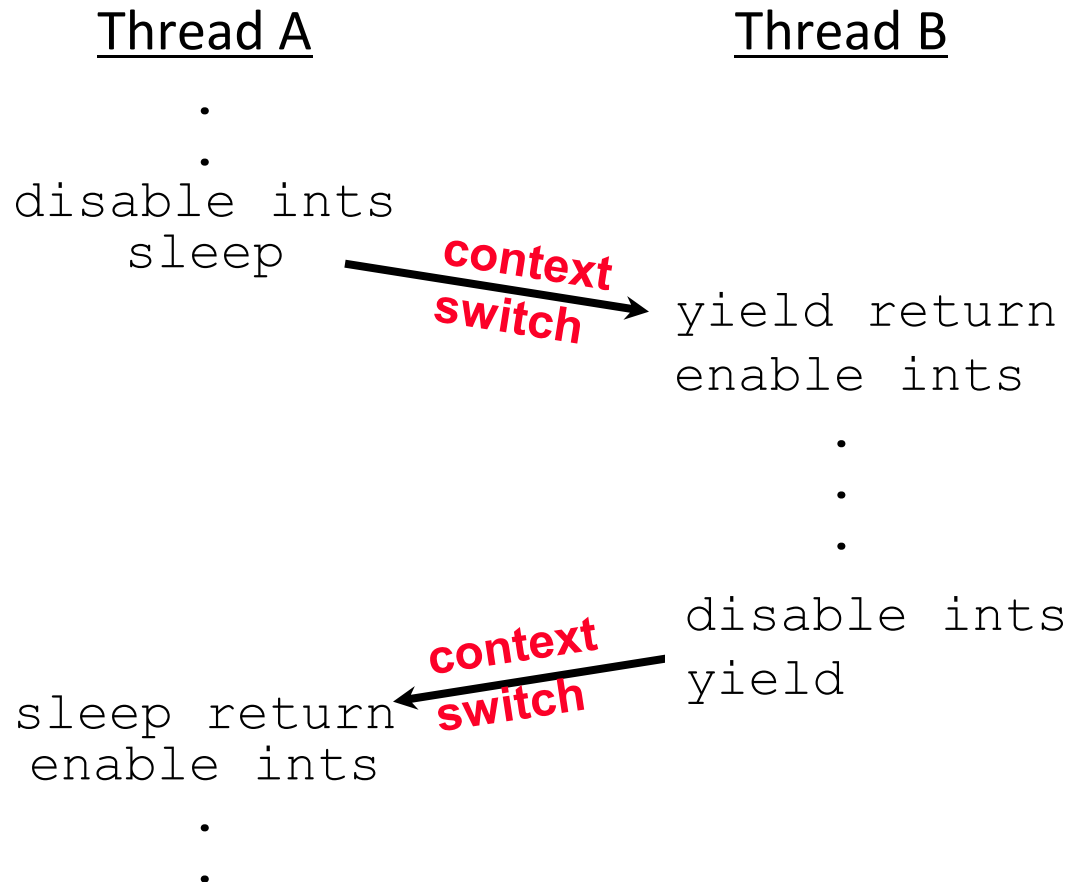
- What about re-enabling `ints` when going to sleep?

	<pre>Acquire() { disable interrupts; if (value == BUSY) { put thread on wait queue and Go to sleep(); } else { value = BUSY; } enable interrupts; }</pre>
Enable Position	
Enable Position	

- Before putting thread on the wait queue?
 - Release puts the thread on the ready queue, but the thread still thinks it needs to go to sleep
 - Misses wakeup
- Want to put it after `sleep()`. But, how?

How to Re-enable After Sleep () ?

- Since `ints` are disabled when you call `sleep`:
 - Responsibility of the next thread to re-enable `ints`
 - When the sleeping thread wakes up, returns to acquire and re-enables interrupts



Interrupt disable and enable across context switches

- An important point about structuring code:
 - In Nachos code you will see lots of comments about assumptions made concerning when interrupts disabled
 - This is an example of where modifications to and assumptions about program state can't be localized within a small body of code
 - In these cases it is possible for your program to eventually “acquire” bugs as people modify code
- Other cases where this will be a concern?
 - What about exceptions that occur after lock is acquired? Who releases the lock?

```
mylock.acquire();  
a = b / 0;  
mylock.release();
```

Atomic Read-Modify-Write instructions

- Problems with previous solution:
 - Can't give lock implementation to users
 - Doesn't work well on multiprocessor
 - » Disabling interrupts on all processors requires messages and would be very time consuming
- Alternative: “richer” atomic instruction **sequences**
 - These instructions read a value from memory and write a new value atomically
 - Hardware is responsible for implementing this correctly
 - » on both uniprocessors (moderately hard, x86: yes, MIPS: partially)
 - » and multiprocessors (requires help from cache coherence protocol)
 - Unlike disabling interrupts, can be used on both uniprocessors and multiprocessors

Examples of Read-Modify-Write

- `test&set (&address) { /* most architectures */
 result = M[address];
 M[address] = 1;
 return result;
}`
- `swap (&address, register) { /* x86 */
 temp = M[address];
 M[address] = register;
 register = temp;
}`
- `compare&swap (&address, reg1, reg2) { /* 68000 */
 if (reg1 == M[address]) {
 M[address] = reg2;
 return success;
 } else {
 return failure;
 }
}`

Implementing Locks with test&set

- Simple solution:

```
int value = 0; // Free
Acquire() {
    while (test&set(value));
}
Release() {
    value = 0;
}
```

```
test&set (&address) {
    result = M[address];
    M[address] = 1;
    return result;
}
```

- Simple explanation:

- If lock is free, test&set reads 0 and sets value=1, so lock is now busy. It returns 0 so while exits
- If lock is busy, test&set reads 1 and sets value=1 (no change). It returns 1, so while loop continues
- When we set value = 0, someone else can get lock

Problem: Busy-Waiting for Lock

- Positives for this solution
 - Machine can receive interrupts
 - User code can use this lock
 - Works on a multiprocessor
- Negatives
 - Inefficient: busy-waiting thread will consume cycles waiting
 - Waiting thread may take cycles away from thread holding lock (no one wins!)
 - **Priority Inversion**: If busy-waiting thread has higher priority than thread holding lock \Rightarrow no progress!



Better Locks using test&set

- Can we build test&set locks without busy-waiting?
 - Can't entirely, but can minimize!
 - Idea: only busy-wait to atomically check lock value

```
int guard = 0;  
int value = FREE;
```



```
Acquire() {  
    // Short busy-wait time  
    while (test&set(guard));  
    if (value == BUSY) {  
        put thread on wait queue;  
        go to sleep() & guard = 0;  
    } else {  
        value = BUSY;  
        guard = 0;  
    }  
}
```

```
Release() {  
    // Short busy-wait time  
    while (test&set(guard));  
    if anyone on wait queue {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    guard = 0;  
}
```

- Note: sleep has to be sure to reset the guard variable
 - Why can't we do it just before or just after the sleep?

Locks using test&set vs. Interrupts

- Compare to “disable interrupt” solution

```
int value = FREE;
```

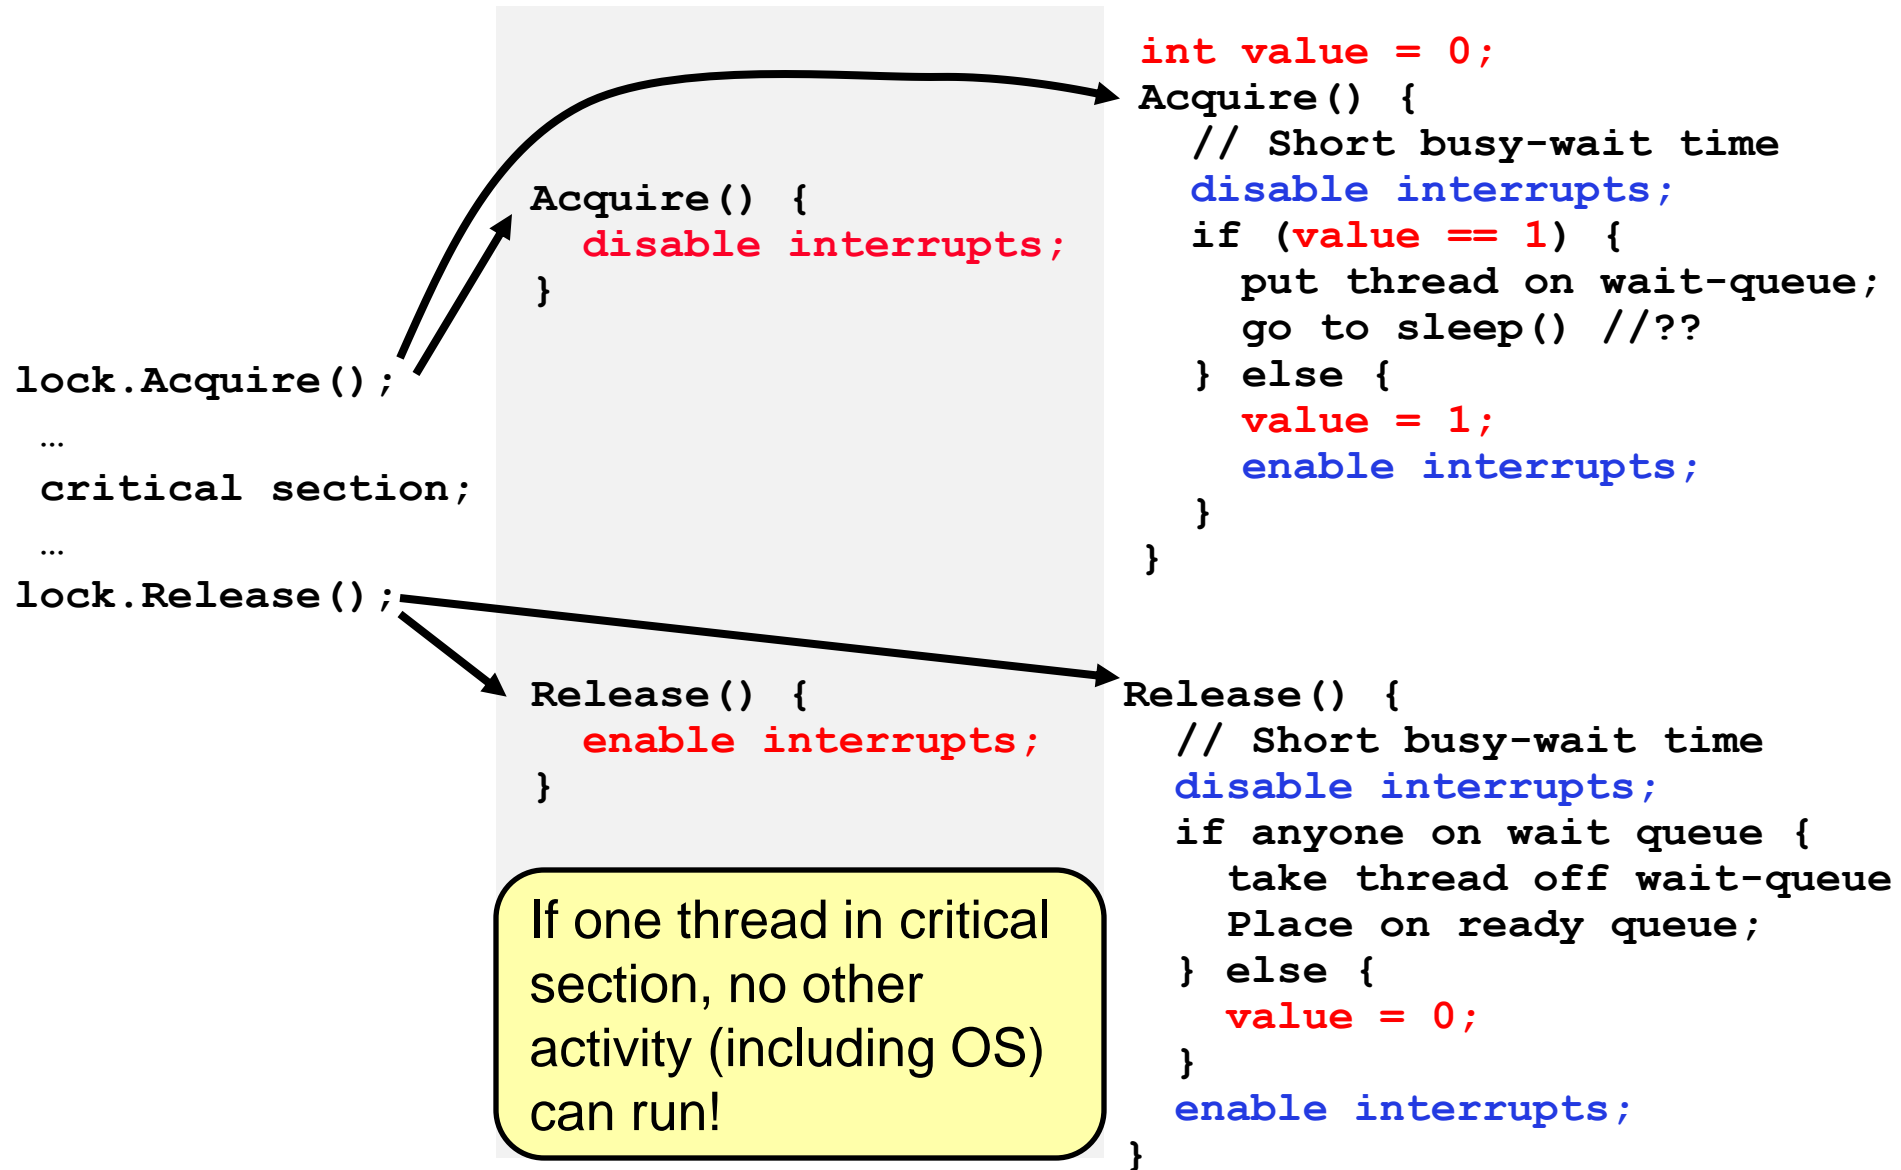


```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

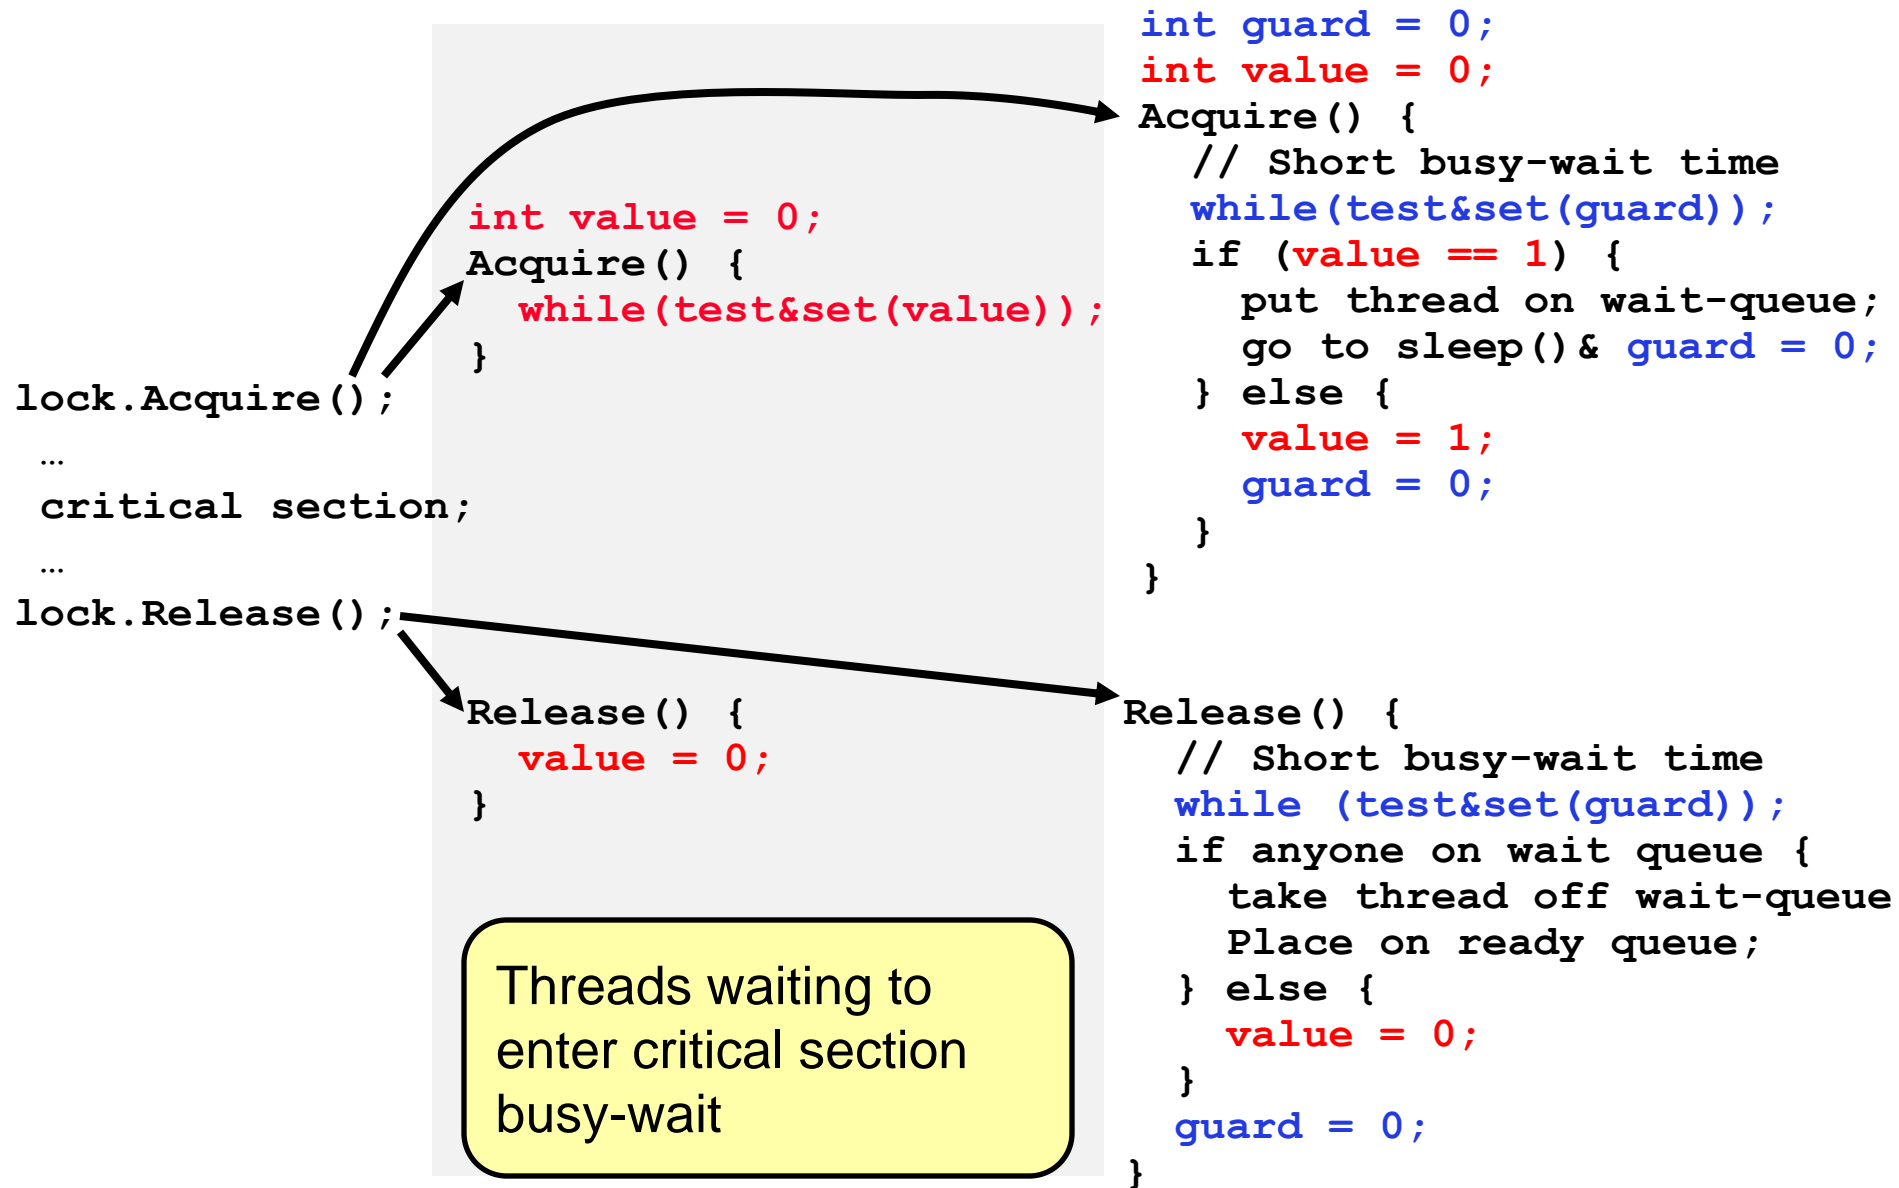
```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

- Basically replace
 - `disable interrupts` → `while (test&set(guard));`
 - `enable interrupts` → `guard = 0;`

Recap: Locks



Recap: Locks



Summary

- Talked about hardware atomicity primitives:
 - Disabling of Interrupts, test&set, swap, comp&swap, load-linked/store conditional
- Showed several constructions of Locks
 - Must be very careful not to waste/tie up machine resources
 - » Shouldn't disable interrupts for long
 - » Shouldn't spin wait for long
 - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable