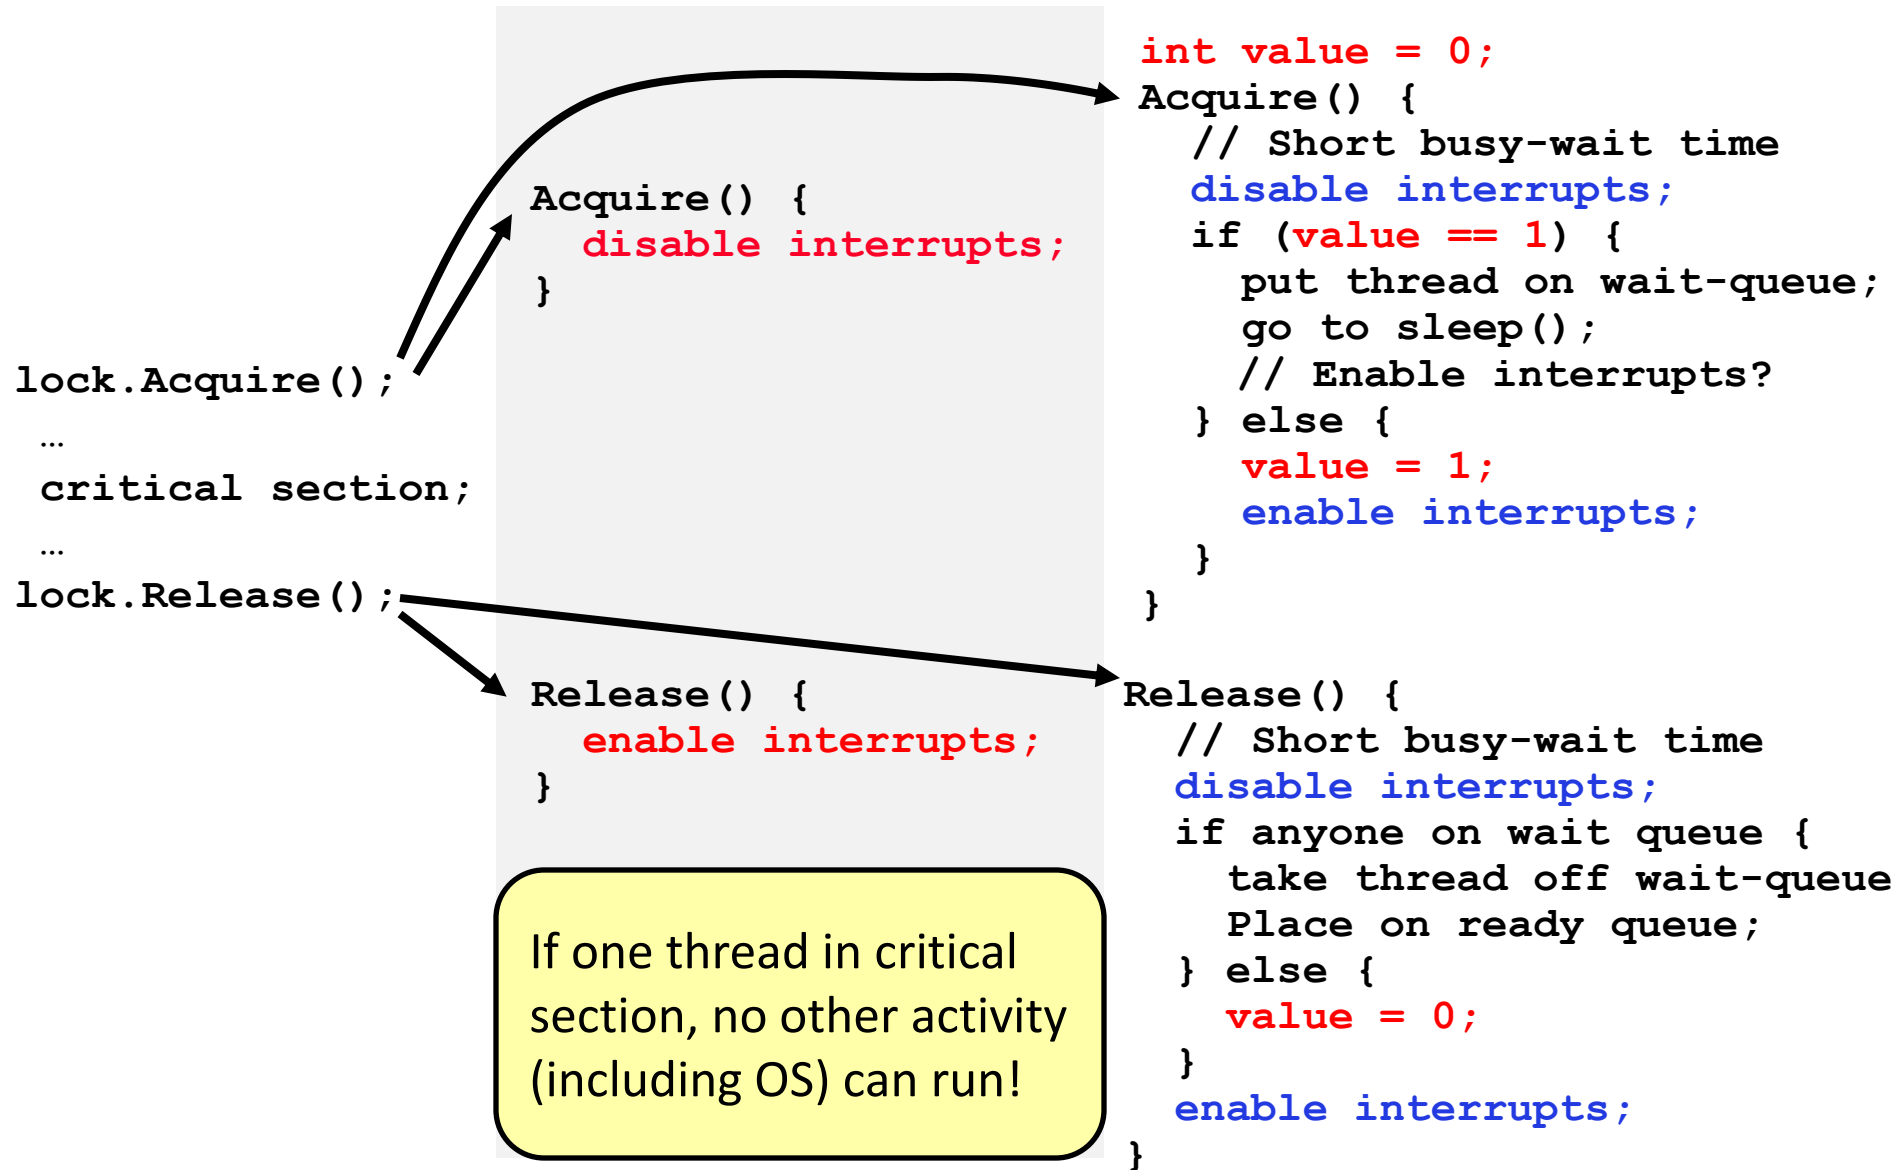


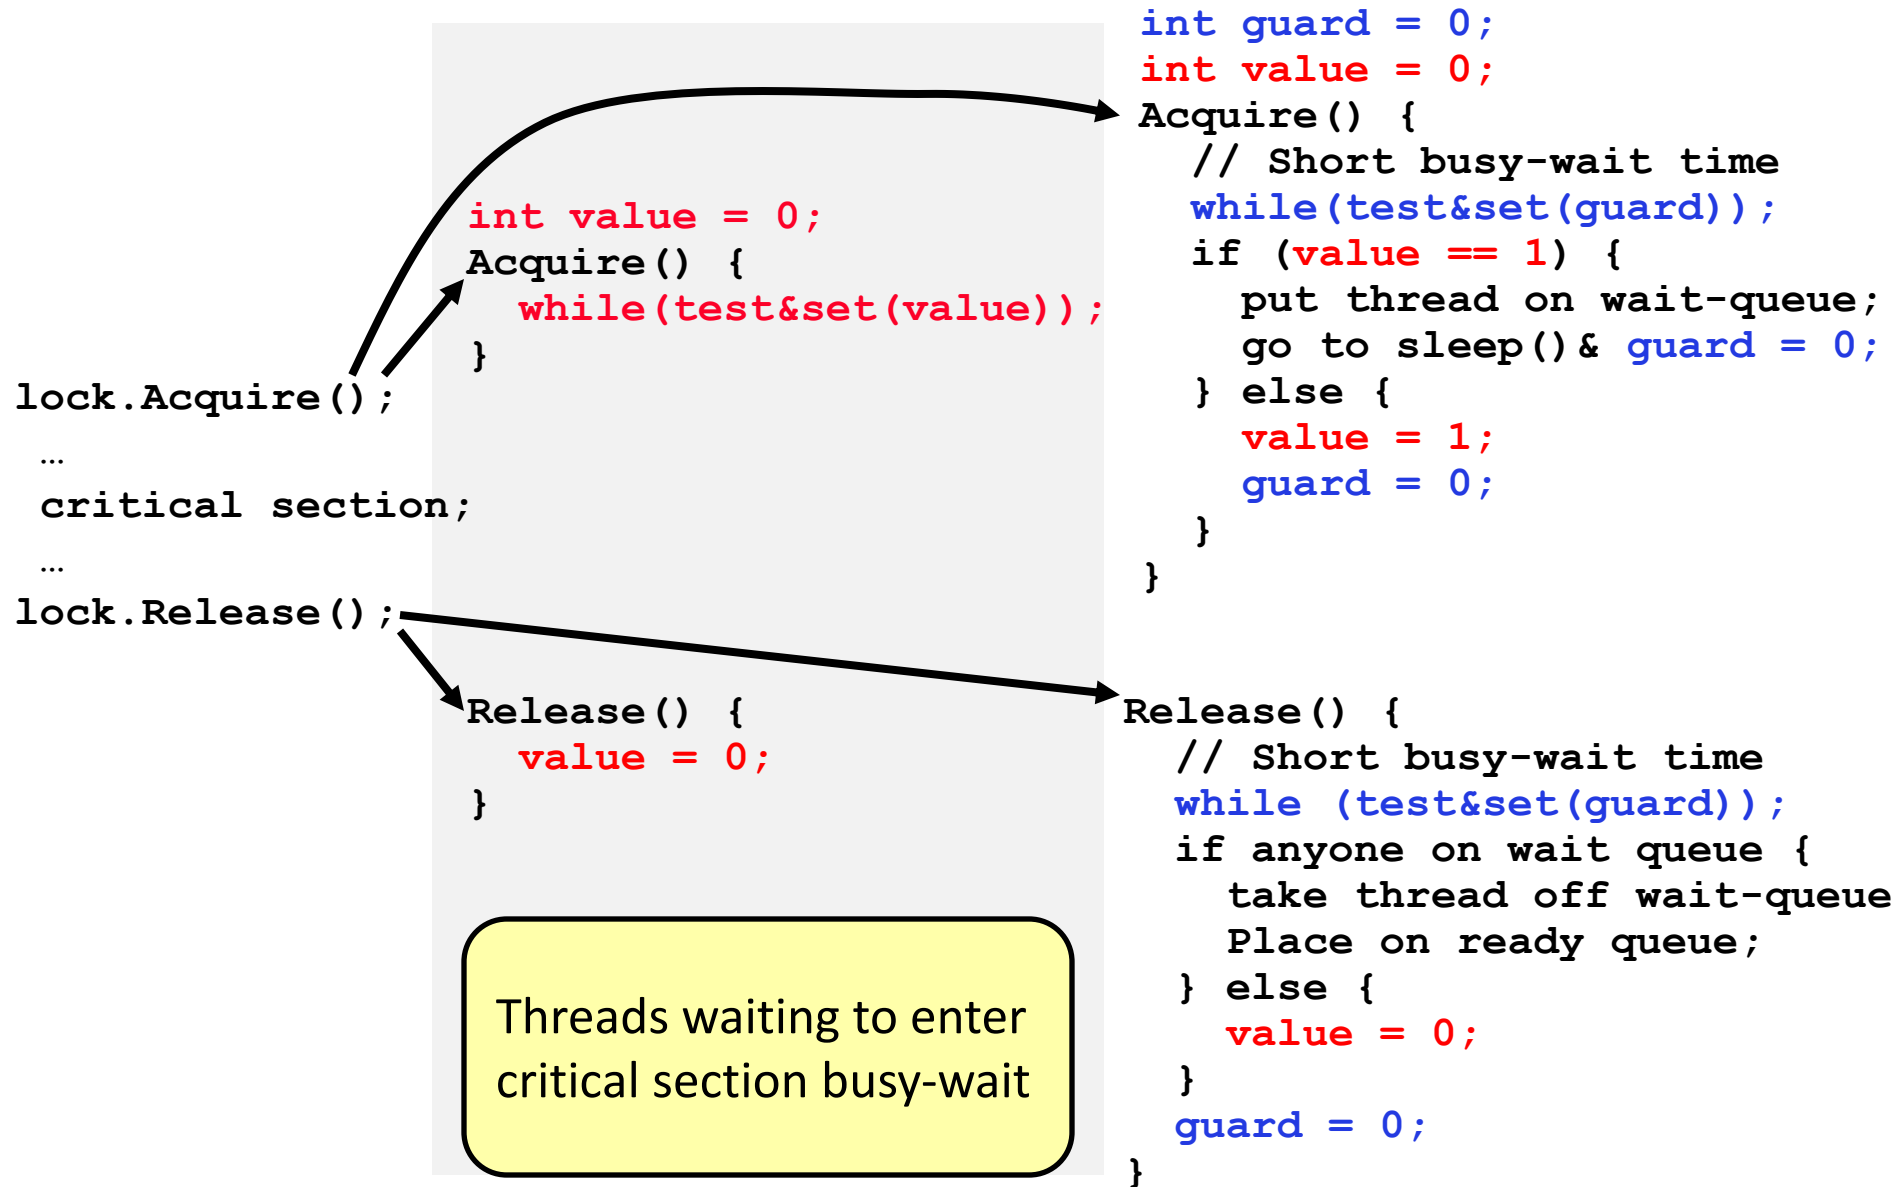
CSE150
Operating Systems
Lecture 8

Semaphores and Monitors

Review: Locks by Disabling Interrupts




Review: Locks using test&set



Communication & Coordination

- Group leader
- Lab sessions
- Regular meeting



Projects: Techniques for Partitioning Tasks

- Functional
 - Person A implements threads, Person B implements semaphores, Person C implements locks...
- Task
 - Person A designs, Person B writes code, Person C tests
 - May be difficult to find right balance, but can focus on each person's strengths (Theory vs systems hacker)
 - Since Debugging is hard, Microsoft has *two* testers for *each* programmer
- Most CSE-150 project teams are functional, but people have had success with task-based divisions

Suggested Documents for You to Maintain

- Organizational Chart
 - Who is responsible for what task?
- Schedule: What is your anticipated timing?
 - This document is critical!
- Meeting notes
 - Document all decisions
 - You can often cut & paste for the design documents
- Design document: the manual plus performance specs
 - This should be the first document generated and the last one finished.
 - Transform your design document to the final project report.



Test Continuously

- Integration tests all the time, not at 11pm on due date!
 - Schedule periodic integration tests
 - » Get everyone in the same room, check out code, build, and test.
 - » Don't wait until it is too late!
- Test early, test later, test again
 - Tendency is to test once and forget; what if something changes in some other part of the code?



Today


- Continue with Synchronization Abstractions
 - Semaphores, monitors, and condition variables

Where are we going with synchronization?

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors Send/Receive
Hardware	Load/Store Disable Ints Test&Set Comp&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
 - Everything is pretty painful if only atomic primitives are load and store
 - Need to provide primitives useful at user-level

Semaphores



- Semaphores are a kind of generalized lock
 - First defined by Dijkstra in late 60s
 - Main synchronization primitive used in original UNIX
- Definition: a Semaphore has a non-negative integer value and supports the following two operations:
 - **P()**: an atomic operation that waits for semaphore to become positive, then decrements it by 1
 - **V()**: an atomic operation that increments the semaphore by 1, waking up a waiting P, if any
 - Note that **P()** stands for “*proberen*” (to test) and **V()** stands for “*verhogen*” (to increment) in Dutch

Semaphores Like Integers Except

- Semaphores are like integers, except
 - No negative values
 - Only operations allowed are P and V – can't read or write value, except to set it initially
 - Operations must be atomic
 - » Two P's together can't decrement value below zero
 - » Thread going to sleep in P won't miss wakeup from V
- Semaphore from railway analogy
 - Here is a semaphore initialized to 2 for resource control:



Two Uses of Semaphores

- Mutual Exclusion (initial value = 1)

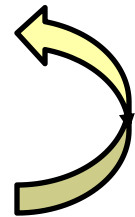
- Also called “Binary Semaphore”.
- Can be used for mutual exclusion:

```
semaphore.P();  
// Critical section goes here  
semaphore.V();
```


- Scheduling Constraints (initial value = 0)

- Allow thread 1 to wait for a signal from thread 2, i.e., thread 2 **schedules** thread 1 when a given **constraint** is satisfied
- Example: suppose you had to implement ThreadJoin which must wait for thread to terminate:

```
Initial value of semaphore = 0  
ThreadJoin {                               // Main program  
    semaphore.P();  
}  
ThreadFinish {                             // Child Thread  
    semaphore.V();  
}
```



Producer-consumer with a bounded buffer



- Problem Definition
 - Producer puts things into a shared buffer
 - Consumer takes them out
 - Need synchronization to coordinate producer/consumer
- Put a fixed-size buffer between producer and consumer
 - Need to synchronize access to this buffer
 - Producer needs to wait if buffer is full
 - Consumer needs to wait if buffer is empty
- Example: Coke machine
 - Producer can put limited number of cokes in machine
 - Consumer can't take cokes out if machine is empty
 - They cannot work at the same time.

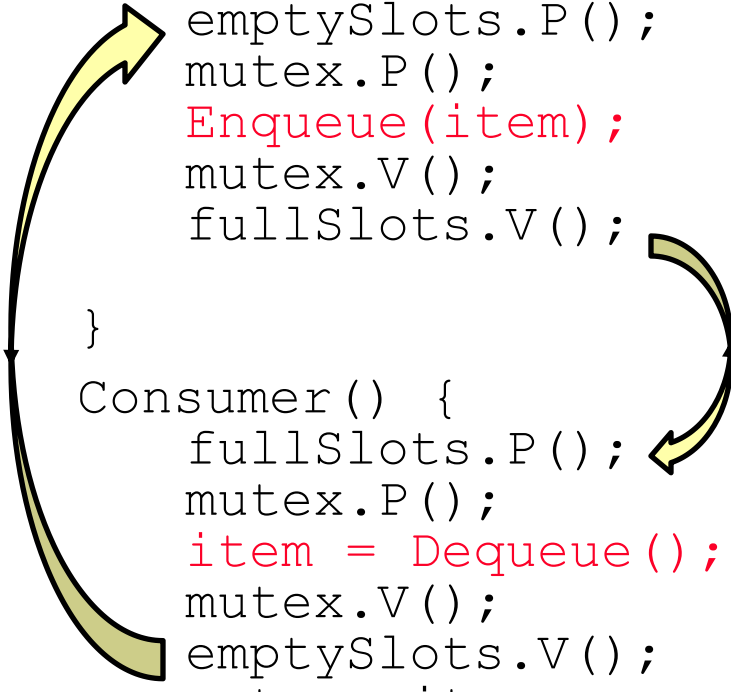


Correctness constraints for solution

- Correctness Constraints:
 - Consumer must wait for producer to fill buffer, if none (scheduling constraint)
 - Producer must wait for consumer to make room in buffer, if all full (scheduling constraint)
 - Only one thread can manipulate buffer queue at a time (mutual exclusion)
- Remember why we need mutual exclusion
 - Because computers are stupid
 - Imagine if in real life: the delivery person is filling the machine and somebody comes up and tries to stick their money into the machine
- General rule of thumb:
Use a separate semaphore for each constraint
 - Semaphore fullSlots; // consumer's constraint
 - Semaphore emptySlots; // producer's constraint
 - Semaphore mutex; // mutual exclusion

Full Solution to Bounded Buffer

```
Semaphore fullSlots = 0;    // Initially, no coke
Semaphore emptySlots = bufSize;
                                // Initially, num empty slots
Semaphore mutex = 1;        // No one using machine
```



```
Producer(item) {
    emptySlots.P();           // Wait until space
    mutex.P();               // Wait until machine free
    Enqueue(item);           // Tell consumers there is
    mutex.V();               // more coke
    fullSlots.V();
}

Consumer() {
    fullSlots.P();           // Check if there's a coke
    mutex.P();               // Wait until machine free
    item = Dequeue();
    mutex.V();
    emptySlots.V();          // tell producer need more
    return item;
}
```

Discussion about Solution

- Is order of P's important?
 - Yes! May cause deadlock when the buffer is full and a producer comes first.

```
Producer(item) {  
    mutex.P();  
    emptySlots.P();  
    Enqueue(item);  
    mutex.V();  
    fullSlots.V();  
}  
  
Consumer() {  
    fullSlots.P();  
    mutex.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```


Discussion about Solution

- Is order of P's important?
 - Yes! May cause deadlock when the buffer is empty.

```
Producer(item) {  
    emptySlots.P();  
    mutex.P();  
    Enqueue(item);  
    mutex.V();  
    fullSlots.V();  
}  
  
Consumer() {  
    mutex.P();  
    fullSlots.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

Discussion about Solution

- Is order of P's important?
 - Yes! May cause deadlock when the buffer is full or empty.

```
Producer(item) {  
    mutex.P();  
    emptySlots.P();  
    Enqueue(item);  
    mutex.V();  
    fullSlots.V();  
}  
  
Consumer() {  
    mutex.P();  
    fullSlots.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

Discussion about Solution

- Is order of V's important?
 - No, except that it might affect scheduling efficiency, when the buffer is empty and the consumer comes first.

```
Producer(item) {  
    emptySlots.P();  
    mutex.P();  
    Enqueue(item);  
    fullSlots.V();  
    mutex.V();  
}  
  
Consumer() {  
    fullSlots.P();  
    mutex.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

Discussion about Solution

- Is order of V's important?
 - No, except that it might affect scheduling efficiency, when the buffer is full.

```
Producer(item) {  
    emptySlots.P();  
    mutex.P();  
    Enqueue(item);  
    mutex.V();  
    fullSlots.V();  
}  
  
Consumer() {  
    fullSlots.P();  
    mutex.P();  
    item = Dequeue();  
    emptySlots.V();  
    mutex.V();  
    return item;  
}
```

Motivation for Monitors and Condition Variables

- Semaphores are a huge step up; just think of trying to do the bounded buffer with only loads and stores
- Problem is that semaphores are low-level:
 - They are used for both mutex and scheduling constraints
 - How do you prove correctness to someone?
 - » Example: the fact that flipping of P's in bounded buffer gives deadlock is not immediately obvious.
- Need a cleaner idea ...

Motivation for Monitors and Condition Variables

- Cleaner idea: Use *locks* for mutual exclusion and *condition variables* for scheduling constraints
- **Monitor**: a lock and zero or more condition variables for managing concurrent access to shared data
 - Some languages like Java provide this natively – an object (instance of a class) whose methods are implemented with mutual exclusion (using “**synchronized**” keyword)
 - » To learn how to achieve synchronization using monitor, we cannot use this keyword in our projects. We need to define our own lock and condition variables.
 - Most others use actual locks and condition variables

Simple Monitor Example (version 1)

- Here is an (infinite) synchronized queue

```
Lock lock;
Queue queue;                                // shared data

AddToQueue(item) {
    lock.Acquire();                          // Get Lock
    queue.enqueue(item);                     // Add item
    lock.Release();                          // Release Lock
}

RemoveFromQueue() {
    lock.Acquire();                          // Get Lock
    item = queue.dequeue();                  // Get next item
    lock.Release();                          // Release Lock
    return(item);                            // Might return null
}
```

- Not very interesting use of “Monitor”
 - Only uses a lock and no condition variables
 - Cannot put consumer to sleep if nothing to dequeue!

Condition Variables

- How do we change the `RemoveFromQueue()` routine to wait until something is on the queue?
- **Condition variable**: a variable `x` that implements:
 - `x.wait()`: Wait on a condition (go to sleep?)
 - `x.signal()`: Wake up one waiter, if any
 - Many threads can call `x.wait()`, they will be queued up waiting for a call to `x.signal()`. That call will start the first waiting thread.
- Some systems also implement:
 - `broadcast()` to wake up all waiting threads
- To support sleeping while waiting inside critical section, we add
 - `x.wait(&lock)`: Atomically release the lock and go to sleep. Re-acquire lock later, before returning.
- Rule: Must hold lock when doing condition variable operations

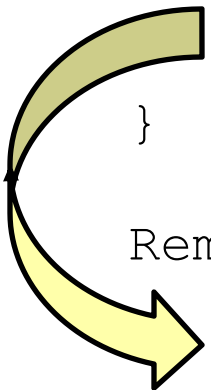
Complete Monitor Example (condition variable)

- Here is an (infinite) synchronized queue

```
Lock lock;
Condition dataready;
Queue queue;                                // shared data

AddToQueue(item) {
    lock.Acquire();                          // Get Lock
    queue.enqueue(item);                     // Add item
    dataready.signal();                      // Signal any waiters
    lock.Release();                          // Release Lock
}

RemoveFromQueue() {
    lock.Acquire();                          // Get Lock
    while (queue.isEmpty()) {                // If nothing, sleep
        dataready.wait(&lock);
    }
    item = queue.dequeue();                  // Get next item
    lock.Release();                          // Release Lock
    return(item);
}
```



Summary

- Semaphores: Like integers with restricted interface
 - Two operations:
 - » `P()` : Wait if zero; decrement when becomes non-zero
 - » `V()` : Increment and wake a sleeping task (if exists)
 - » Can initialize value to any non-negative value
 - Use separate semaphore for each constraint
- Monitors: A lock plus one or more condition variables
 - Always acquire lock before accessing shared data
 - Use condition variables to wait inside critical section
 - » Three Operations: `Wait()`, `Signal()`, and `Broadcast()`

Next

- Reading: Chapter 7.