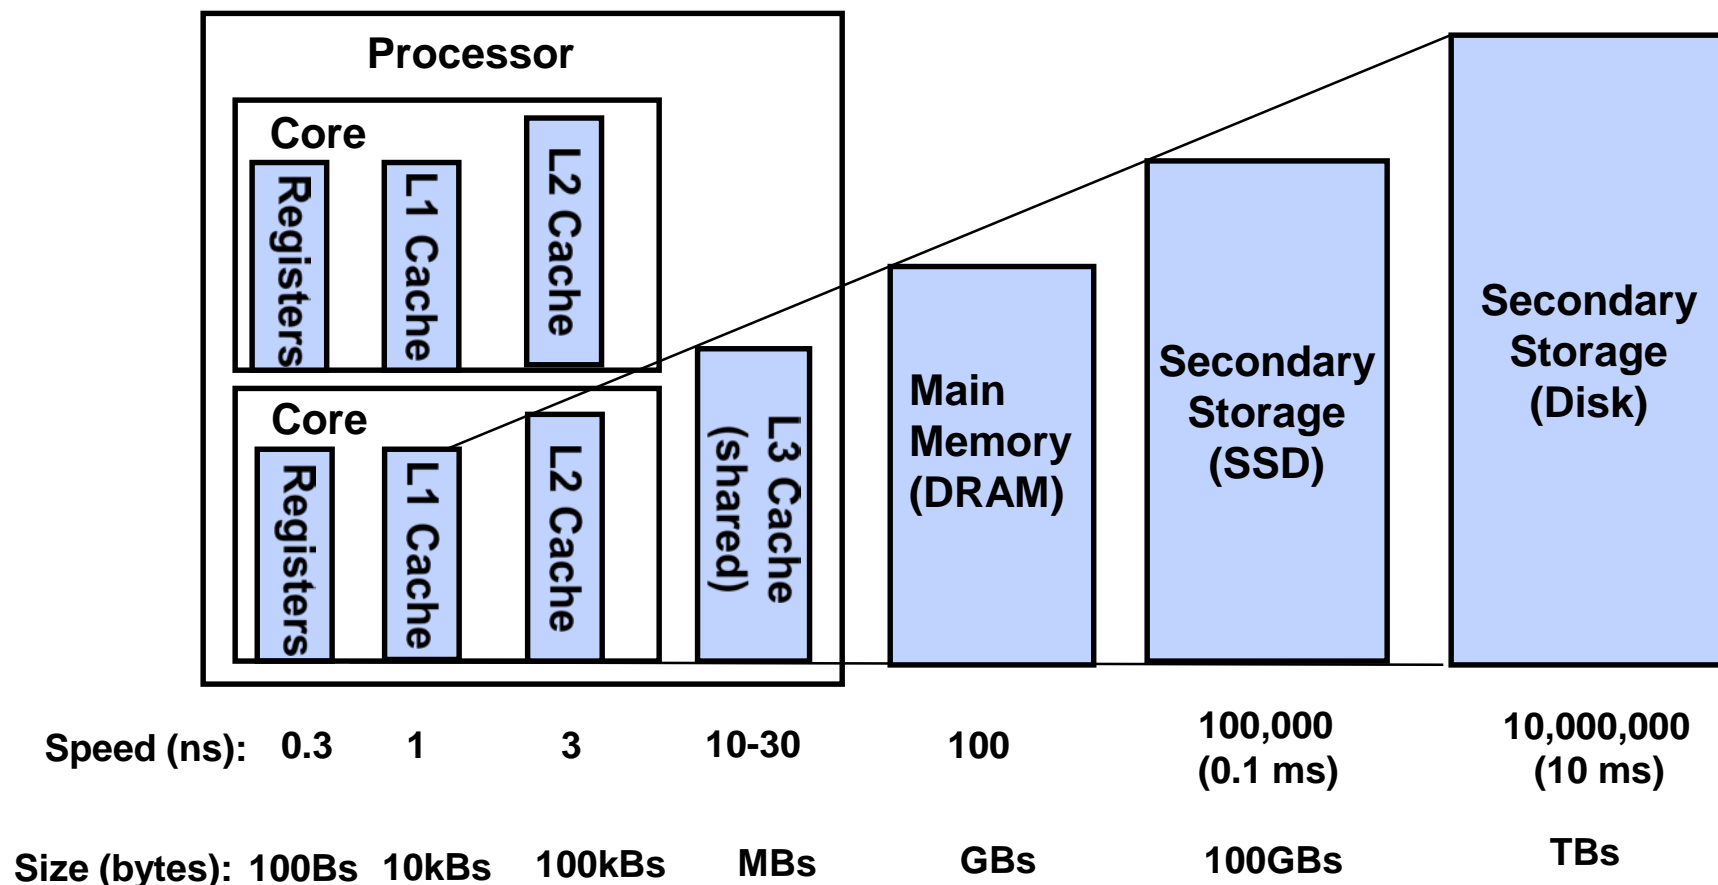


CSE150
Operating Systems
Lecture 15


Caches and TLBs

Memory Hierarchy

- Take advantage of the principle of locality to:
 - Present as much memory as in the cheapest technology
 - Provide access at speed offered by the fastest technology




Caching Concept




- **Cache**: a repository for copies that can be accessed more quickly than the original
 - Make frequent case fast and infrequent case less dominant
- Caching at different levels
 - Can cache: memory locations, address translations, pages, file blocks, file names, network routes, etc...
- Only good if:
 - Frequent case frequent enough and
 - Infrequent case not too expensive
- Important measure: Average Access time =
 $(\text{Hit Rate} \times \text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Time})$

Example

- Data in memory, no cache:



- Data in memory, 10ns cache:




Average Access time =


(Hit Rate x HitTime) + (Miss Rate x MissTime)

- HitRate + MissRate = 1

Why Does Caching Help? Locality!




- **Temporal Locality** (Locality in Time):
 - Keep recently accessed data items closer to processor
- **Spatial Locality** (Locality in Space):
 - Move contiguous blocks to the upper levels



Caching Questions

- 8 byte cache
- 32 byte memory
- 1 block = 1 byte
- Assume CPU accesses 01100



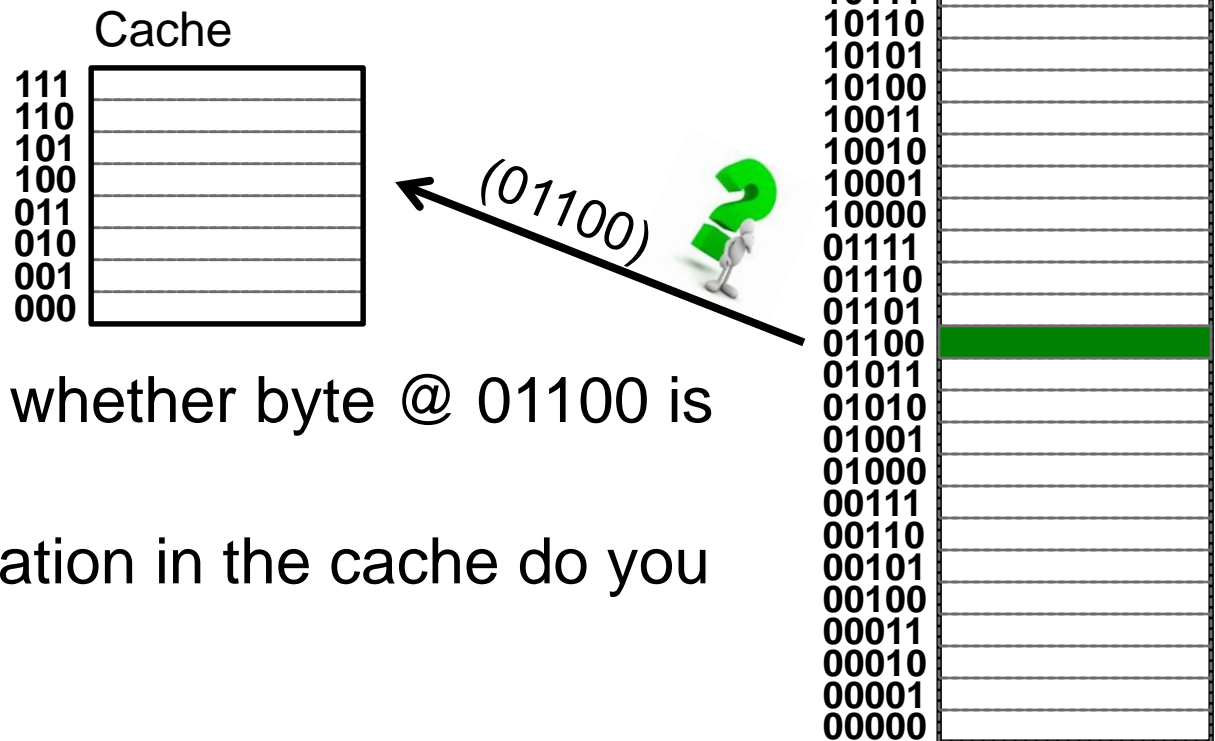
1. How do you know whether byte @ 01100 is cached?

Physical Memory

11111	
11110	
11101	
11100	
11011	
11010	
11001	
11000	
10111	
10110	
10101	
10100	
10011	
10010	
10001	
10000	
01111	
01110	
01101	
01100	
01011	
01010	
01001	
01000	
00111	
00110	
00101	
00100	
00011	
00010	
00001	
00000	

Caching Questions


- 8 byte cache
- 32 byte memory
- 1 block = 1 byte
- Assume CPU accesses 01100



1. How do you know whether byte @ 01100 is cached?
2. If not, at which location in the cache do you place the byte?

Caching Questions

- 8 byte cache
- 32 byte memory
- 1 block = 1 byte
- Assume CPU accesses 01100



1. How do you know whether byte @ 01100 is cached?
2. If not, at which location in the cache do you place the byte?
3. If cache full, which cached byte do you evict?

Physical Memory

11111	
11110	
11101	
11100	
11011	
11010	
11001	
11000	
10111	
10110	
10101	
10100	
10011	
10010	
10001	
10000	
01111	
01110	
01101	
01100	
01011	
01010	
01001	
01000	
00111	
00110	
00101	
00100	
00011	
00010	
00001	
00000	

Simple Example: Direct Mapped Cache

- Each byte (block) in physical memory is cached to a **single** cache location
 - Least significant bits of address (last 3 bits) index the cache
 - (00**100**), (01**100**), (10**100**), (11**100**) cached to 100

Index Tag Cache

111
110
101
100
011
010
001
000


- How do you know which byte is cached?
 - Cache stores the most significant two bits (i.e., **tag**) of the cached byte

Physical Memory

11111	
11110	
11101	
11 100	
11011	
11010	
11001	
11000	
10111	
10110	
10101	
10 100	
10011	
10010	
10001	
10000	
01111	
01110	
01101	
01 100	
01011	
01010	
01001	
01000	
00111	
00110	
00101	
00 100	
00011	
00010	
00001	
00000	

Simple Example: Direct Mapped Cache

- Each byte (block) in physical memory is cached to a **single** cache location
 - Least significant bits of address (last 3 bits) index the cache
 - (00**100**), (01**100**), (10**100**), (11**100**) cached to 100




- How do you know whether (01100) is cached?
 - Check **tag** associated with index 100
- At which cache location do you place (01100)?
 - 100
- If cache full, which cached byte do you evict?
 - 100

Simple Example: Direct Mapped Cache

- Each byte (block) in physical memory is cached to a **single** cache location
 - Least significant bits of address (last 3 bits) index the cache
 - (00**100**), (01**100**), (10**100**), (11**100**) cached to 100


Index Tag Cache

111
110
101
100
011
010
001
000




- How do you know which byte is cached?
 - Cache stores the most significant two bits (i.e., **tag**) of the cached byte

Physical Memory




Simple Example: Direct Mapped Cache

- Each byte (block) in physical memory is cached to a **single** cache location
 - Least significant bits of address (last 3 bits) index the cache
 - (00**100**), (01**100**), (10**100**), (11**100**) cached to 100



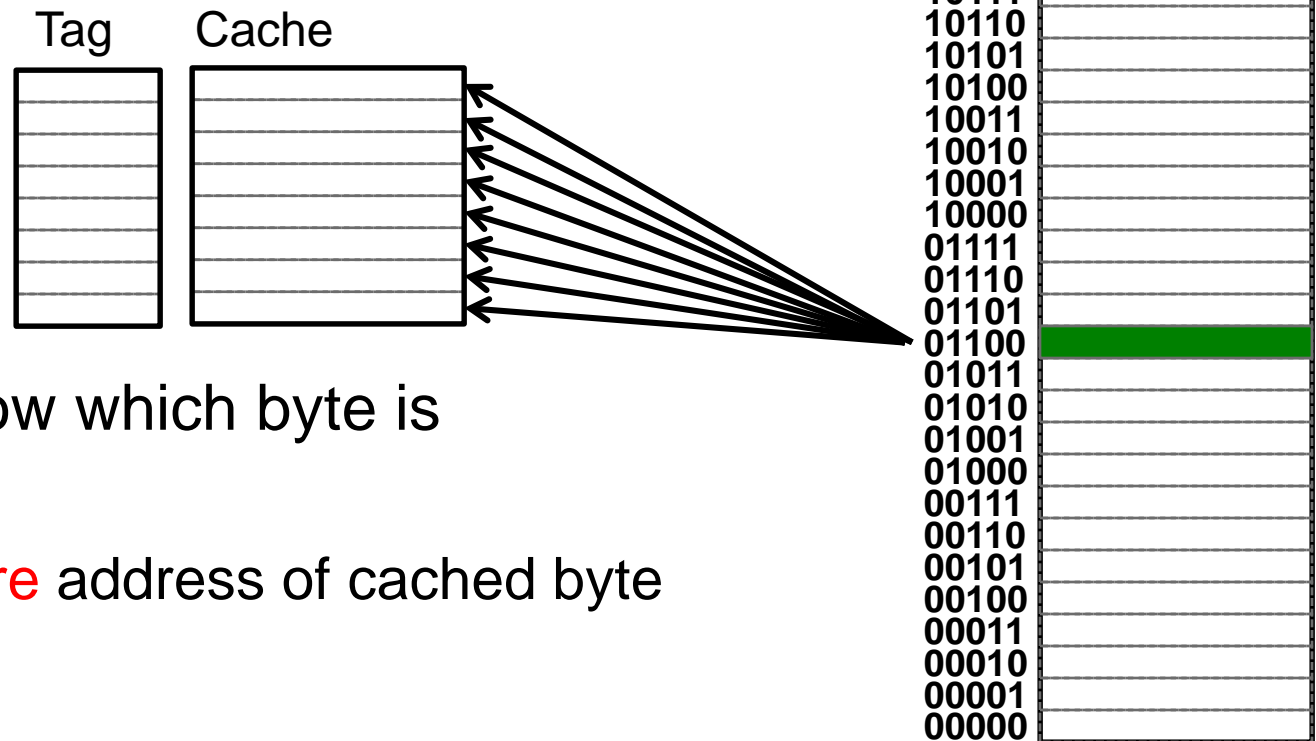
Physical Memory



- How do you know whether (01100) is cached?
 - Check **tag** associated with index 100
- At which cache location do you place (01100)?
 - 100
- If cache full, which cached byte do you evict?
 - 100

Simple Example: Fully Associative Cache

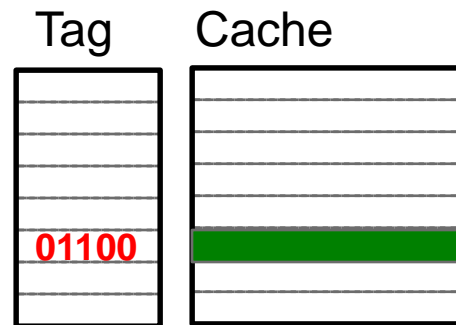
- Each byte can be stored at **any** location in the cache



- How do you know which byte is cached?
 - Tag store **entire** address of cached byte

Simple Example: Fully Associative Cache

- Each byte can be stored at **any** location in the cache



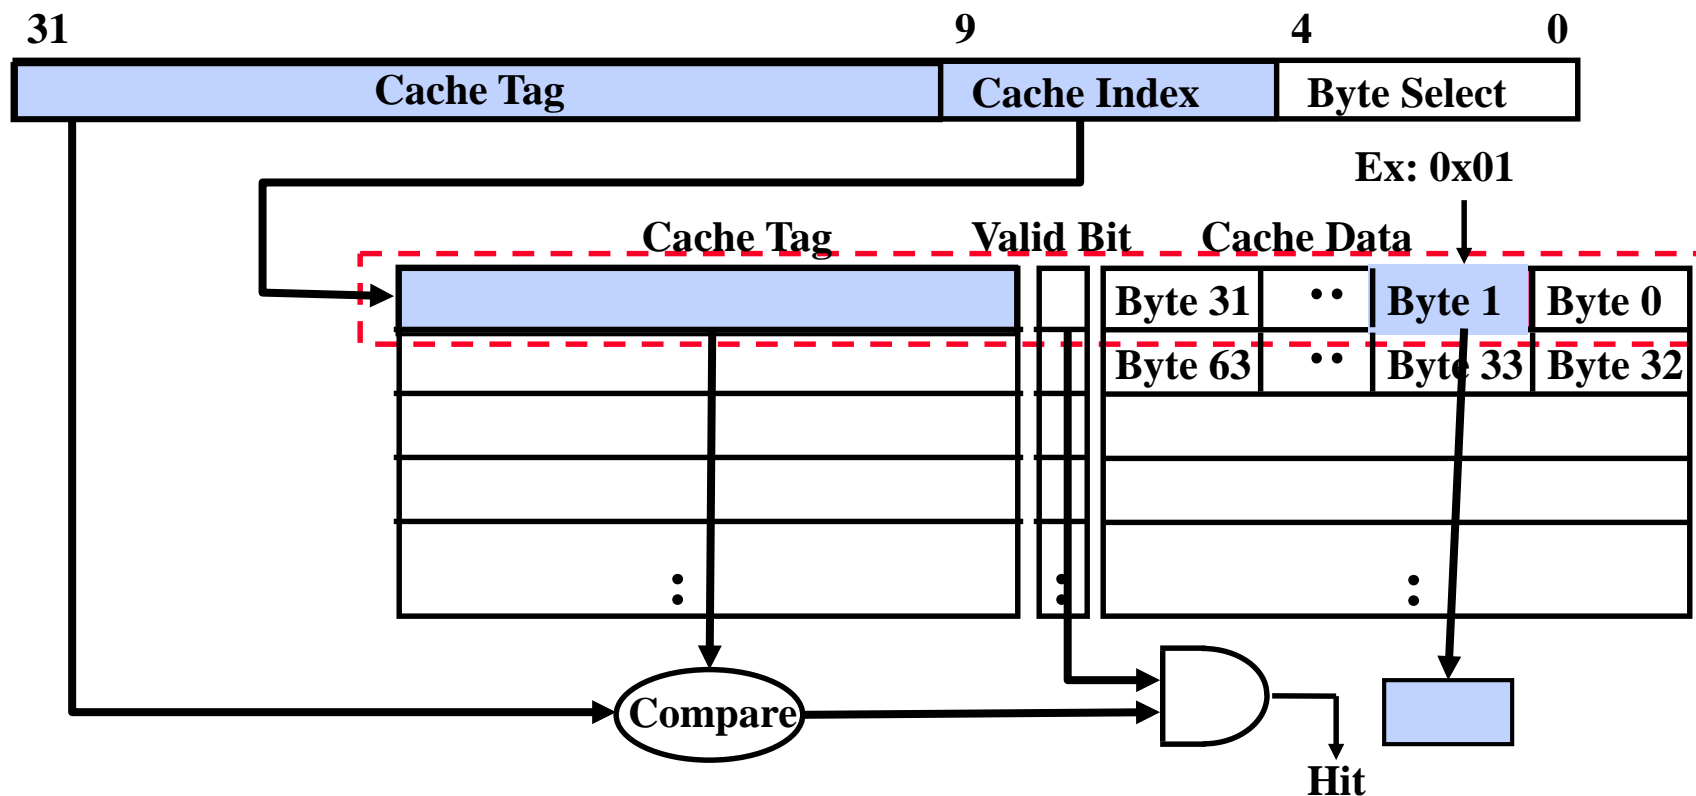
Physical Memory

11111	
11110	
11101	
11100	
11011	
11010	
11001	
11000	
10111	
10110	
10101	
10100	
10011	
10010	
10001	
10000	
01111	
01110	
01101	
01100	
01011	
01010	
01001	
01000	
00111	
00110	
00101	
00100	
00011	
00010	
00001	
00000	

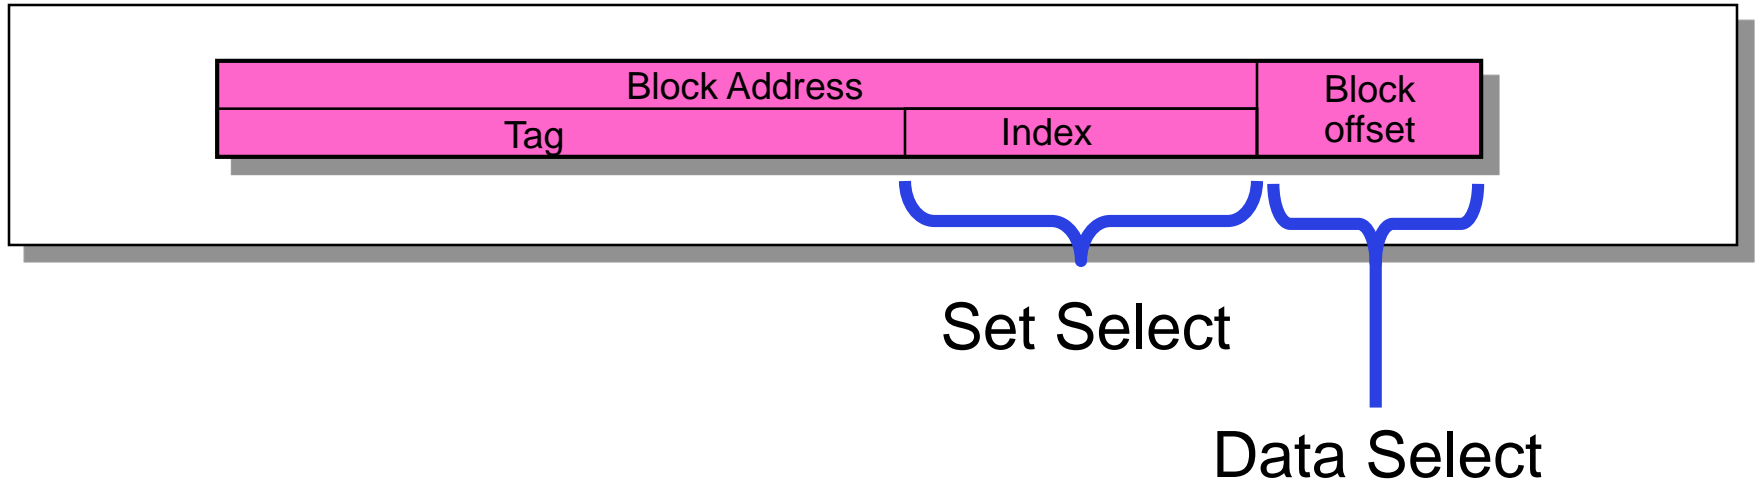
- How do you know whether (01100) is cached?
 - Check **tag** of all cache entries
- At which cache location do you place (01100)?
 - Any
- If cache full, which cached byte do you evict?
 - Specific eviction policy

Direct Mapped Cache

- Cache index selects a cache block
- “Byte select” selects byte within cache block
 - Example: Block Size=32B blocks
- Cache tag fully identifies the cached data
- Data with same “cache index” shares the same cache entry
 - Conflict misses



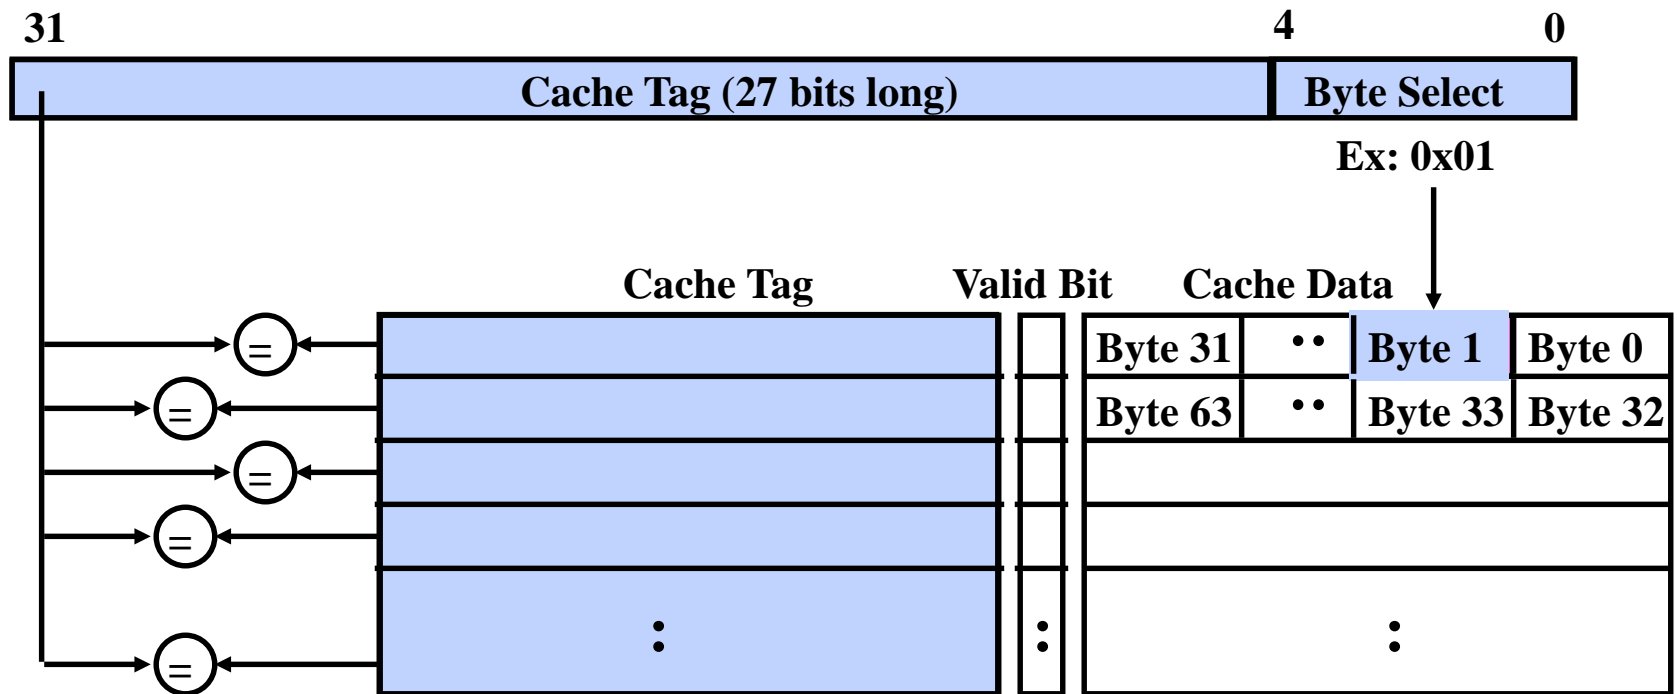
How is a Block found in a Direct Mapped Cache?



- Index Used to Lookup Candidates in Cache
 - Index identifies the set
- Tag used to identify actual copy
 - If no candidates match, then declare cache miss
- Block is minimum quantum of caching
 - Data select field used to select data within block
 - Many caching applications don't have data select field

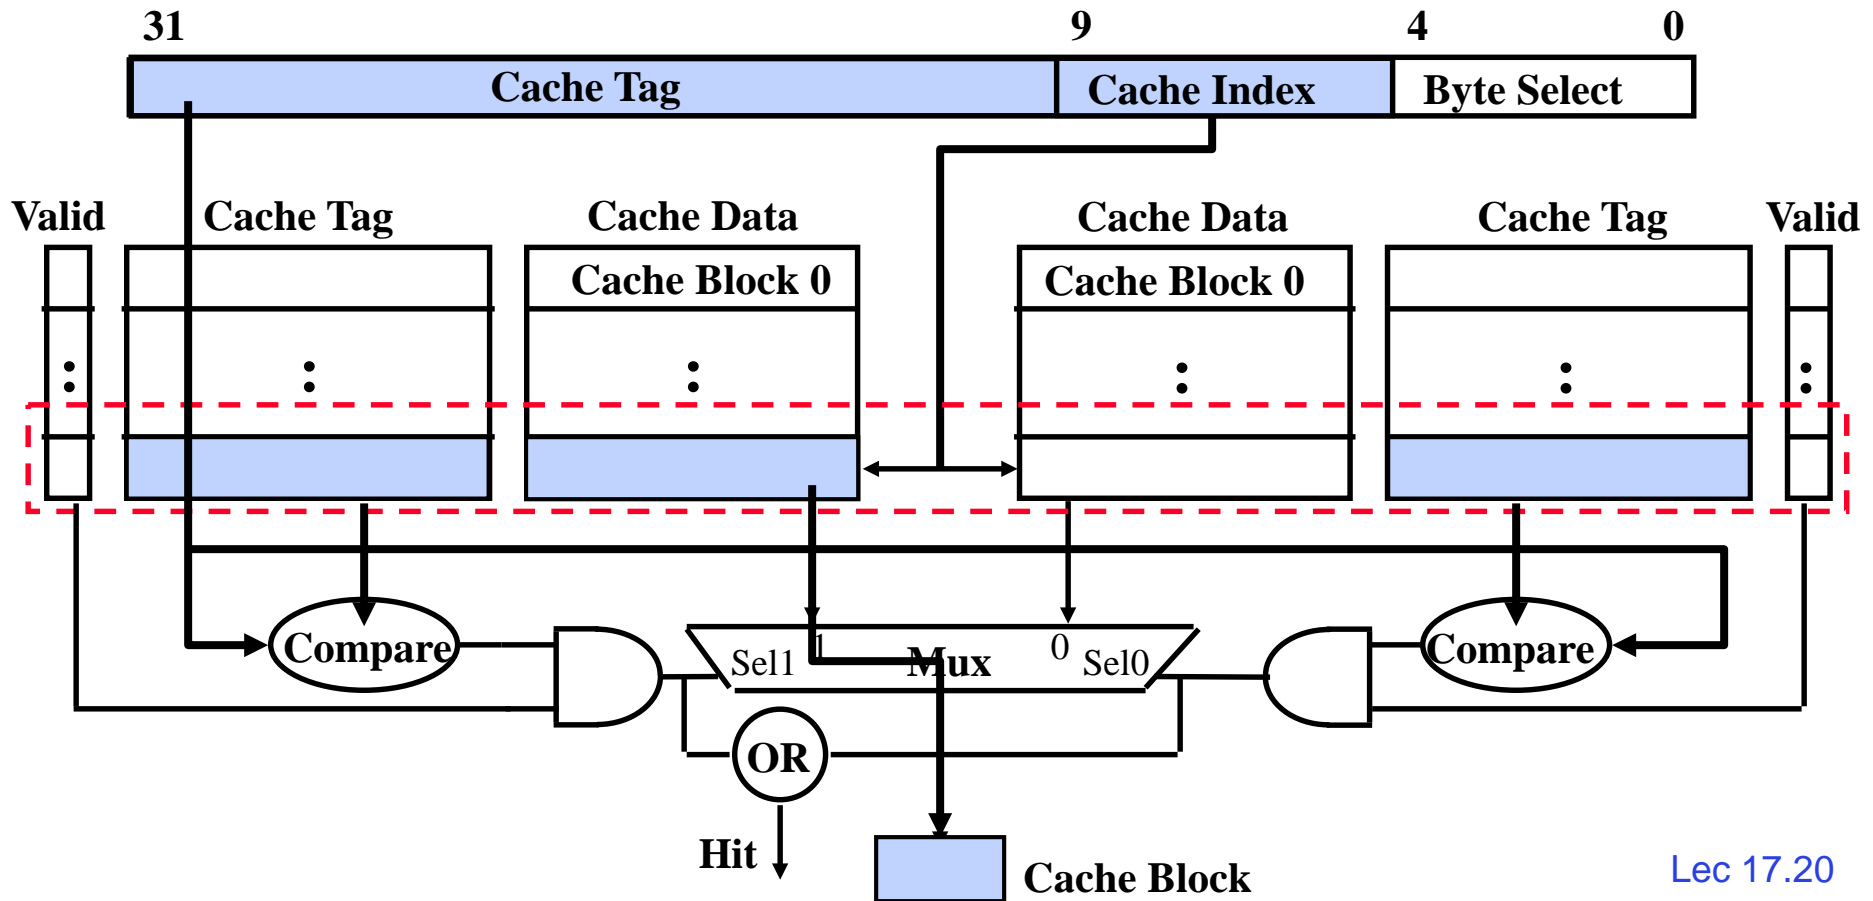
Fully Associative Cache

- **Fully Associative:** Every block can hold any line
 - Address does not include a cache index
 - Compare Cache Tags of all Cache Entries in Parallel
- Example: Block Size=32B blocks
 - We need N 27-bit comparators



Set Associative Cache


- **N-way set associative:** N entries per Cache Index
 - N direct mapped caches operates in parallel
- **Example: Two-way set associative cache**
 - Two tags in the set are compared to input in parallel
 - Data is selected based on the tag result



Where does a Block Get Placed in a Cache?

- Example: Block 12 placed in 8 block cache

32-Block Address Space:






Block no. 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1

Direct mapped:
block 12 (01100)
can go only into
block 4 (12 mod 8)

Set associative:
block 12 can go
anywhere in set 0

Fully associative:
block 12 can go
anywhere



Sources of Cache Misses

- **Compulsory** (cold start): first reference to a block
 - “Cold” fact of life: not a whole lot you can do about it
 - Note: When running “billions” of instruction, Compulsory Misses are insignificant
- **Capacity**:
 - Cache cannot contain all blocks access by the program
 - Solution: increase cache size
- **Conflict** (collision):
 - Multiple memory locations mapped to same cache location
 - Solutions: increase cache size, or increase associativity
- **Two others**:
 - **Coherence** (Invalidation): other process (e.g., I/O) updates memory
 - **Policy**: Due to non-optimal replacement policy

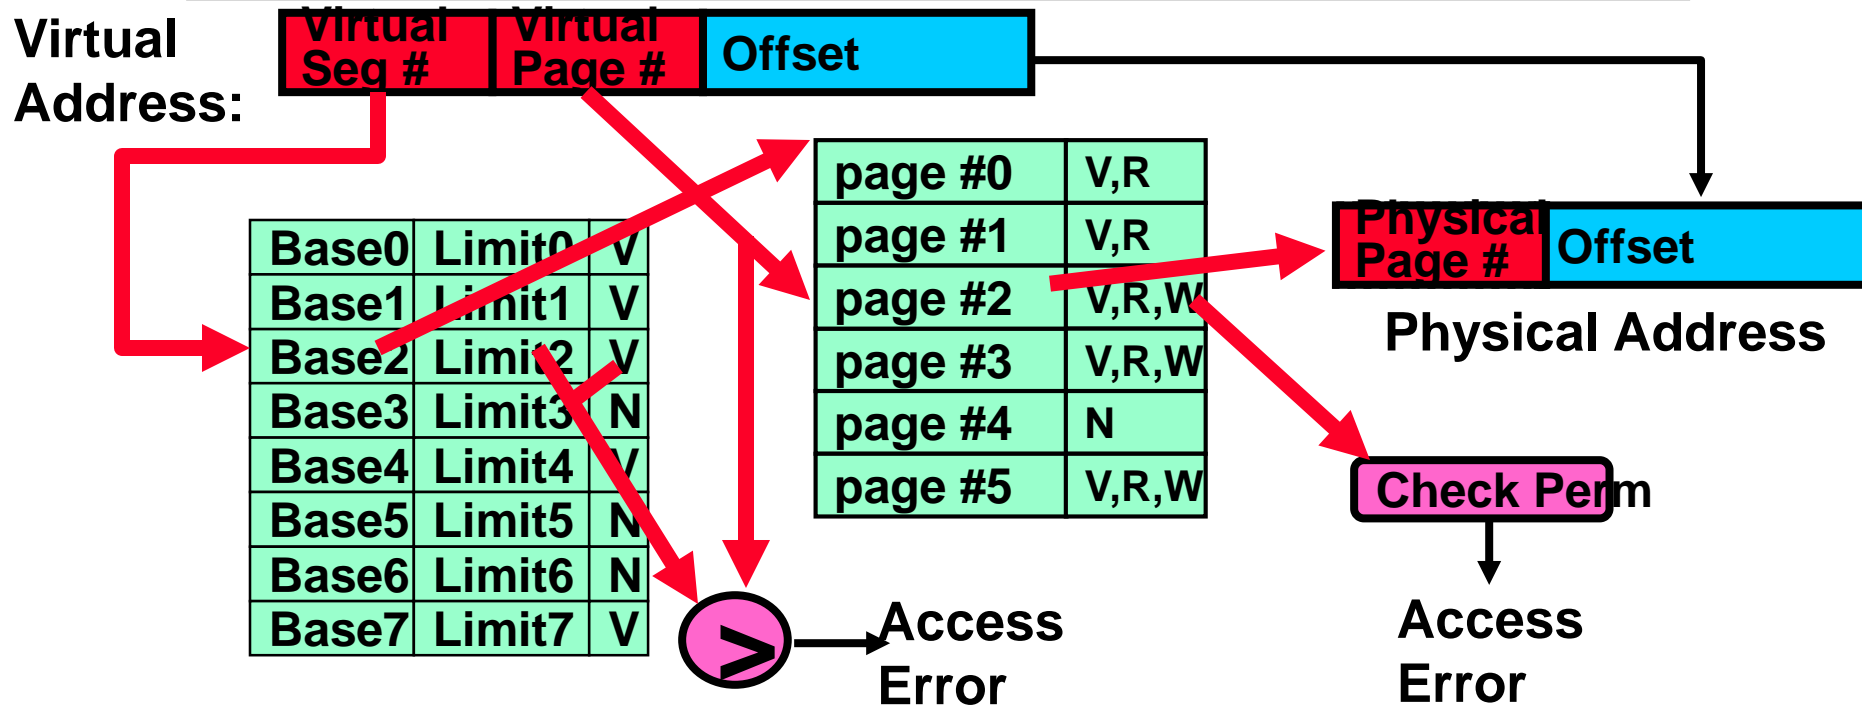
Which Block Should be Replaced on a Miss?

- Easy for Direct Mapped: Only one possibility
- Set Associative or Fully Associative:
 - Random
 - LRU (Least Recently Used)

What Happens on a Write?

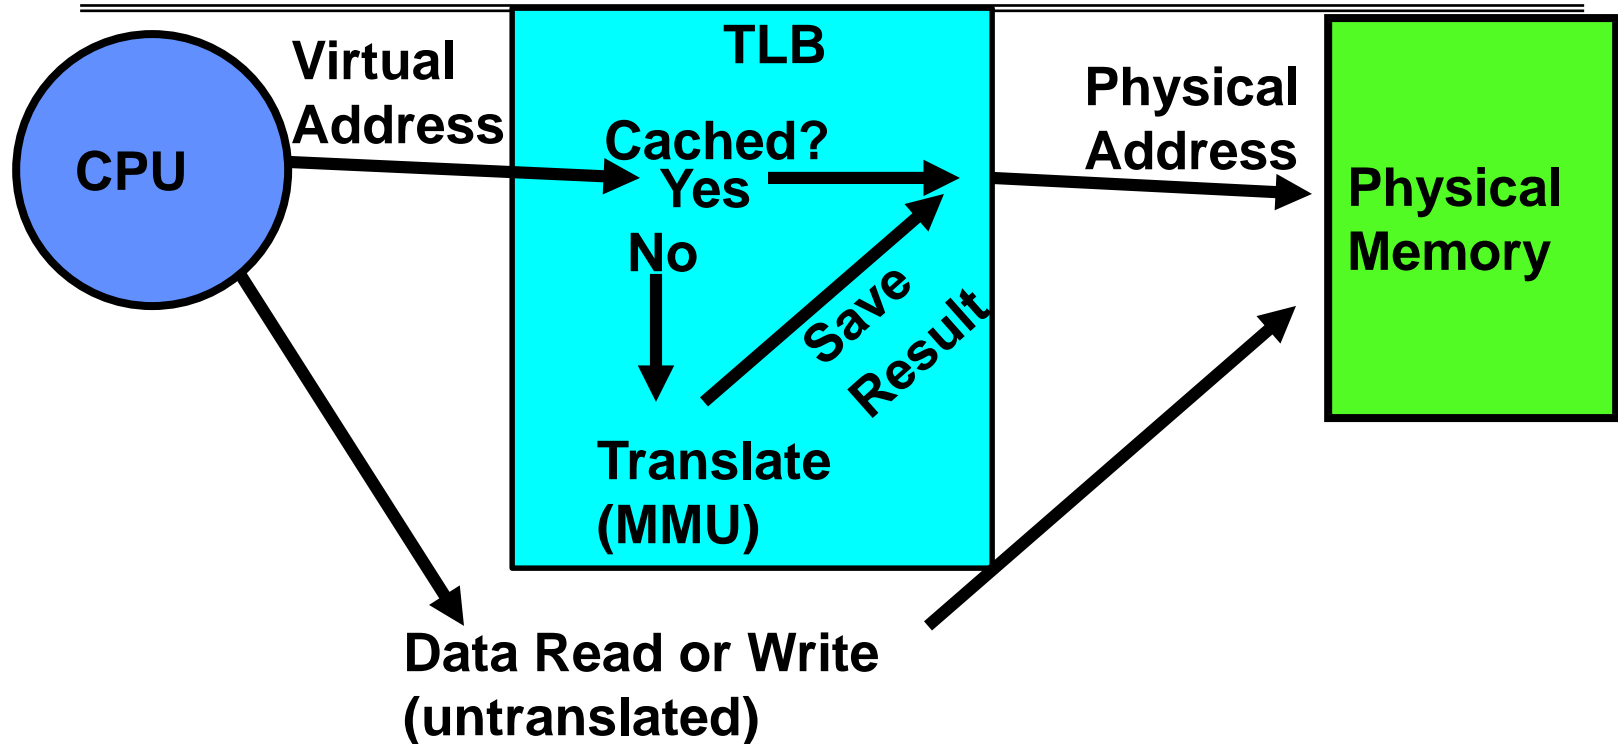
- **Write through:** The information is written both to the block in the cache and to the block in the lower-level memory
- **Write back:** The information is written only to the block in the cache.
 - Modified cache block is written to main memory only when it is replaced
- Pros and Cons of each?
 - WT:
 - » PRO: read misses cannot result in writes
 - » CON: processor held up on writes unless writes buffered
 - WB:
 - » PRO: repeated writes not sent to DRAM
processor not held up on writes
 - » CON: More complex

Another Major Reason to Deal with Caching



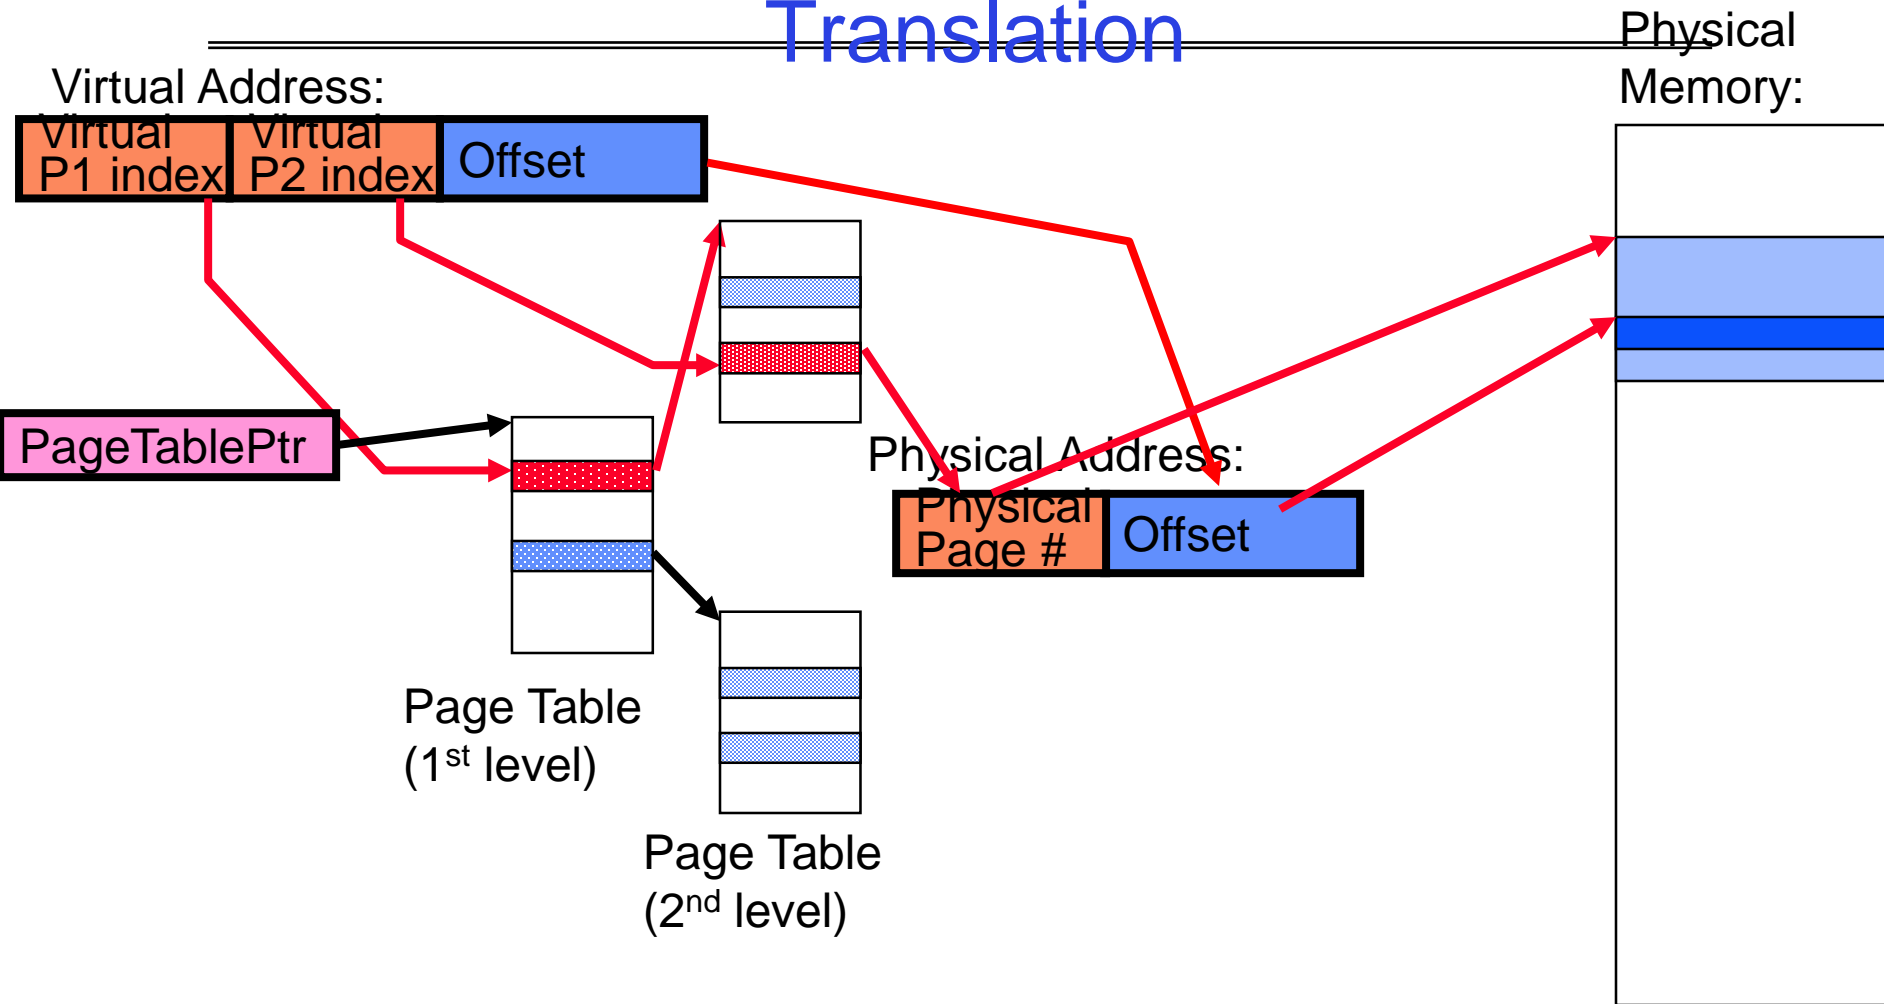
- Cannot afford to translate on every access
 - At least three DRAM accesses per actual DRAM access
- Solution? Cache translations!
 - Translation Cache: TLB (“Translation Lookaside Buffer”)

Caching Applied to Address Translation

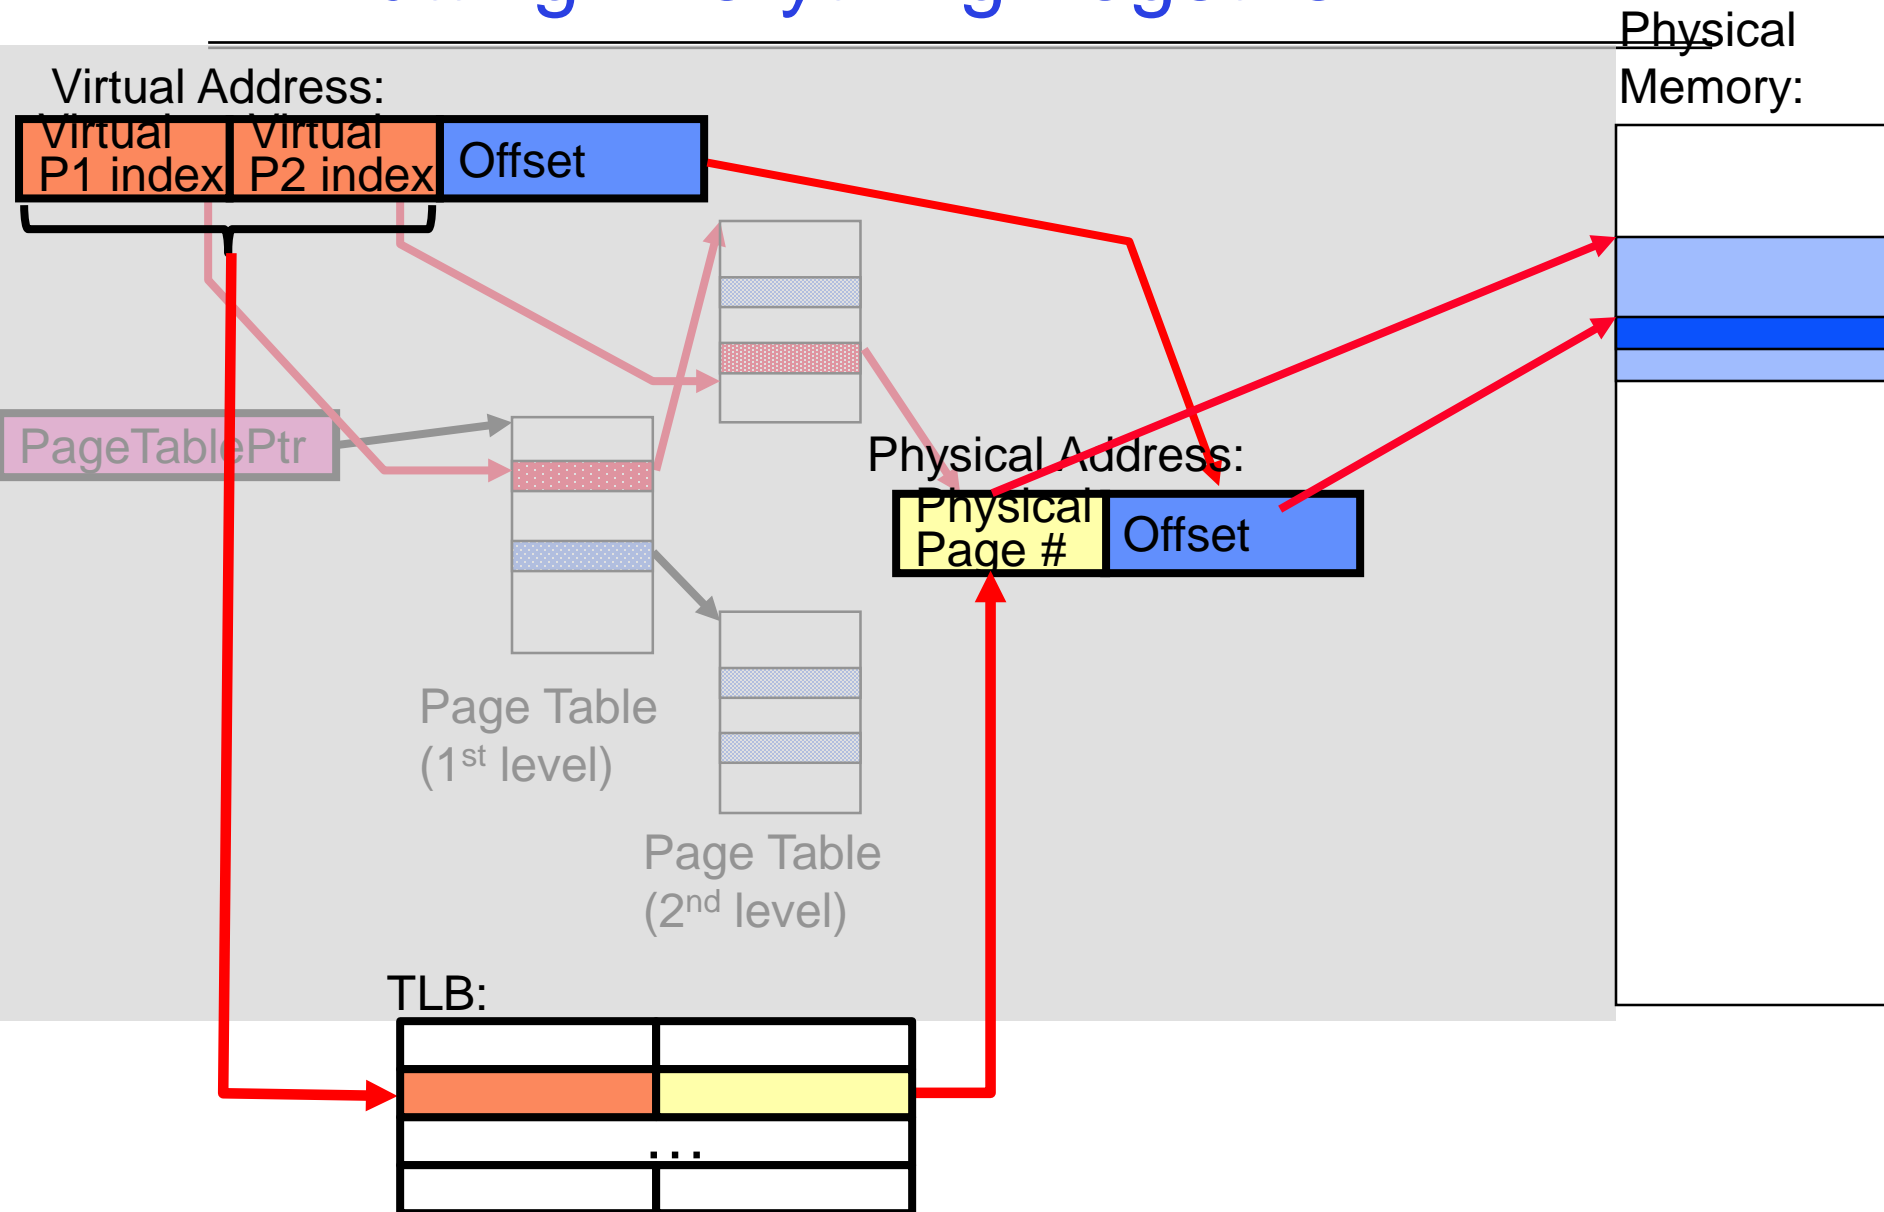


- Question is one of page locality: does it exist?
 - Instruction accesses spend a lot of time on the same page (since accesses sequential)
 - Stack accesses have definite locality of reference
 - Data accesses have less page locality, but still some...

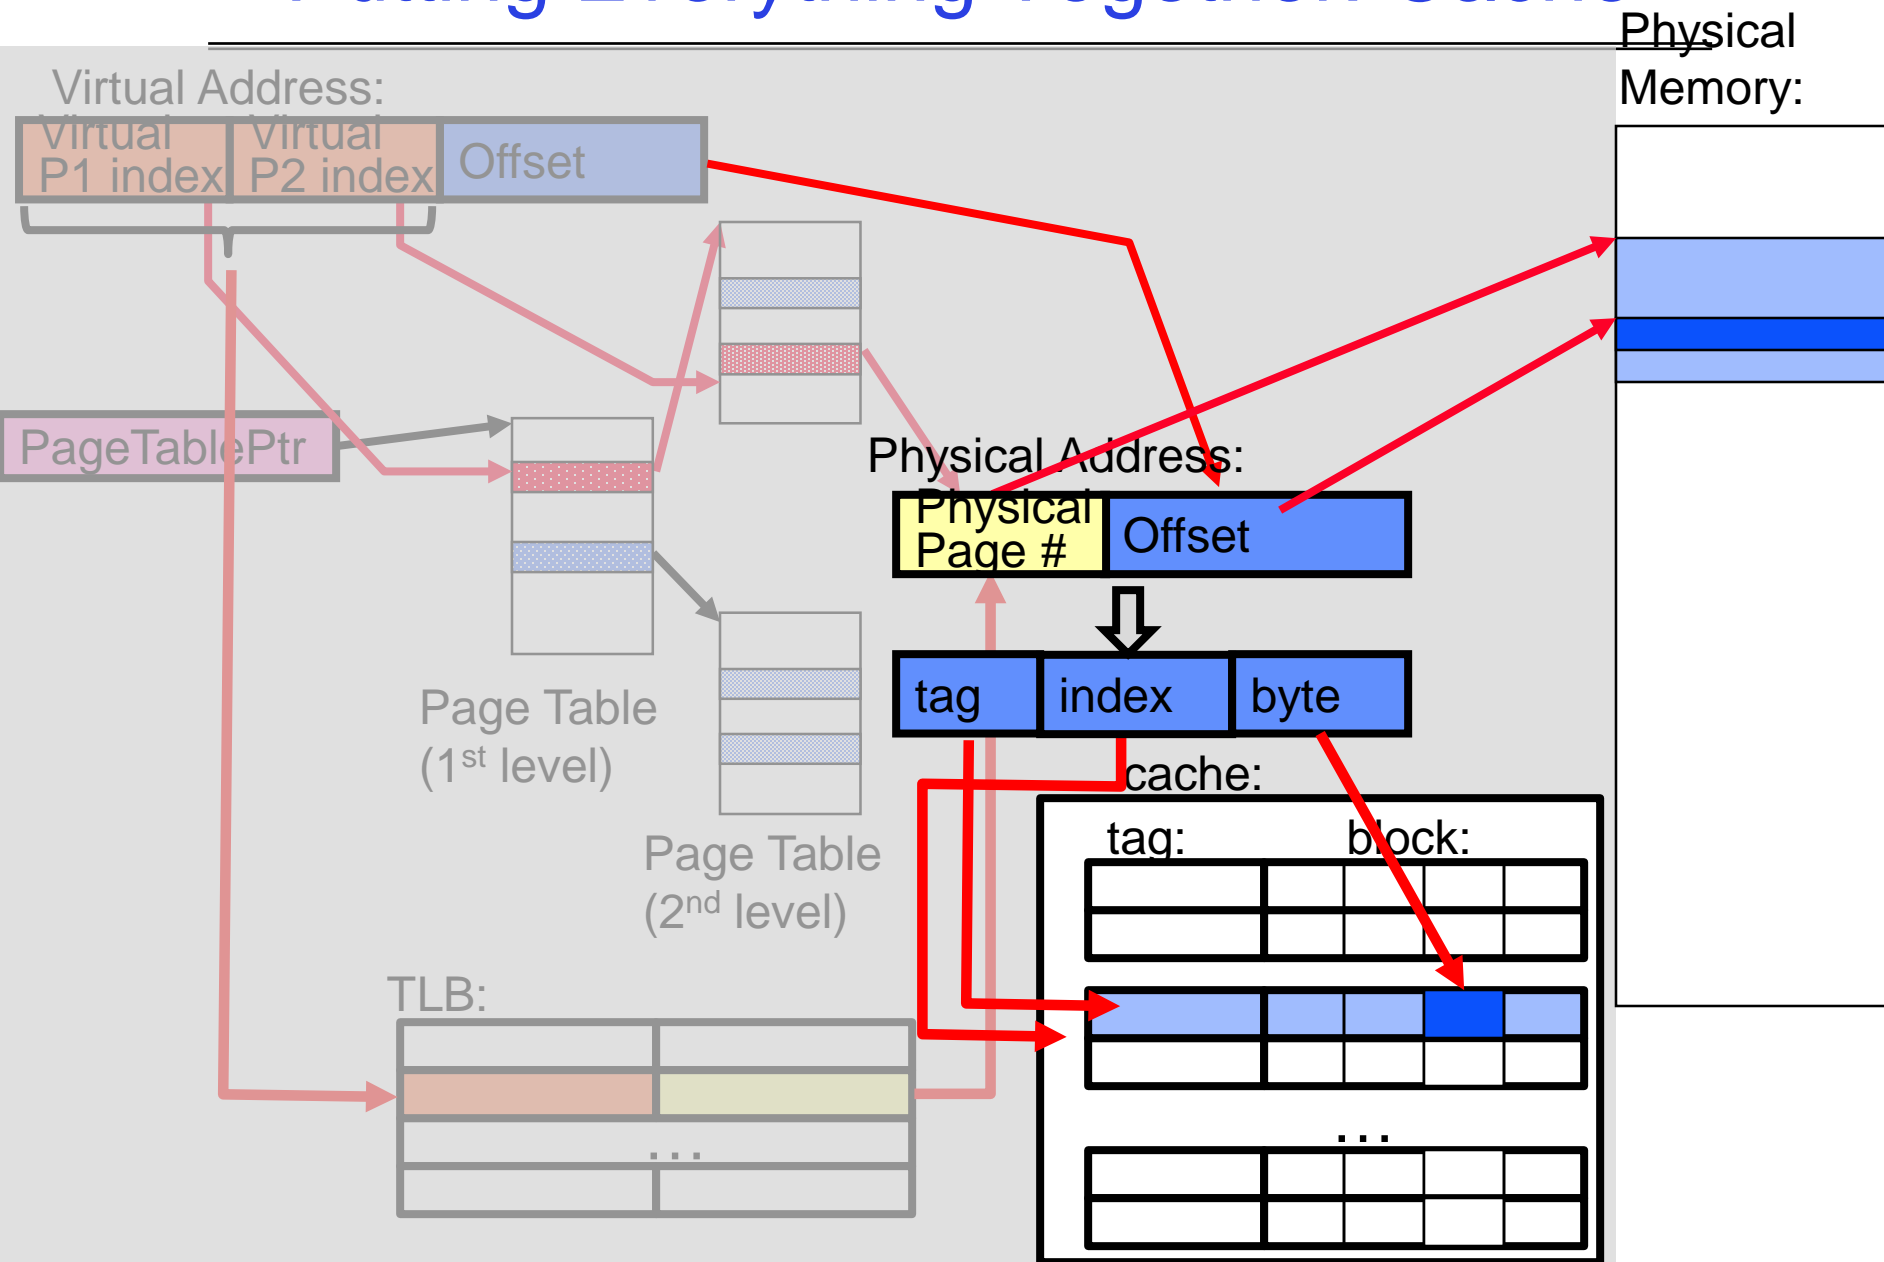
Putting Everything Together: Address Translation



Putting Everything Together: TLB



Putting Everything Together: Cache



What Actually Happens on a TLB Miss?

- Hardware traversed page tables:
 - On TLB miss, hardware in MMU looks at current page table to fill TLB (may walk multiple levels)
 - » If PTE valid, hardware fills TLB and processor never knows
 - » If PTE marked as invalid, causes Page Fault, after which kernel decides what to do afterwards
- Software traversed Page tables
 - On TLB miss, processor receives TLB fault
 - Kernel traverses page table to find PTE
 - » If PTE valid, fills TLB and returns from fault
 - » If PTE marked as invalid, internally calls Page Fault handler
- Most chip sets provide hardware traversal
 - Modern operating systems tend to have more TLB faults since they use translation for many things

What happens on a Context Switch?

- Need to do something, since TLBs map virtual addresses to physical addresses
 - Address Space just changed, so TLB entries no longer valid!
- Options?
 - Invalidate entire TLB (flush): simple but might be expensive
 - » What if switching frequently between processes?
 - Include ProcessID in TLB
 - » This is an architectural solution: needs hardware

Summary (1/2)

- The Principle of Locality:
 - Program likely to access a relatively small portion of the address space at any instant of time.
 - » **Temporal Locality**: Locality in Time
 - » **Spatial Locality**: Locality in Space
- Three (+1) Major Categories of Cache Misses:
 - **Compulsory Misses**: sad facts of life. Example: cold start misses.
 - **Conflict Misses**: increase cache size and/or associativity
 - **Capacity Misses**: increase cache size
 - **Coherence Misses**: Caused by external processors or I/O devices

Summary (2/2)

- Cache Organizations:
 - Direct Mapped: single block per set
 - Set associative: more than one block per set
 - Fully associative: all entries equivalent
- TLB is cache on address translations
 - Fully associative to reduce conflicts

Next week

- Chapter 9, including but not limited to the following key words.
 - Page replacement
 - Kernel user