

CSE150  
Operating Systems  
Lecture 19

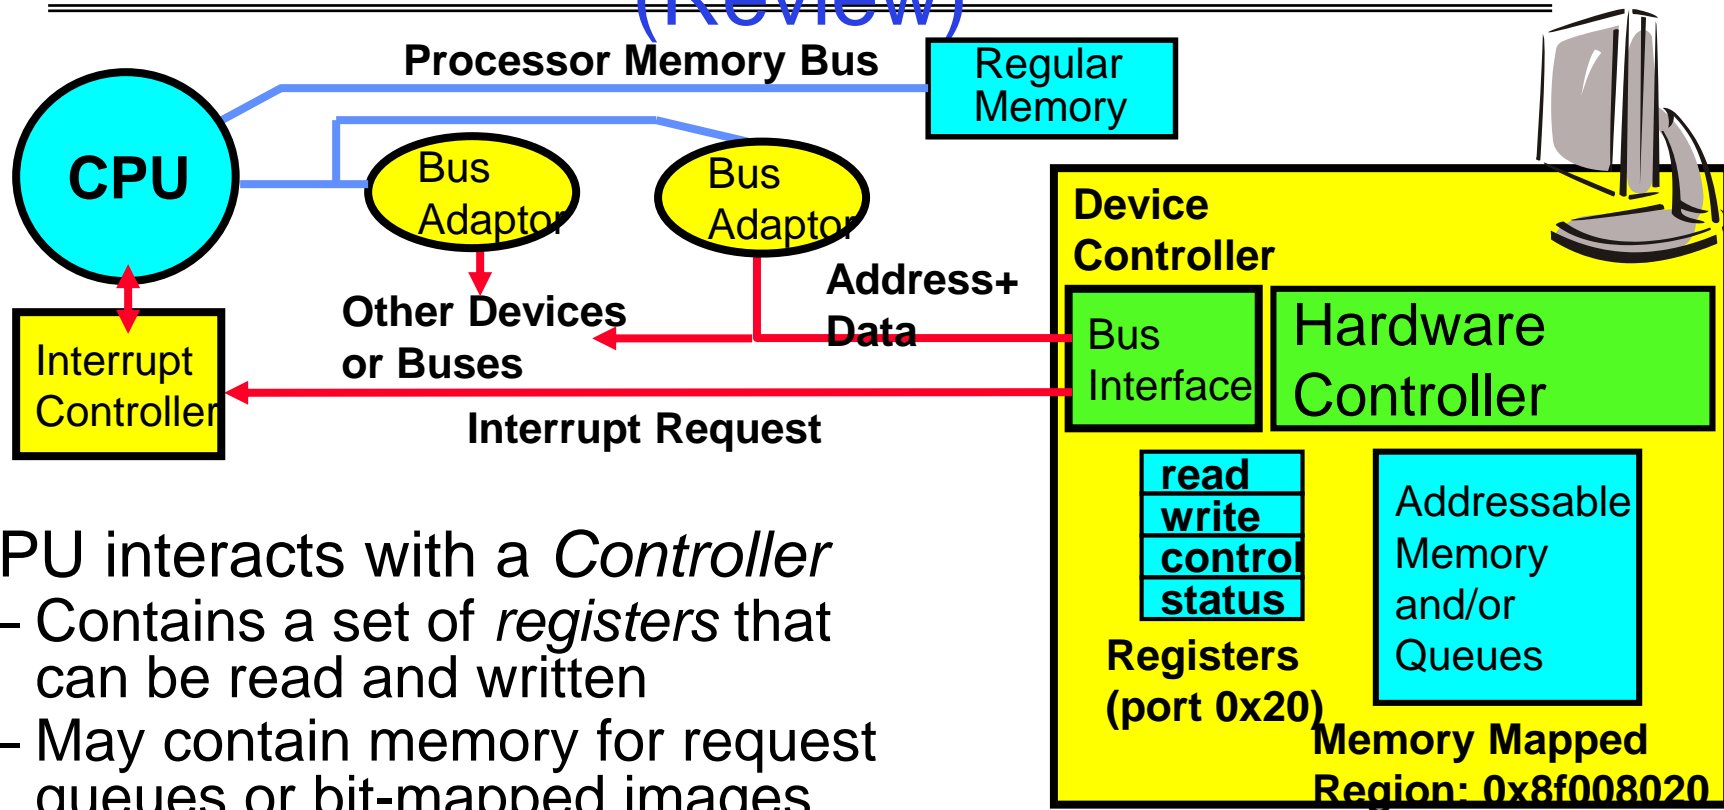
Disk/SSDs  
and File Systems

# Want Standard Interfaces to Devices (Review)

---

- **Block Devices:** e.g., disk drives, tape drives, DVD-ROM, flash
  - Access blocks of data
  - Commands include `open()`, `read()`, `write()`, `seek()`
  - Raw I/O or file-system access
  - Memory-mapped file access possible
- **Character/Byte Devices:** e.g., keyboards, mice, serial ports, some USB devices
  - Single characters at a time
  - Commands include `get()`, `put()`
  - Libraries layered on top allow line editing
- **Network Devices:** e.g., Ethernet, Wireless, Bluetooth
  - Different enough from block/character to have own interface
  - Unix and Windows include **socket** interface
    - » Separates network protocol from network operation
    - » Includes `select()` functionality

# How Does the Processor Talk to Devices? (Review)





- CPU interacts with a *Controller*
  - Contains a set of *registers* that can be read and written
  - May contain memory for request queues or bit-mapped images
- Regardless of the complexity of the connections and buses, processor accesses registers in two ways (IA):
  - **I/O instructions:** in/out instructions (e.g., Intel's 0x21, AL)
  - **Memory mapped I/O:** load/store instructions
    - » Registers/memory appear in physical address space
    - » I/O accomplished with load and store instructions

# Today

---

- Disks and SSDs
  - Hardware Performance Parameters
  - Disk Scheduling
- Important Storage Policies and Patterns
- File System Structures


# I/O Performance




**Response Time = Queue + I/O device service time**

- Performance of I/O subsystem
  - Metrics: Response Time, Throughput
  - Contributing factors to latency:
    - » Software paths (can be loosely modeled by a queue)
    - » Hardware controller
    - » I/O device service time
- Queuing behavior:
  - Can lead to big increases of latency as utilization approaches 100%
  - Solutions?

# I/O Performance



**Response Time = Queue + I/O device service time**



- Solutions?
  - Make everything faster 😊
  - Decouple systems
    - » multiple independent buses
    - » or tree-structured buses with higher root bandwidth
  - Buffering (as long as you don't have to wait for it) and spooling
    - » Give the processor something to do that gets the data “closer” to its endpoint.

# Hard Disk Drives (HDDs)



Western Digital Drive


<http://www.storagereview.com/guide/>

IBM Personal Computer/AT (1986)

30 MB hard disk - \$500

30-40ms seek time

0.7-1 MB/s (est.)




**Read/Write Head  
Side View**



**IBM/Hitachi Microdrive**

# Properties of a Magnetic Hard Disk





- Properties
  - Independently addressable element: **sector**
    - » OS always transfers groups of sectors together—“**blocks**”
  - A disk can access directly any given block either sequentially or randomly.
- Zoned bit recording
  - Constant bit density: more bits (sectors) on outer tracks
  - Speed varies with track location



# Magnetic Disk Characteristic


- Read/write: three-stage process:
  - **Seek time**: position the head/arm over the proper track (into proper cylinder)
  - **Rotational latency**: wait for the desired sector to rotate under the read/write head
  - **Transfer time**: transfer a block of bits (sector) under the read-write head
- Disk Latency = Queuing Time + Controller time + Seek Time + Rotation Time + Transfer Time



- Highest Bandwidth:
  - Transfer large group of blocks sequentially from one track

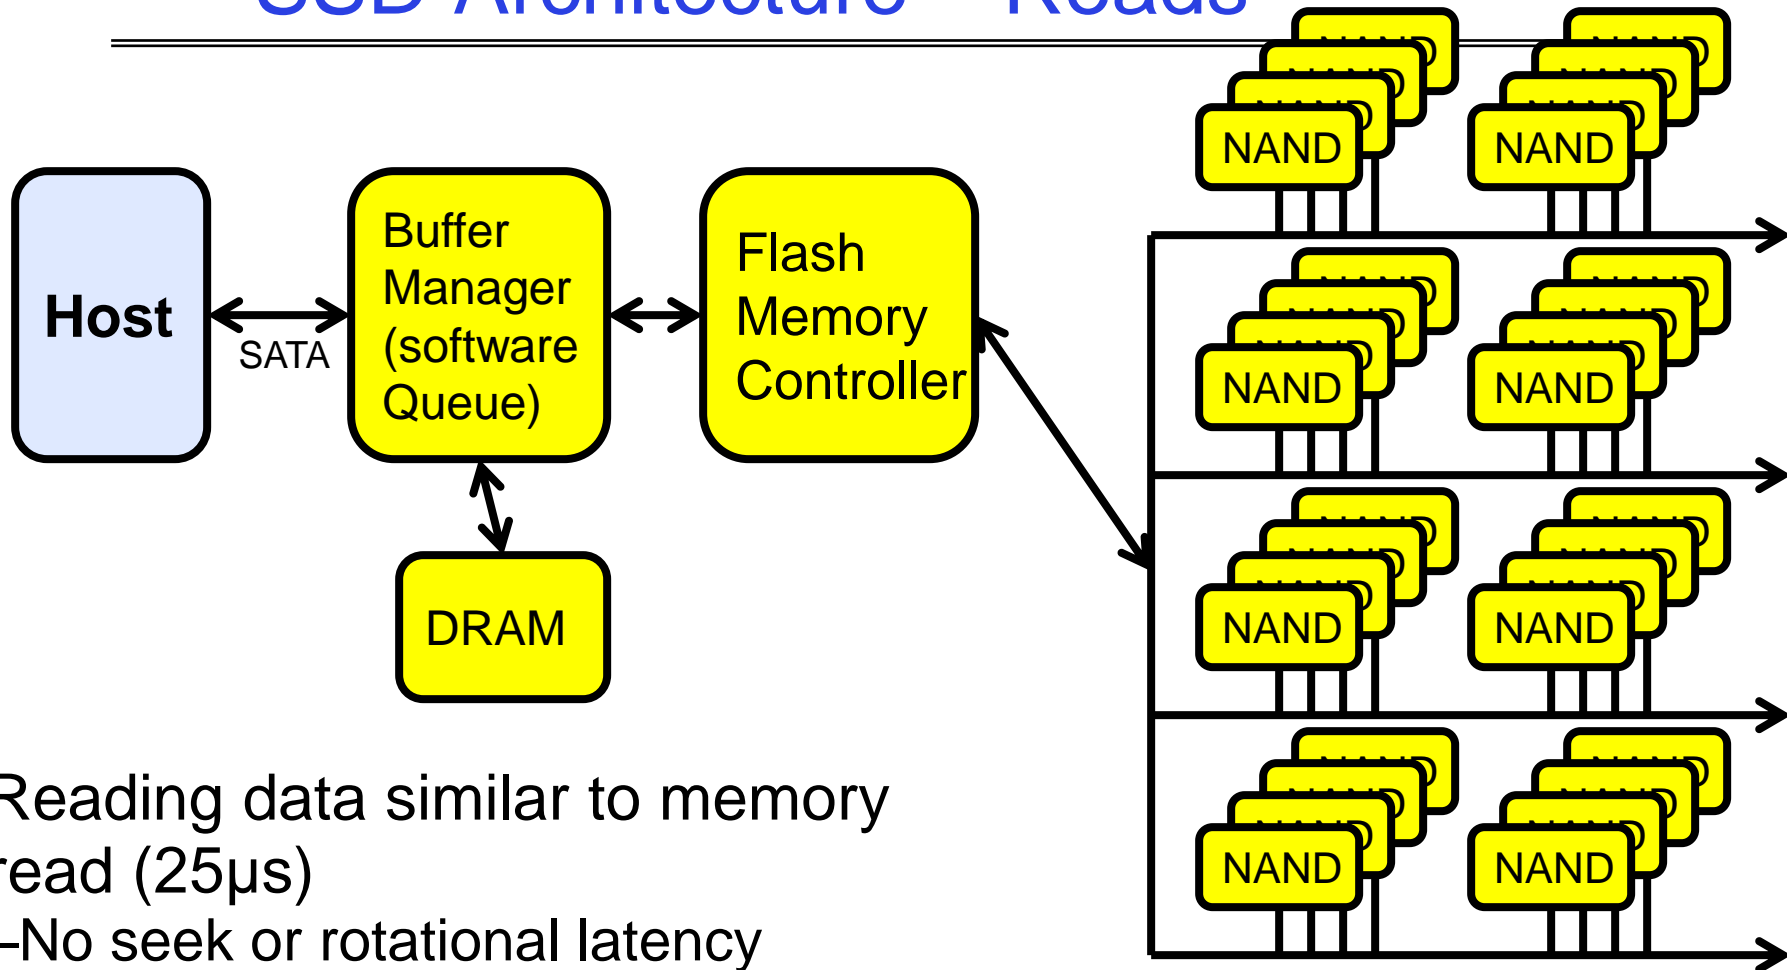
# Solid State Disks (SSDs)

---



- 1995 – Replace rotating magnetic media with non-volatile memory (battery backed DRAM)
- 2009 – Use NAND Multi-Level Cell (2-bit/cell) flash memory
  - Sector (4 KB page) addressable, but stores 4-64 “pages” per memory block
- No moving parts (no rotate/seek motors)
  - Eliminates seek and rotational delay (0.1-0.2ms access time)
  - Very low power and lightweight

# SSD Architecture – Reads



Reading data similar to memory read ( $25\mu\text{s}$ )

- No seek or rotational latency

- Transfer time: transfer a 4KB page

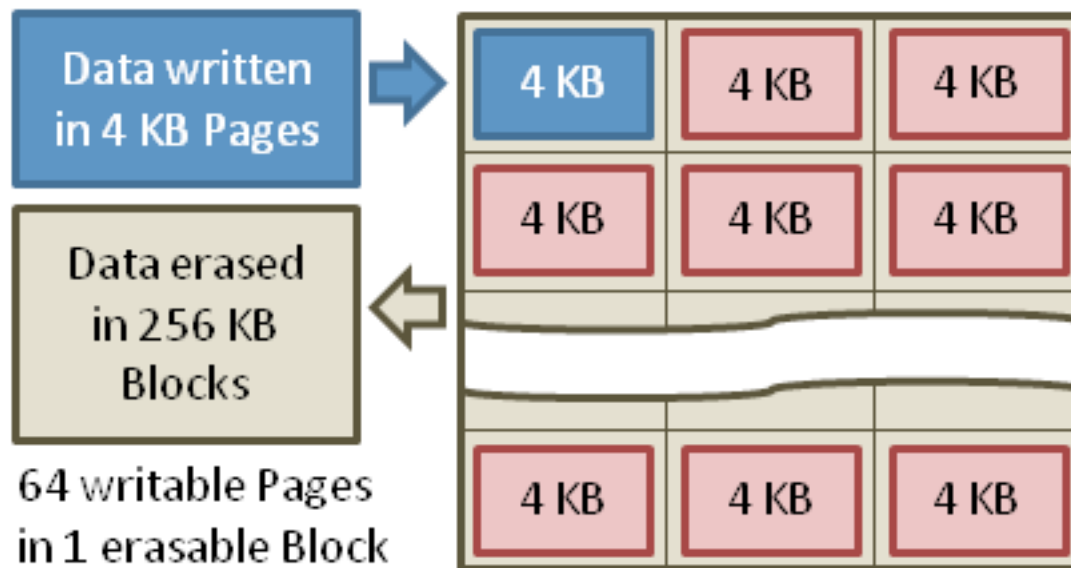
  - » Limited by controller and disk interface (SATA: 300-600MB/s)

- Latency = Queuing Time + Controller time + Xfer Time**

- Highest Bandwidth:** Sequential OR Random reads

# SSD Architecture – Writes (I)

- Writing data is complex! ( $\sim 200\mu\text{s}$  –  $1.7\text{ms}$ )
  - Can only write empty pages in a block
  - Erasing a block takes  $\sim 1.5\text{ms}$
  - Controller maintains pool of empty blocks by combining used pages (read, erase, write).




Typical NAND Flash Pages and Blocks

[https://en.wikipedia.org/wiki/Solid-state\\_drive](https://en.wikipedia.org/wiki/Solid-state_drive)

# SSD Architecture – Writes (II)

---

- Write A, B, C, D




[https://en.wikipedia.org/wiki/Solid-state\\_drive](https://en.wikipedia.org/wiki/Solid-state_drive)

# SSD Architecture – Writes (II)

---

- Write A, B, C, D
- Write E, F, G, H and A', B', C', D'
  - Record A, B, C, D as obsolete



[https://en.wikipedia.org/wiki/Solid-state\\_drive](https://en.wikipedia.org/wiki/Solid-state_drive)

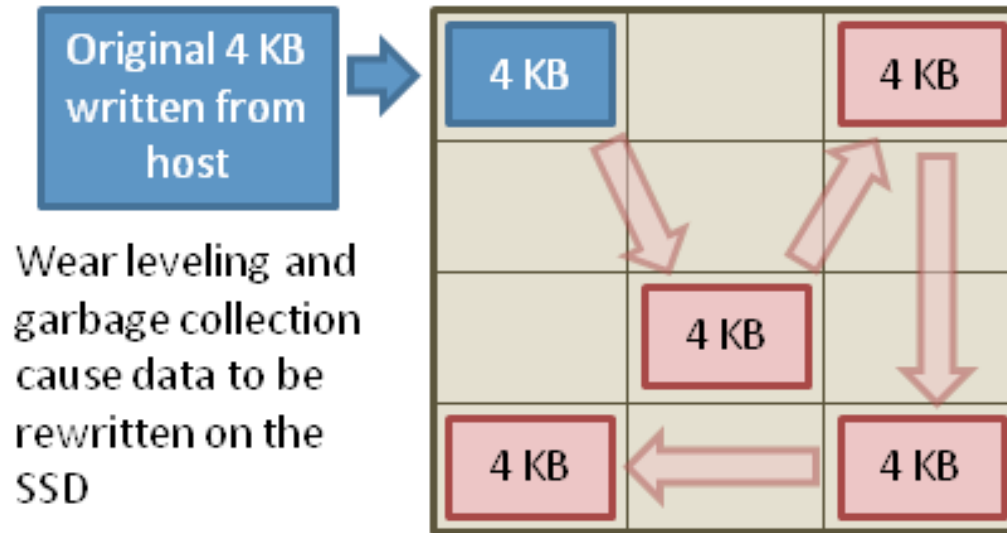
# SSD Architecture – Writes (II)

- Write A, B, C, D
- Write E, F, G, H and A', B', C', D'
  - Record A, B, C, D as obsolete
- Controller *garbage collects* obsolete pages by copying valid pages to new block
- Typical steady state behavior when SSD is almost full
  - One erase every 64 or 128 writes



## SSD Architecture – Writes (III)

- Write and erase cycles require “high” voltage
  - Damages memory cells, limits SSD lifespan
  - Controller uses error correction.



- Result is very workload dependent performance
  - Latency = Queuing Time + Controller time (Find Free Page) + Xfer Time
  - Highest BW: Seq. OR Random writes (limited by empty pages)

Rule of thumb: writes 10x more expensive than reads, and erases 10x more expensive than writes



# SSD Summary

---

- Pros (vs. hard disk drives):
  - Low latency, high throughput (eliminate seek/rotational delay)
  - No moving parts:
    - » Very light weight, low power (0.3x disk), silent, very shock insensitive
  - Read at memory speeds (limited by controller and I/O bus)
- Cons
  - Smaller storage (0.5x disk), expensive (7~10x disk)
    - » Hybrid alternative: combine small SSD with large HDD
  - Cannot update a single page in a block.
  - Asymmetric block write performance: read pg/erase/write pg
    - » Controller garbage collection (GC) algorithms have major effect on performance