



CSE150
Operating Systems
Lecture 22

Distributed Systems and Networking

Centralized vs Distributed Systems




Client/Server Model



Peer-to-Peer Model

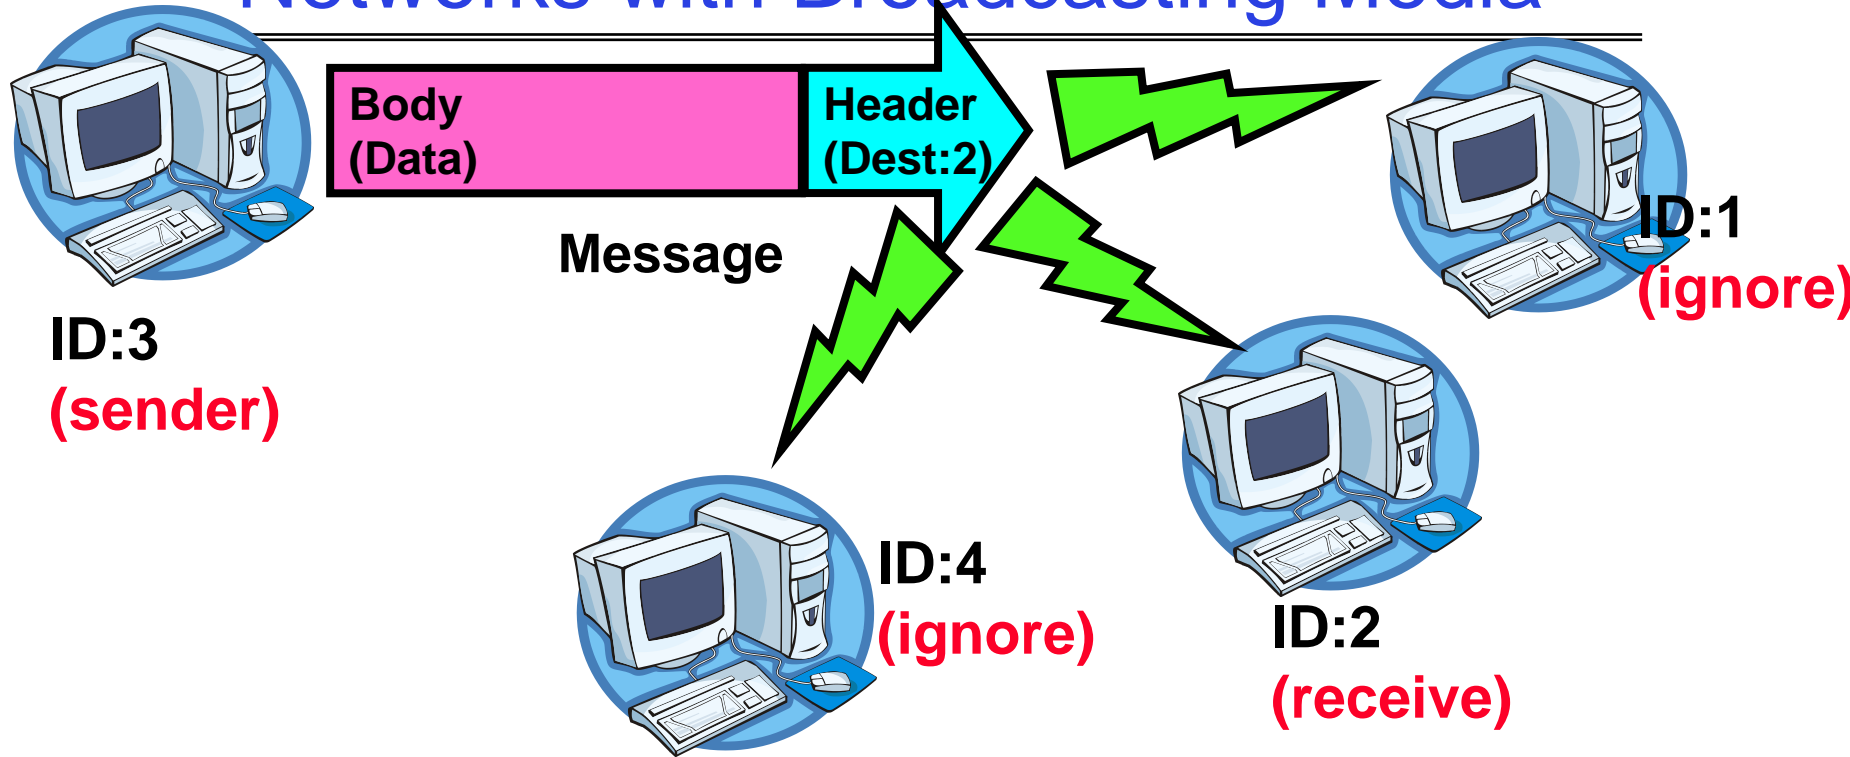
- **Centralized System:** System in which major functions are performed by a single physical computer
 - Originally, everything on single computer
 - Later: client/server model
- **Distributed System:** physically separate computers working together on some task
 - Early model: multiple servers working together
 - » Probably in the same room or building
 - » Often called a “cluster”
 - Later models: peer-to-peer/wide-spread collaboration

Networking Definitions



- **Network:** physical connection that allows two computers to communicate
- **Packet:** unit of transfer, sequence of bits carried over the network
 - Network carries packets from one CPU to another
 - Destination gets interrupt when packet arrives
- **Protocol:** agreement between two parties as to how information is to be transmitted

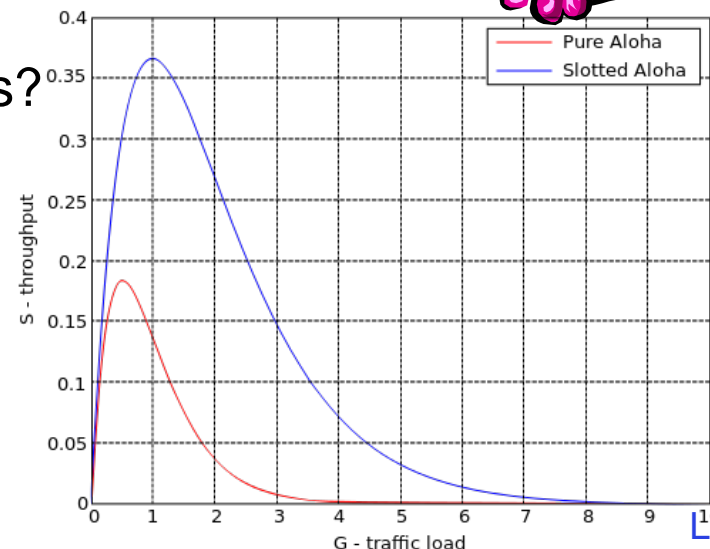

Networks with Broadcasting Media




- **Delivery:** When you broadcast a packet, how does a receiver know who it is for? (packet goes to everyone!)
 - Put header on front of packet: [Destination | Packet]
 - Everyone gets packet, discards if not the target
 - In Ethernet, this check is done in hardware
 - » No OS interrupt if not for particular destination

Network Arbitration

- **Multiple Access:** Act of negotiating use of shared medium
 - What if two senders try to broadcast at same time?
 - » Collision
- Aloha network (70's): packet radio within Hawaii
 - Blind broadcast, with **checksum** at end of packet. If received correctly (not garbled), send back an **acknowledgement**.
 - If ACK not received correctly, discard or re-transmits.
 - If two senders try to send at same time, both get garbled, both simply re-send **later**.
 - » Backoff a random duration.
 - Problem:
 - » Stability: what if load increases?
 - » Inefficiency




Networks with Point-to-Point Media




- Why have a shared bus at all? Why not simplify and only have point-to-point links + routers/switches?
 - Originally wasn't cost-effective
 - Now, easy to make high-speed switches and routers that can forward packets from a sender to a receiver.
- **Point-to-point network:** a network in which every physical wire is connected to only two computers
- **Switch:** a bridge that transforms a shared-bus (broadcast) configuration into a point-to-point network.
- **Router:** a device that acts as a junction between two networks to transfer data packets among them.

Hierarchical Networking: The Internet

- How can we build a network with millions of hosts?
 - Hierarchy! Not every host connected to every other one
 - Use a network of Routers to connect subnets together
 - » Routing is often by prefix: e.g. first router matches first 8 bits of address, next router matches more, etc.




The Problem



- Re-implement every application for every technology?
- No! But how does the Internet design avoid this?

Solution: Intermediate Layers

- Introduce intermediate layers that provide **set of abstractions** for various network functionality & technologies
 - A new app/media implemented only once



Software System Modularity

Partition system into modules & abstractions:

- Well-defined interfaces give flexibility
 - **Hides** implementation - thus, it can be freely changed
 - Extend functionality of system by adding new modules
- E.g., libraries encapsulating set of functionality
- E.g., programming language + compiler abstracts away not only how the particular CPU works ...

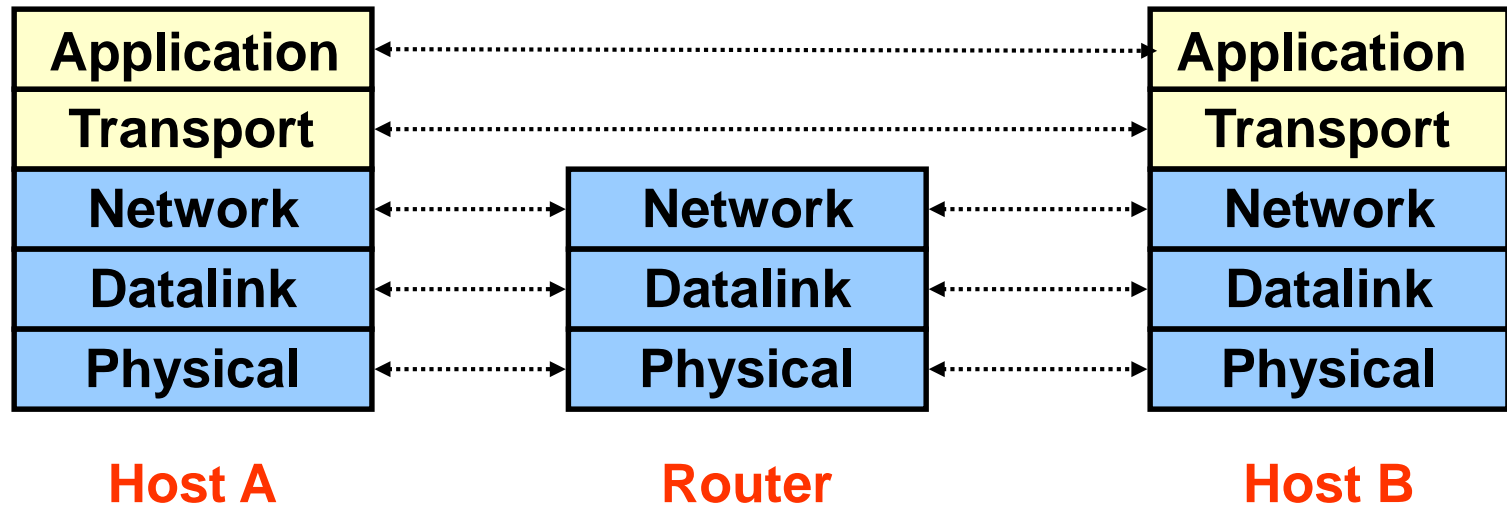
Network System Modularity

Like software modularity, but:

- Implementation distributed across many machines (routers and hosts)
- Must decide:
 - How to break system into modules:
 - » **Layering**

Internet Architecture: The Five Layers

- Lower three layers implemented everywhere
- Top two layers implemented only at hosts
- Logically, layers interact with peer's corresponding layer




Layering: A Modular Approach

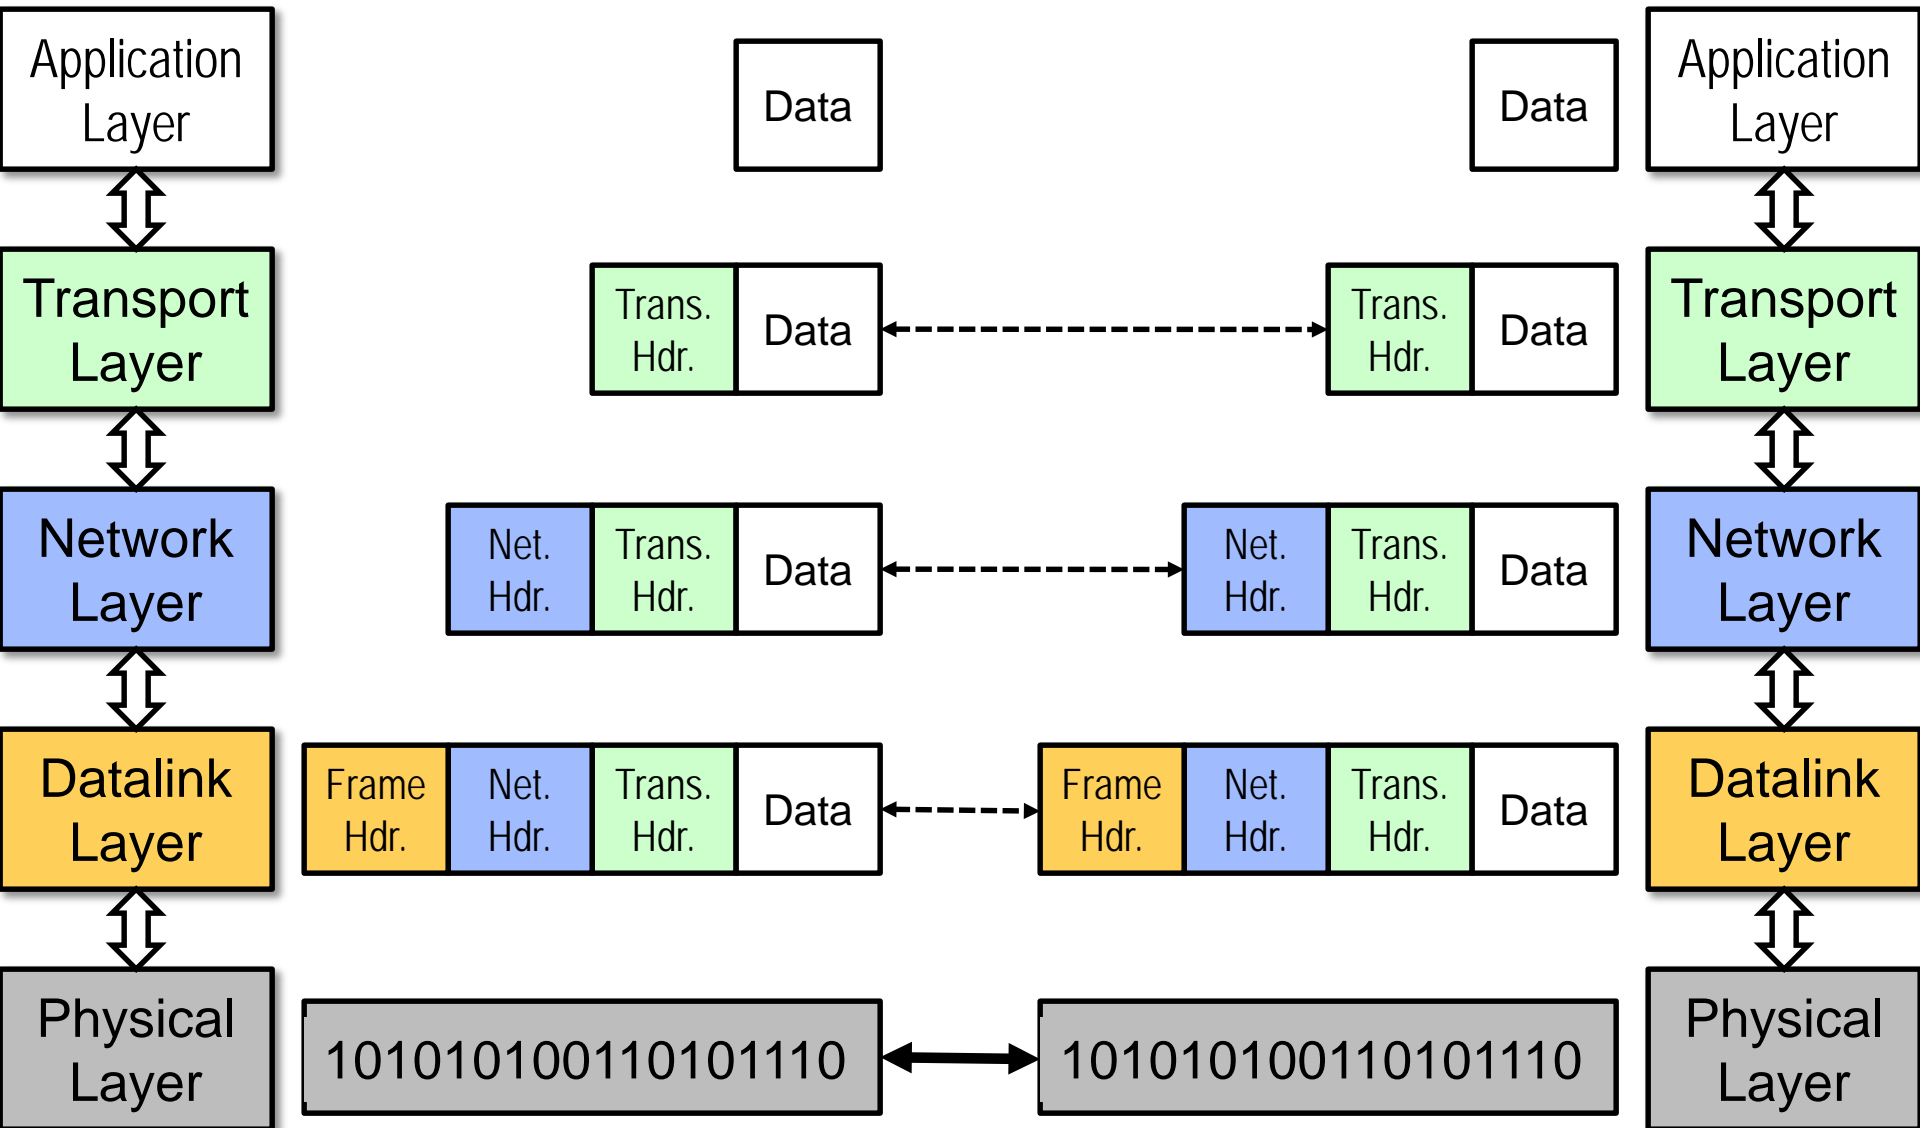
- Partition the system
 - Each layer **solely** relies on services from layer below
 - Each layer **solely** exports services to layer above
- Interface between layers defines interaction
 - Hides implementation details
 - Layers can change without disturbing other layers

Physical Communication

- Communication goes down to physical network
- Then from network peer to peer
- Then up to relevant layer



Layering: Packets in Envelopes

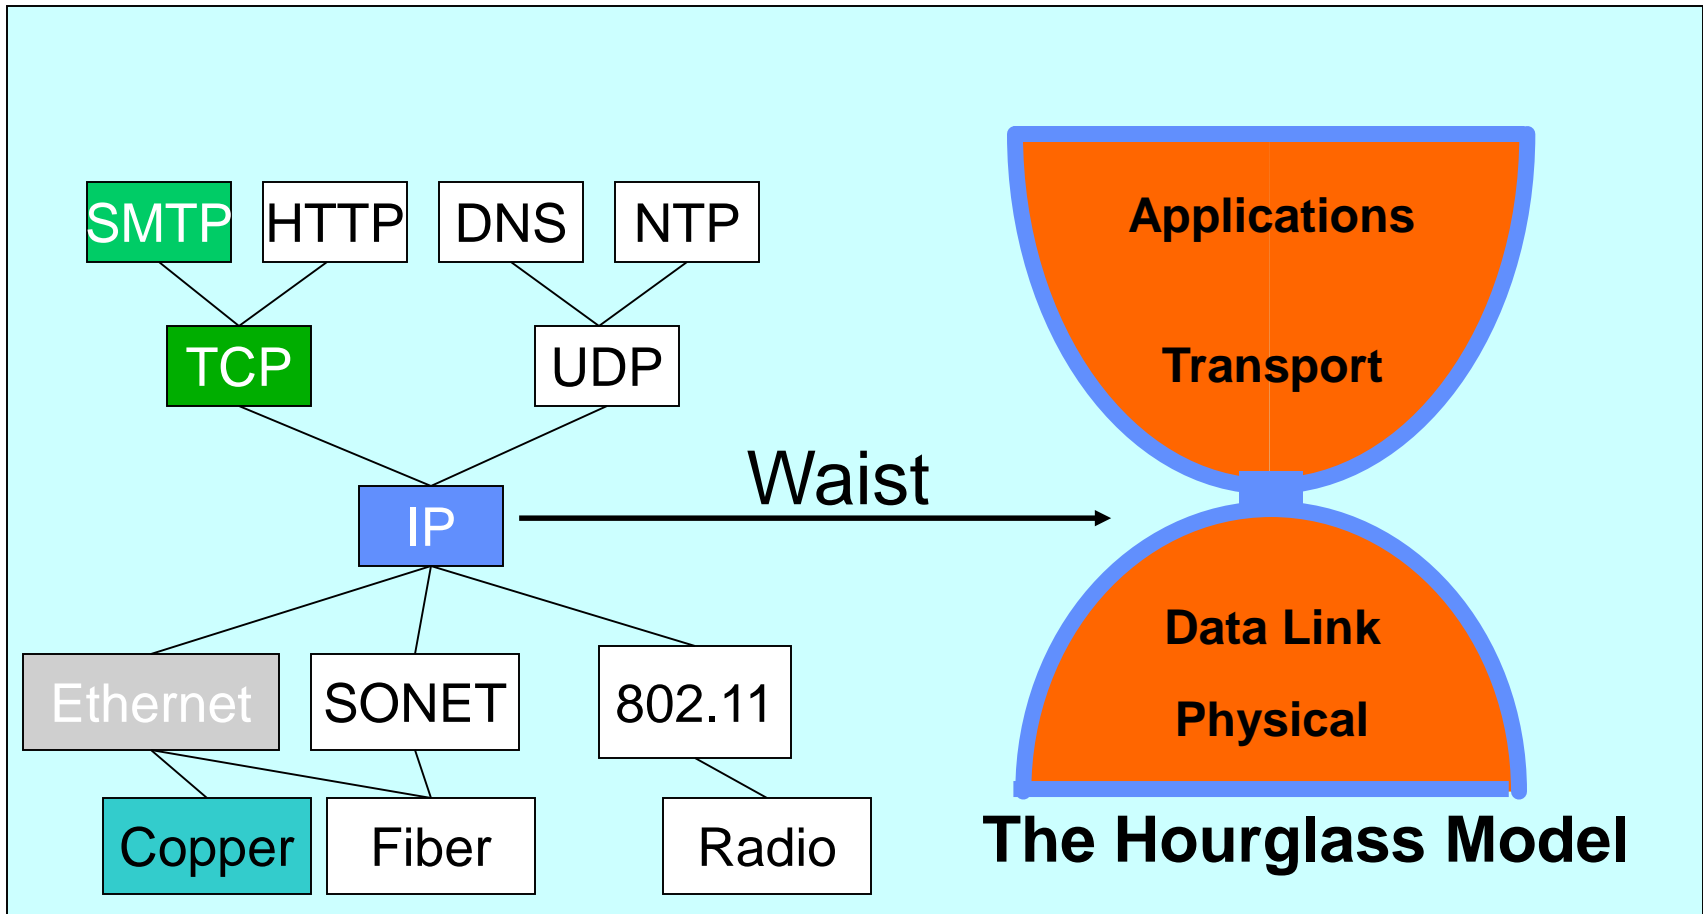


Application Layer (7 - not 5!)

Application
Present. Session
Transport
Network
Datalink
Physical

- **Service:** any service provided to the end user
- **Interface:** depends on the application
- **Protocol:** depends on the application
- Examples: Skype, SMTP (email), HTTP (Web), Halo, BitTorrent ...
- What happened to layers 5 & 6?
 - “Session” and “Presentation” layers
 - Part of **OSI** architecture, but not Internet architecture
 - Their functionality is provided by application layer

The Internet *Hourglass*




There is just **one** network-layer protocol, **IP**
The “narrow waist” facilitates **interoperability**

Internet Protocol (IP)

Application
Present. Session
Transport
Network
Datalink
Physical

- Internet Protocol: Internet's network layer
- Service it provides: “Best-Effort” Packet Delivery
 - Tries it's “best” to deliver one packet to its **destination**
 - Packets may be lost
 - Packets may be corrupted
 - Packets may be delivered out of order



Implications of Hourglass

Single Internet-layer module (**IP**):

- Allows arbitrary networks to interoperate
 - Any network technology that supports IP can exchange packets
- Allows applications to function on all networks
 - Applications that can run on IP can **use any network**
- Supports simultaneous innovations above and below IP

Transport Layer (4)

Application
Present. Session
Transport
Network
Datalink
Physical

- **Service:**
 - Provide end-to-end communication between **processes**
 - **Demultiplexing** of communication between hosts
 - Possible other services:
 - » **Reliability** in the presence of errors
 - » **Timing** properties
 - » **Rate adaption** (flow-control, congestion control)
- **Interface:** send message to “specific process” at given destination; local process receives messages sent to it
- **Protocol:** port numbers, perhaps implement reliability, flow control, packetization of large messages, framing
- Prime Examples: TCP (Transmission Control Protocol) and UDP (User Datagram Protocol)

Internet Transport Protocols

Application
Present.
Session
Transport
Network
Datalink
Physical

- Datagram service (**UDP**)
 - No-frills extension of “best-effort” IP
 - Multiplexing/Demultiplexing among processes
- Reliable, in-order delivery (**TCP**)
 - Connection set-up & tear-down
 - Discarding corrupted packets (segments)
 - Retransmission of lost packets (segments)
 - Flow control
 - Congestion control
- Services **not available**
 - Delay and/or bandwidth guarantees
 - Sessions that survive change-of-IP-address

TCP Service


- 1) Open connection: 3-way handshaking
- 2) Reliable byte stream transfer from (IPa, TCP_Port1) to (IPb, TCP_Port2)
 - Indication if connection fails: Reset
- 3) Close (tear-down) connection

Open Connection: 3-Way Handshaking

- Goal: agree on a set of parameters, i.e., the start sequence number for each side
 - Starting sequence number: sequence of first byte in stream
 - Starting sequence numbers are random

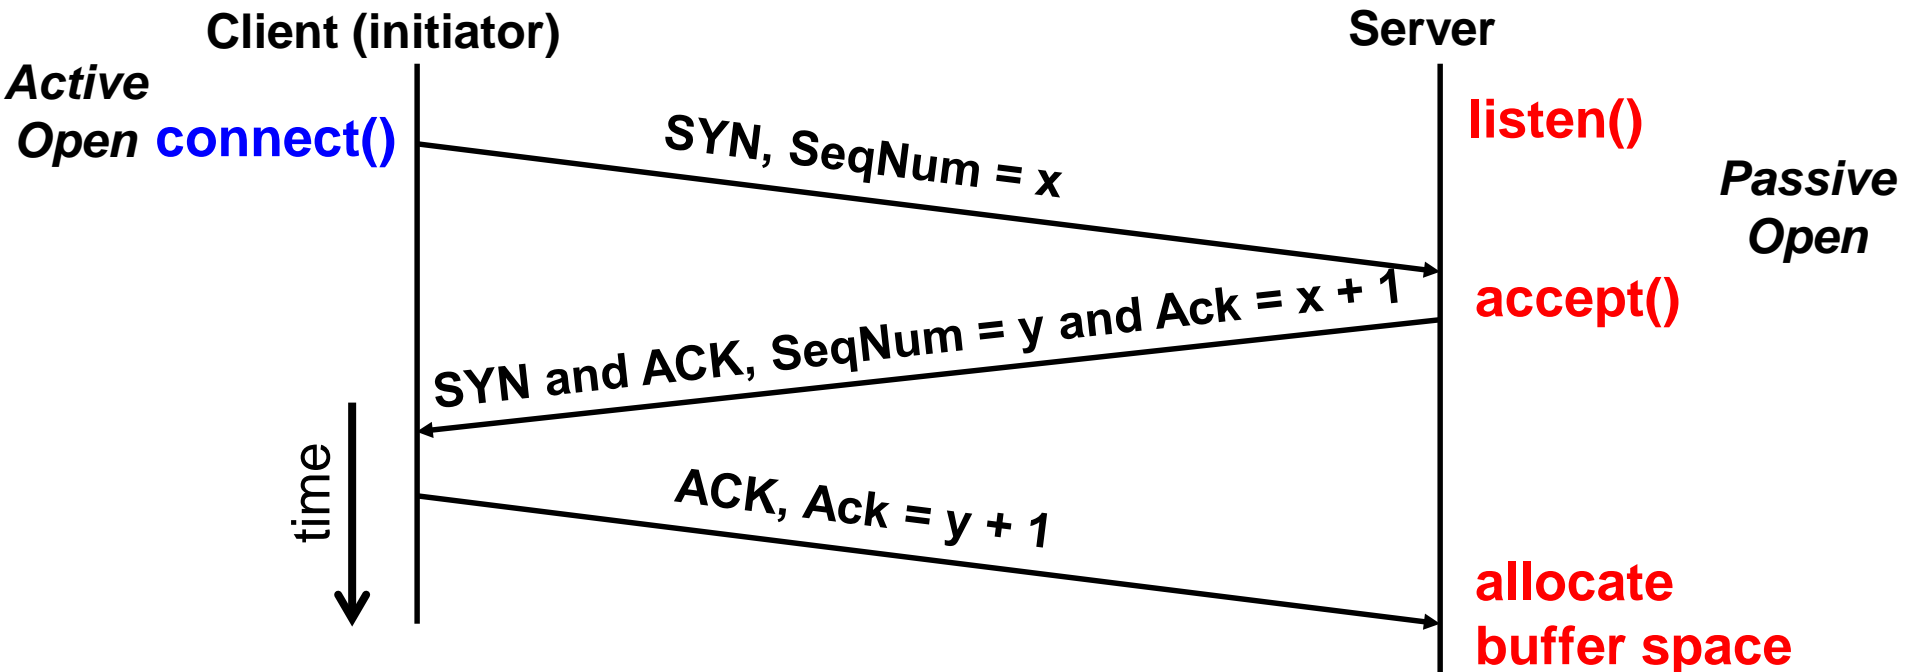
Open Connection: 3-Way Handshaking

- Server waits for new connection calling **listen()**
- Sender call **connect()** passing socket which contains server's IP address and port number
 - OS sends a special packet (SYN) containing a proposal for first sequence number, x



Open Connection: 3-Way Handshaking

- If it has enough resources, server calls **accept()** to accept connection, and sends back a SYN ACK packet containing
 - Client's sequence number incremented by one, $(x + 1)$
 - A sequence number proposal, y , for first byte server will send

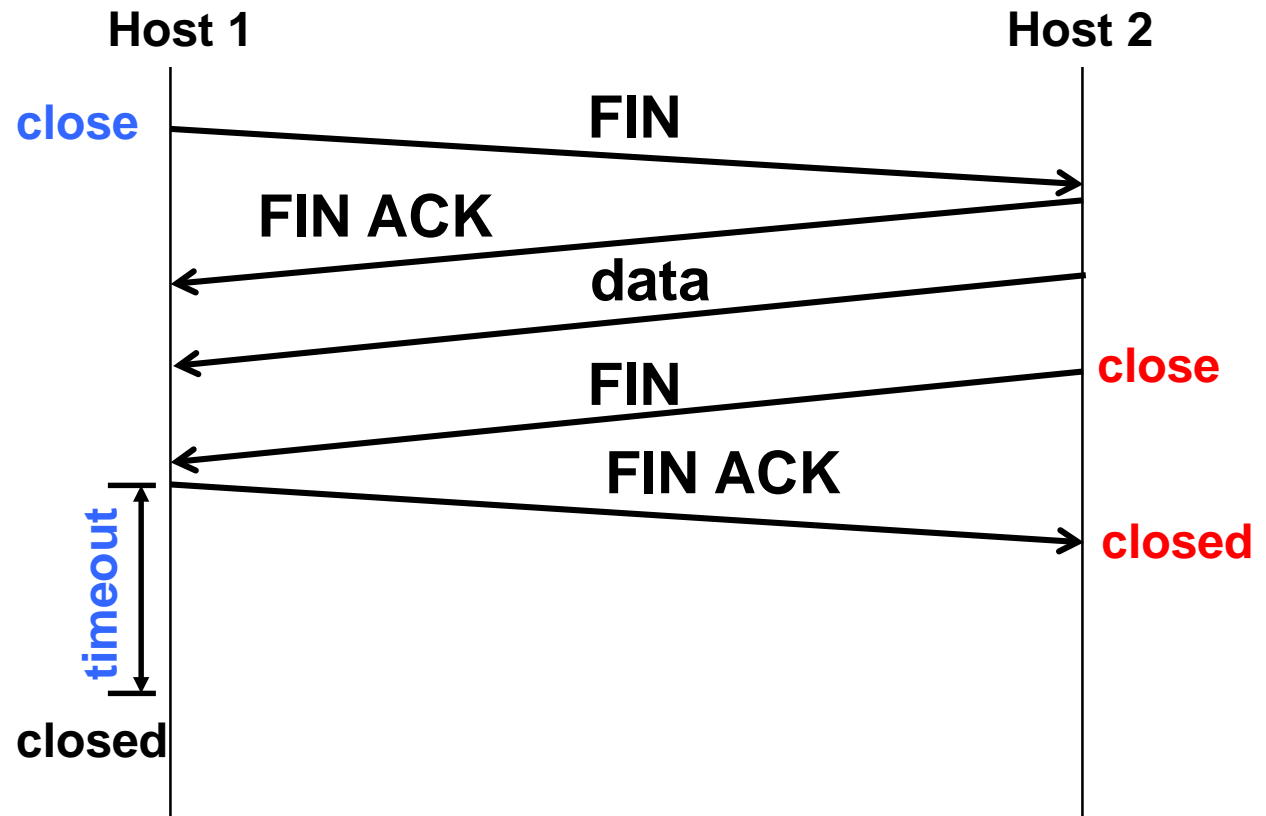


3-Way Handshaking (cont'd)

- Three-way handshake adds 1 RTT delay
- Why?
 - Full-duplex transmissions
 - Congestion control: SYN (40 byte) acts as cheap probe

Close Connection


- Goal: both sides agree to close the connection
- 4-way connection tear down



- Can retransmit **FIN ACK** if it is lost
- Protects against delayed packets from other connection

Socket API

- Base level Network programming interface



BSD Socket API

- Created at UC Berkeley (1980s)
- Most popular network API
- Ported to various OSes, various languages
 - Windows Winsock, BSD, OS X, Linux, Solaris, ...
 - Socket modules in Java, Python, Perl, ...
- Similar to Unix file I/O API
 - In the form of *file descriptor* (sort of handle).
 - Can share same `read()`/`write()`/`close()` system calls

Sockets in concept

Client

Create Client Socket



Connect it to server (host:port)



write request

read response



Close Client Socket

Server

Create Server Socket



Bind it to an Address (host:port)



Listen for Connection



Accept connection

Connection Socket



read request


write response



Close Connection Socket



Close Server Socket



Summary

- Centralized vs Distributed Systems
- Network layering